NEW PROBLEMS OF PHILOSOPHY

ROUTLEDGE

SKEPTICISM

ANNALISA COLIVA AND DUNCAN PRITCHARD

SKEPTICISM

Skepticism is one of the perennial problems of philosophy: from antiquity, to the early modern period of Descartes and Hume, and right through to the present day. It remains a fundamental and widely studied topic and, as Annalisa Coliva and Duncan Pritchard show in *Skepticism*, it presents us with a paradox with important ramifications not only for epistemology but also for many other core areas of philosophy.

This book provides a thorough grounding in contemporary debates about skepticism, exploring the following key topics:

- the core skeptical arguments, with a particular focus on Cartesian and Humean radical skepticism
- the epistemic principles that are held to underlie skeptical arguments, such as the Closure and Underdetermination principles
- the content externalism of Putnam, Davidson, and Chalmers, and how it might help us respond to radical skepticism
- the epistemic externalism/internalism distinction and how it relates to the skeptical problematic
- contextualism in epistemology and its anti-skeptical import
- the various interpretations of a Wittgensteinian hinge epistemology
- the viability of epistemological disjunctivism, including whether it can be combined with hinge epistemology as part of a dual response to radical skepticism
- liberal and conservative responses to the Humean skeptical paradox.

Both authors are prominent figures who work on skepticism, and so one novelty of the book is that it provides an insight into their own contrasting responses to this philosophical difficulty. With the addition of annotated further reading and a glossary, this is an ideal starting point for anyone studying the philosophy of skepticism, along with students of epistemology, metaphysics, and contemporary analytic philosophy.

Annalisa Coliva is Full Professor, Chancellor Fellow, and Chair of the Department of Philosophy at the University of California, Irvine, USA. Her books include *Moore and Wittgenstein: Scepticism, Certainty and Common Sense* (2010), *Extended Rationality: A Hinge Epistemology* (2015), and *The Varieties of Self-Knowledge* (2016). With Maria Baghramian she is the author of *Relativism* (Routledge, 2019).

Duncan Pritchard is Distinguished Professor of Philosophy and Director of the Center for Knowledge, Technology & Society at the University of California, Irvine, USA. His books include *Epistemic Luck* (2005), *The Nature and Value of Knowledge* (with Millar and Haddock, 2010), *Epistemological Disjunctivism* (2012), *Epistemic Angst* (2015), *Scepticism: A Very Short Introduction* (2019), and *What Is This Thing Called Knowledge?* (4th edn, Routledge, 2018).

Taylor & Francis
Taylor & Francis Group
http://taylorandfrancis.com

NEW PROBLEMS OF PHILOSOPHY

Series Editor: José Luis Bermúdez

"New Problems of Philosophy is developing a most impressive line-up of topical volumes aimed at graduate and upper-level undergraduate students in philosophy and at others with interests in cutting edge philosophical work. Series authors are players in their respective fields and notably adept at synthesizing and explaining intricate topics fairly and comprehensively."

— *John Heil, Monash University, Australia, and Washington University in St. Louis, USA*

"This is an outstanding collection of volumes. The topics are well chosen and the authors are outstanding. They will be fine texts in a wide range of courses."

— *Stephen Stich, Rutgers University, USA.*

The *New Problems of Philosophy* series provides accessible and engaging surveys of the most important problems in contemporary philosophy. Each book examines either a topic or theme that has emerged on the philosophical landscape in recent years, or a longstanding problem refreshed in light of recent work in philosophy and related disciplines. Clearly explaining the nature of the problem at hand and assessing attempts to answer it, books in the series are excellent starting points for undergraduate and graduate students wishing to study a single topic in depth. They will also be essential reading for professional philosophers. Additional features include chapter summaries, further reading and a glossary of technical terms.

The Metaphysics of Identity
André Gallois

Consciousness
Rocco J. Gennaro

Abstract Entities
Sam Cowling

Embodied Cognition, Second Edition
Lawrence Shapiro

Self-Deception
Eric Funkhouser

Relativism
Maria Baghramian and Annalisa Coliva

Empathy
Heidi Maibom

Moral Disagreement
Richard Rowland

Perception
Adam Pautz

Thinking and Perceiving
Dustin Stokes

Emotion
Charlie Kurth

Skepticism
Annalisa Coliva and Duncan Pritchard

For more information about this series, please visit: https://www.routledge.com/New-Problems-of-Philosophy/book-series/NPOP

SKEPTICISM

Annalisa Coliva and
Duncan Pritchard

Routledge
Taylor & Francis Group
LONDON AND NEW YORK

Cover image: © skyNext / Getty images

First published 2022
by Routledge
4 Park Square, Milton Park, Abingdon, Oxon OX14 4RN

and by Routledge
605 Third Avenue, New York, NY 10158

Routledge is an imprint of the Taylor & Francis Group, an informa business

Library of Congress Cataloguing-in-Publication Data
Names: Coliva, Annalisa, 1973- author. | Pritchard, Duncan, author.
Title: Skepticism / Annalisa Coliva and Duncan Pritchard.
Description: Abingdon, Oxon ; New York, NY : Routledge, 2022. |
Includes bibliographical references and index.
Identifiers: LCCN 2021046262 (print) |
LCCN 2021046263 (ebook) | ISBN 9780367178314 (hardback) |
ISBN 9780367178352 (paperback) | ISBN 9780429057946 (ebook)
Subjects: LCSH: Skepticism.
Classification: LCC B837 .C65 2022 (print) | LCC B837 (ebook) |
DDC 149/.73--dc23
LC record available at https://lccn.loc.gov/2021046262
LC ebook record available at https://lccn.loc.gov/2021046263

ISBN: 978-0-367-17831-4 (hbk)
ISBN: 978-0-367-17835-2 (pbk)
ISBN: 978-0-429-05794-6 (ebk)

DOI: 10.4324/9780429057946

Typeset in Joanna and Scala Sans
by MPS Limited, Dehradun

CONTENTS

ACKNOWLEDGMENTS

We would like to thank Adam Carter and an anonymous referee for extensive comments on an initial draft of the book. We would also like to thank Tony Bruce and Adam Johnson at Routledge for making this project possible.

Taylor & Francis
Taylor & Francis Group
http://taylorandfrancis.com

INTRODUCTION

Skeptical doubt about whether we know as much as think we do is as old as philosophy itself, and under its various guises it has been a hugely influential force throughout the history of ideas. In order to make such a sprawling topic manageable for a book-length treatment, our focus will be on the contemporary debate regarding an extreme form of skepticism known as *radical skepticism*. In particular, we will be looking at two kinds of radical skeptical argument that have been especially influential. While both of them have historical antecedents, our interest will be in their expression in the present day.

In line with this emphasis on the skeptical problem as it is manifest in the contemporary literature, our target will specifically be radical skepticism about our knowledge of the physical world that we occupy, what is known as 'external world' skepticism. How do we know that we live in a world of tables and chairs and other mundane physical objects, of a kind that we seem to interact with on a regular basis? Interestingly, skepticism of this kind can also be extended to (at least aspects of) our bodily selves, as the skeptical worries that challenge our knowledge of, say, chairs apply with just as much force when it comes to our knowledge that—to take a well-trodden example—we have hands. Skepticism has had many other targets besides our knowledge of the external world, such as our knowledge of the existence of other minds beside one's own, or the existence of

DOI: 10.4324/9780429057946-1

the past, or of moral values. Nonetheless, external world skepticism can surely lay claim to being by far the most discussed skeptical difficulty in contemporary debate. Moreover, getting clear about its structure and possible responses should help us to clarify and eventually respond to distinct forms of skepticism that target other areas of human thought.

We follow the contemporary literature in primarily conceiving of the skeptical challenge not as a philosophical position occupied by a real-life skeptic, but rather as a *paradox*. Paradoxes arise when there is, it seems, a deep tension within our ordinary concepts within a certain domain. The paradox of free will, for example, concerns how our ordinary ways of thinking about what free will is, and about relevant features of the world around us (such as how one event might cause another), seem to entail that genuine freedom of the will is impossible, even despite our strong conviction that we do have it. In presenting the free will puzzle as a paradox, one is not thereby claiming that we lack free will, only that there is a deep tension in our thinking about free will. Denying that we have free will may be one *response* to the paradox of free will, but whatever response one offers, the fact remains that there seems to be something inherently paradoxical about our ordinary ways of thinking in this regard. Relatedly, even if one holds that there is something intrinsically contradictory about *defending* the claim that one lacks free will (as defending a thesis surely presupposes that one has free will), this wouldn't entail that there is nothing paradoxical about our ordinary ways of thinking about free will (only that there is a particular difficulty inherent to a specific response to that paradox).

Similar points hold if we think of radical skepticism as a paradox. *Qua* paradox, radical skepticism purports to be exposing deep tensions within our own epistemological concepts, pre-theoretically understood. This means that the skeptical conclusions it generates arise out of premises and inferential principles that seem entirely compelling, as they are rooted in our own epistemic practices. That is, if radical skepticism really is the paradox that it purports to be, then it arises out of commitments that we *all* share, rather than being the product of particular theoretical commitments held by a skeptical adversary. In this sense, the radical skeptical paradox is a problem for all of us. Moreover, notice that radical skepticism *qua* paradox is importantly different to radical skepticism *qua* position. As with responding to the free will paradox by denying free will, responding to the radical skepticism paradox by embracing the radical skeptical claim that we

don't know very much is just one way of responding to the paradox (albeit the most negative response). Similarly, even if it is true that there is something self-defeating about actually embracing the radical skeptical conclusion (just as there might be something self-defeating about claiming that one lacks free will), it wouldn't follow that the paradox of radical skepticism is thereby answered (any more than a parallel observation regarding free will answers the free will paradox).

Presenting skepticism as a paradox, rather than as a position, thus has some methodological advantages, as it enables us to disentangle the problem in general from whatever difficulties face a particular, skeptical, response to that problem. Besides, it should be conceded that actual radical skeptics have not been 'thick on the ground', either throughout the history of philosophy or in present-day philosophy (though as we will note in Chapter 1, there is a rich tradition of 'lived' skepticism, albeit of a less radical variety, that extends from the Pyrrhonian skeptics, and taking in such important figures as Hume and Montaigne). Moreover, although skepticism has been taken seriously by a range of important philosophers, both historically and in the contemporary literature, their number has always been dwarfed by those who have treated skepticism as something that can easily be dismissed, or which deserves to be ignored altogether. Nonetheless, the skeptical paradox, in its various guises, proves incredibly difficult to handle once its structure is articulated, as we show, and so is not so easily set to one side.

Relatedly, focusing too much on the dialectic between skeptics and their opponents—for example, maintaining the inescapability of our belief in the existence of an external world, or insisting that the skeptic incurs a significant 'burden of proof'—risks downplaying the philosophical interest and relevance, as well as difficulty, of the skeptical paradox. By contrast, focusing on the discussion of the premises and the inferential principles at play in the derivation of the skeptical paradox proves much more instructive, since it mobilizes our most deep-seated intuitions and beliefs regarding their *prima facie* acceptability. Whenever philosophy forces us to confront such fundamental issues, it is bound to produce results that reverberate across our philosophical thinking more generally—for example, by requiring revisions with respect to our deep-seated beliefs and inferential practices, or by suggesting subtle distinctions in the way we pre-theoretically or theoretically conceive of key notions, such as knowledge, justification, certainty, and so on.

In particular, we differentiate between responses to the skeptical paradox in terms of whether they attempt to *undercut* it or *override* it. Undercutting proposals aim to show that there is in fact no paradox in play, despite appearances to the contrary. Our basic epistemological commitments, at least when properly understood, do not generate all the elements needed to generate the paradox. Overriding anti-skeptical strategies, in contrast, are more concessive, in that they allow that there is a fundamental tension in our epistemological concepts, and thus that the paradox is genuine. Nonetheless, they offer revisionary proposals that enable us to resist the radical skeptical conclusion that the paradox generates.

We argue in Chapter 1 that the external world skeptical paradox comes in various forms. We focus on two versions that have been particularly influential in contemporary epistemology and explore their underlying structure. We refer to these formulations of the skeptical paradox as 'Cartesian' and 'Humean,' but despite their names, we make no claim that such arguments are accurate renditions of what Descartes and Hume respectively held. (Note that we are also not claiming that they are completely distinct forms of skepticism either).

The Cartesian skeptical paradox makes essential appeal to radical skeptical scenarios. These are logically and metaphysically possible scenarios, indistinguishable from ordinary life, where most of our beliefs are false, such as the possibility that one is a disembodied brain-in-a-vat (BIV) being 'fed' deceptive experiences about the world. Being indistinguishable from normal experiences, the challenge posed by radical skeptical scenarios is that we do not seem to have any means of excluding them.

The Cartesian skeptical paradox can be formulated in two different ways. The first turns on the *Closure principle* for knowledge (or justification), according to which, roughly, if one knows one proposition (such as that one has hands), and one knows that it entails a second proposition (e.g., that one is not a handless BIV), then one also knows that second proposition. The second turns on the *Underdetermination principle*, according to which in order to have knowledge of a proposition (such as that one has hands), one must have epistemic support for one's belief in that proposition which favors it over known alternatives (such as that one is not a handless BIV). Either way, the general idea is that our inability to exclude radical skeptical scenarios seems to generate the conclusion that we are unable to have much of the knowledge of the world around us that we normally credit to ourselves.

The Humean skeptical paradox, in contrast, does not necessarily appeal to radical skeptical scenarios (at least not explicitly anyway). The problem it raises instead primarily concerns the apparent *epistemic circularity* involved in arguments designed to produce knowledge (or justified belief) that there is an external world. In particular, the claim is that any form of reasoning we might employ in order to derive this anti-skeptical conclusion would already presuppose that this conclusion was known, and hence would be epistemic circular. If that's right, then it seems there simply is no way of reasoning our way to the knowledge of the external world.

In Chapter 2, we examine semantic and epistemic externalist responses to radical skepticism. The former kind of response aims to use content externalism to challenge the coherence of the kind of radical error-possibility that Cartesian skepticism appeals to. Hilary Putnam offers a version of this line that maintains that one cannot think the thought that one is a BIV without thereby thinking falsely (because if one were a BIV, then one wouldn't be able to think a thought with *that* content at all). For similar reasons, Donald Davidson maintains that our beliefs are in their nature veridical, in that to have a system of contentful beliefs at all entails that these beliefs must be mostly true. We also consider in this regard a recent proposal made by David Chalmers. Although motivated by similar considerations, it generates a markedly different conclusion. It claims that even if we were the victim of a radically skeptical scenario, we would still retain much of our knowledge regardless, as our beliefs would then have very different contents. For example, we are wrong to suppose that the BIV has mostly false beliefs; it is rather that his beliefs have a very different content to our (non-BIV) beliefs, such that they are true regardless. For example, his beliefs about his 'hands' are generally true as they are concerned with the kind of 'vat-hands' possessed by a BIV. We argue that the kind of intellectual comfort provided by these semantic approaches to radical skepticism is, on closer inspection, rather insubstantial.

We then turn to a certain kind of epistemic externalist response to radical skepticism. Like content externalism, this also challenges the role of radical skeptical scenarios in Cartesian skepticism, though it does so more directly, by insisting that we can satisfy the epistemic conditions for knowing that they are false (i.e., rather than appealing to claims about meaning in this regard). In particular, it maintains that the appeal of thinking that we are unable to have such anti-skeptical knowledge essentially trades on epistemic internalist claims that we should reject.

We examine the prospects of such a response, and find them wanting. In particular, we pose a dilemma for this proposal. The first horn of the dilemma is to opt for a moderate form of epistemic externalism which is not highly revisionary (and hence plausible), but which is thereby also ineffective against radical skepticism. The second horn of the dilemma is to embrace an extreme form of epistemic externalism. While this can be used to block radical skepticism, it is also so revisionary of our epistemic concepts that this is a hollow victory, as the 'knowledge' we regain from the radical skeptic is so distinct from (and inferior to) that which we were trying to defend from skeptical attack. Either way, such epistemic externalist responses to radical skepticism are not found to be compelling.

In Chapter 3, we consider the principle of Closure for epistemic operators under known entailment, which is a key principle in the derivation of one form of the Cartesian skeptical paradox. We look at the criticisms of this principle raised by Fred Dretske and Robert Nozick. These are based on accounts of knowledge which have been developed in response to Edmund Gettier's attack on the tripartite conception of knowledge, according to which knowledge consists in having true and justified—that is, rationally supported—beliefs. We are critical of these revisionary conceptions of knowledge, yet we think that discussion of the status of the Closure principle is an important component of any considered response to the skeptical paradox. Accordingly, we distinguish various possible interpretations of it, some of which are more plausible than others, and examine their relationship with another epistemic inferential principle in the vicinity—*viz.*, the principle of transmission of justification from justified premises to entailed conclusions. Upon scrutiny, we maintain that, whatever the fate of Closure a response is owed to the challenge that its denial will entail a commitment to 'abominable conjunctions,' such as 'I know I have a hand, but I don't know that I am not a handless brain in a vat.'

In the same chapter, we look at semantic contextualist responses, put forward in recent years by Stewart Cohen, Keith DeRose, and David Lewis. This maintains that the truth of knowledge ascriptions can vary from context to context—including, crucially, as a result of whether skeptical alternatives are considered—depending on which contextually relevant epistemic standards are in play. Thus, in low standards contexts where such alternatives are not at issue, knowledge ascriptions can express truths. In contrast, in high standards contexts where skeptical alternatives are in play, these very same knowledge ascriptions would be false. Yet, in both contexts, closure would

be preserved (since there is no single context where a knowledge ascription of the entailing proposition expresses a truth while a knowledge ascription of the entailed proposition expresses a falsehood). We raise several worries in connection with this revisionary proposal regarding 'knowledge,' and cast doubts on the plausibility of coherently applying this contextualist line of argument to the skeptical paradox.

In Chapter 4, we turn to *hinge epistemology*. This is a growing trend in contemporary epistemology that builds on Wittgenstein's insights in his final notebooks, published as *On Certainty* (1969). After introducing the general themes behind hinge epistemology, we then turn to consider a version of it that has been put forward by Duncan Pritchard, and which targets the Closure-based version of the Cartesian skeptical paradox (we will be examining some other types of hinge epistemology in Chapters 6 and 7, once we have covered some key themes in contemporary treatments of radical skepticism). Key to this proposal is the idea that any system of rational evaluation presupposes a backdrop of basic certainties—the *hinge commitments*—which are themselves immune to rational evaluation (and so are arationally held). For Pritchard, Wittgenstein's myriad hinges are all in fact manifestations of an overarching *über hinge commitment* that we cannot be massively and fundamentally mistaken about all our beliefs (as we would be if we were the victim of a Cartesian skeptical scenario). Moreover, Pritchard argues that the distinctive nature of the propositional attitude involved in our hinge commitments is not one of belief, at least in the sense of belief that is relevant for our purposes (i.e., that propositional attitude that is a constituent part of rationally grounded knowledge). This allows Pritchard to defend the compatibility of his response to Cartesian skepticism with the validity of closure, understood as a principle that would generate knowledge, and thus belief (in the relevant sense), of the conclusion. One lacks knowledge of the denials of skeptical hypotheses, as we are hinge committed to them, but equally as hinge commitments one cannot come to have a belief in them via closure-based reasoning. Such commitments are thus not counterexamples to the Closure principle, properly understood, but rather fall outside of its scope.

In Chapter 5, we turn to *epistemological disjunctivism*, as originally articulated by John McDowell and recently developed in detail by Pritchard. This proposal draws on earlier work by a number of figures on *metaphysical disjunctivism* regarding perceptual experience, whereby the content of one's perceptual experiences is fundamentally different in 'good cases' involving

ordinary veridical perception as opposed to indistinguishable 'bad cases,' where one is radically deceived (such as radical skeptical scenarios). Epistemological disjunctivism advances an epistemic counterpart of this claim. In particular, it maintains that one's perceptual belief can enjoy factive rational support in the good case (i.e., rational support that entails the believed proposition), even though one's belief in the corresponding bad case can at most enjoy non-factive rational support. Moreover, this factive rational support is held to be reflectively accessible, and hence satisfies the rubric of epistemic internalism. Epistemological disjunctivism is thus able to offer a direct response to the underdetermination-based version of the Cartesian skeptical paradox, in that such factive reasons offer decisive favoring rational support for one's everyday perceptual beliefs over radical skeptical alternatives.

Epistemological disjunctivism can also in principle offer a direct response to the Closure-based formulation of the Cartesian skeptical paradox too, as one could maintain that such a decisive favoring rational basis ought to suffice for knowing the denials of radical skeptical scenarios. Pritchard argues that there might be grounds to resist this further anti-skeptical claim, however, and instead canvasses support for what he describes as a *biscopic* proposal that weaves together epistemological disjunctivism with hinge epistemology. In particular, the essential idea is that such a combination would be able to entwine the epistemological disjunctivist treatment of underdetermination-based radical skepticism with the response to closure-based radical skepticism offered by hinge epistemology.

In Chapter 6, we return to Humean skepticism. We start by considering G. E. Moore's common-sensical response to idealism, which has traditionally been considered also as a response to skepticism, for reasons we investigate. We explore the complex issue of the relationship between common sense and philosophy, which underlies Moore's response, but also Hume's and Reid's attitude toward the 'hinge' belief that there is an external world, broadly manifest to us through perception.

We then turn to more contemporary discussions. We first look at Jim Pryor's defense of Moore's proof of the existence of an external world. Key to Pryor's *liberal* (or *dogmatist*) position is the idea that perceptual experiences, internalistically construed, can afford one immediate justification for ordinary empirical beliefs (absent defeaters), without the need for any (independently justified) collateral assumptions, such that there is an

external world. Therefore, far from being circular or question-begging, Moore's argument for the existence of an external world is epistemically sound. Pryor concedes, however, that it is dialectically ineffective against skeptical opponents who deny that (we know that) there is an external world, and therefore deny that we do know that we have hands.

The critical discussion of Pryor's position paves the way for a consideration of Crispin Wright's alternative *conservative* (or *skeptical*) view. This maintains that the relevant collateral assumption is needed in order to possess a perceptual justification for the premises of Moore's argument. For Wright—who takes himself to be building on Wittgenstein's insights in *On Certainty*, and who is therefore offering another variety of hinge epistemology—no evidential, truth-conducive justification, empirical or *a priori*, can be obtained for such an assumption. If, as conservatives and Humean skeptics hold, such a justification were needed in order to possess a perceptual justification for one's ordinary empirical beliefs based on perception, and there are only evidential justifications, then the skeptical outcome would ensue that our ordinary beliefs do not amount to knowledge.

Crucially, however, Wright maintains that there is a species of non-evidential justification—*entitlements*, as he calls them—that are applicable for our hinge assumptions. Hence, no skeptical conclusion follows regarding the status of our ordinary empirical beliefs. Wright's position is criticized both with respect to the architecture of perceptual justification it proposes and with respect to the very notion of entitlement. The chapter closes with a discussion of some recent attempts—due to Ralph Wedgwood and Ernest Sosa—to provide *a priori* justifications for these hinge assumptions. Both are found wanting *vis-à-vis* the revision of the notion of the *a priori* they propose and with respect to their effectiveness against radical skepticism.

In the seventh and last chapter, we consider some other formulations of hinge epistemology. We start with P. F. Strawson's naturalist proposal, which we find wanting both in terms of exegesis of Wittgenstein's *On Certainty* and as a systematic response to Humean skepticism. We then expand on Wittgenstein's own normative approach to hinges, by presenting the main tenets of the so-called 'framework' readings of *On Certainty*. The kind of semantic revision imposed by Wittgenstein, whereby hinges are to be understood as rules, and hence are not subject to semantic or epistemic evaluation, together with his claim that skeptical doubts are nonsensical, make for a much stronger but also more contentious response to skepticism than Strawson's proposal.

We then turn to Michael Williams' epistemological contextualism, which, by building on some of Wittgenstein's insights, qualifies as a form of hinge epistemology, at least by our lights. Williams, like Wittgenstein, thinks that our hinges are context-variable. Thus, what is a hinge within one context of inquiry, and is therefore not subject to doubt or verification within that context, may become the object of investigation in a different context of inquiry. Contrary to Wittgenstein, however, he allows for the legitimacy of the philosophical context, whose underlying commitment to foundationalism and to epistemological realism are nonetheless problematical. We criticize Williams' contextualism particularly with respect to its anti-skeptical effectiveness and its propensity to allow for forms of epistemic relativism.

We close the chapter by considering Annalisa Coliva's constitutivist version of hinge epistemology. Coliva seeks a middle course between Pryor's liberal and Wright's conservative view of perceptual justification. Her 'moderate' account holds, *contra* conservatives, that perceptual justification does not require independent justification of the assumption that there is an external world. *Contra* liberals, however, it requires more than simply having a certain course of experience, absent defeaters. For it requires the assumption that there is an external world, broadly manifest to us through perception. Yet, even though such an assumption is neither justified nor justifiable, either evidentially or through entitlements, it is neither epistemically arational, nor arbitrary. According to Coliva, it is rather constitutive of human, and culturally unbound, epistemic rationality. The Humean skeptical paradox is thus blocked by showing that it rests on the mistaken, conservative account of perceptual justification, and too narrow a notion of rationality, one that does not recognize that the latter extends beyond perceptually justified beliefs to those hinges that make the acquisition of perceptual justification, and the deployment of epistemic rationality, possible in the first place. Along the way, a discussion of the relationship between the moderate account of perceptual justification, the principle of transmission of justification, the Closure principle and the challenge from 'abominable conjunctions' is also provided, as well as a discussion of the relationship between the constitutive position and epistemic relativism.

As is apparent from this summary of the content of the book, while each chapter is self-standing and can be used to guide students through the relevant literature, we develop a narrative throughout the whole book. A narrative which, notwithstanding the differences between the two authors,

concurs in according pride of place to Wittgenstein's insights in *On Certainty* and to their developments within the broad family of views that fall within hinge epistemology. We believe this to be the main novelty of the present book.

As is also apparent from the previous summary, the skeptical paradox, particularly in its Humean version, has important connections with relativism. For Humean skepticism may be taken to claim that our ordinary empirical beliefs rest on assumptions, which are not epistemically justified and are therefore arbitrary, from an epistemic point of view, even if they may be psychologically or culturally inescapable. If so, then, the possibility arises that such assumptions may be different—either *de facto*, in different cultures, for instance, or at least in principle. Moreover, if it could be shown that such assumptions may license different verdicts regarding the knowability of ordinary empirical beliefs, we would then witness a powerful form of epistemic relativism whereby different epistemic systems, while being epistemically on a par, would return different and incompatible verdicts regarding the knowability of our ordinary empirical beliefs. One of us has recently explored the issue at length in another book for this very series, with Maria Baghramian, and we recommend reading the first and central chapters of *Relativism* (Baghramian and Coliva 2020), if one is interested in delving deeper into the issue. Relatedly, as part of a much-needed initiative for combating extremism on campus, we have recently developed two Massive Open Online Courses (MOOCs) on skepticism (led by Pritchard) and relativism (led by Coliva) at the University of California, Irvine:

- https://www.coursera.org/learn/skepticism
- https://www.coursera.org/learn/relativism

These MOOCs can be profitably taken alongside this book and the sister book in the series devoted to relativism.

Finally, a note on the text. The book has been jointly planned by the two authors, and throughout the book materials have been re-worked by them together. While the Introduction, Glossary and Chapter 1 are jointly authored, Duncan Pritchard is the main author of Chapters 2, 4, and 5, and Annalisa Coliva is the main author of Chapters 3, 6, and 7.

1

THE SKEPTICAL PARADOX

1.1 Introduction

Many of you will have seen the movie *The Matrix*, directed by the Wachowskis. As Morpheus says to the main character, Neo:

> The Matrix is everywhere. It is all around us. Even now, in this very room. You can see it when you look out your window or when you turn on your television. You can feel it when you go to work, when you go to church, when you pay your taxes. It is the world that has been pulled over your eyes to blind you from the truth.

DOI: 10.4324/9780429057946-2

Morpheus, the leader of a team for the defense of humanity, explains to Neo that the world he has lived in since his birth is the Matrix, an illusory simulated reality built by the machines in 1999 in order to control humankind. Actually, they are now in 2199, and the final stage of the conflict between human beings and machines, which will decide the fate of our species and who will dominate the world, is about to take place.

It is already disturbing to think that everything that one has experienced in one's life up to this moment is the result of a computer simulation. Yet, it is even more disconcerting to realize that, if experience is indistinguishable regardless of its cause—i.e., whether by manipulation by machines or interaction with reality—then there is nothing to prevent even Neo's encounter with Morpheus when he 'reveals' the truth to him about the Matrix from being nothing more than a counterfeit. Much of the intrigue of *The Matrix* depends precisely on the fact that one can never really know whether Neo is really fighting against the machines to save humankind, or whether this is just another manipulation by the machines.

Like many other science fiction movies (just think about Ridley Scott's *Blade Runner* that poses the problem of personal identity), *The Matrix* also raises a philosophical problem, concerning skepticism about our knowledge of the world around us (the so-called 'external' world). This problem was first raised by Descartes in his *Meditations on First Philosophy* ([1641–1642] 1996), and has been a mainstay of philosophy ever since. The issue that it raises is not just a technical one about the nature of knowledge and related epistemic concepts like justification. Rather, this is a puzzle that has existential force, in the sense that it poses a concern that potentially exposes the absurdity of our existence. How can we even make sense of our lives if we cannot be assured that the experiences we are having of the world around us—of our families, our friends, our work, our interests, and so on—are genuine? This is why it is so important to have a proper understanding of both the problem of radical skepticism and the various ways in which one might go about responding to this problem.

1.2 Skepticism

Roughly, to be skeptical means to doubt that one can *know* certain facts or states of affairs. Therefore, skepticism is primarily epistemological rather than ontological. For instance, a skeptic isn't committed to the claim that there are no physical objects (like an idealist), or to the claim that there aren't any

other people (like a solipsist). Rather, for reasons we shall see, she denies that we can *know* that there are physical objects or that we can *know* that there are other minds (i.e., other people). Skepticism thus leads to a form of *motivated* agnosticism regarding our knowledge, where this is agnosticism that is supported by arguments. In particular, this form of agnosticism must be kept distinct from the agnosticism of those who have never considered an issue before, and hence which is not motivated by arguments.

There are many forms of skepticism in philosophy. For example, there is the aforementioned skepticism about the 'external' world: whether we have knowledge of material objects, such as tables, chairs, human bodies, trees, and whatever else that exists independently of our experience of them. There is also a related skepticism about the reliability of our senses: how can we be sure that our senses do not systematically mislead us, since on any given occasion it is through them that we verify whether we have perceived correctly? There is skepticism about other people's minds, which concerns the possibility of knowing whether the people around us really have feelings, emotions, beliefs, and desires, rather than being mere robots who behave as if they had those mental states without actually having them. There is skepticism about induction, which concerns the possibility of knowing that the uniformities in nature that we have observed in the past will continue to take place in the future. One can also be skeptical about the past itself: can we actually know that what we believe to have happened in the remote past really happened a long time ago? How can we know that the world was not created five minutes ago along with everything that makes us believe that it took place a long time before? And there are many other kinds of skepticism too.

In this book we will focus on skepticism about our knowledge of the external world—i.e., about the existence of physical objects such as tables, chairs, human bodies, etc. In fact, this is currently the most discussed type of skepticism within analytical philosophy. Furthermore, it has structural similarities with other forms of skepticism. Thereby, it is reasonable to assume that if it is possible to find a good line of response against external world skepticism, then it may prove useful to deal with other types of skepticism as well.

Pyrrhonian skepticism

Although skepticism of various kinds goes right back to antiquity, skepticism about the external world specifically is a relatively recent and a clearly dated philosophical argument: it begins with Descartes' *Meditations on*

First Philosophy. Arguably, its late origin is not accidental. Ancient skepticism had in fact a practical purpose: through the suspension of judgment one was trying to achieve a form of ataraxy—that is, a kind of serenity that involves the absence of intellectual upheaval. This is the sense in which in antiquity there were people who implemented skepticism as a way of life, as a kind of ethical stance.

This was precisely the kind of ancient skepticism advanced by Pyrrho of Elis, who lived around the third century BCE, and which is consequently known as *Pyrrhonian skepticism*. Pyrrho traveled with Alexander the Great on his expeditions of conquest eastwards, and in the process absorbed the teachings of Indian and Persian scholars, especially with regard to early Buddhist thought (particularly what is now known as the Madhyamaka school of Buddhism). Although he didn't absorb the mystical side of these teachings, he was fascinated by the idea that a radical suspension of judgment might be the path to a kind of philosophical enlightenment. Very roughly, for the Pyrrhonians, it is dogmatism that is the enemy of the good life, and systemic and ongoing doubt is the path away from dogmatism. As a consequence, Pyrrho and his disciples explored the extent to which one could genuinely live a life of thorough-going agnosticism, and so they developed skeptical arguments—the skeptical modes, as they are known (we will review some of them in Section 1.5, while discussing Agrippa's trilemma)—which would lead one to doubt. Of course, since they were concerned to avoid dogmatism of any kind, the Pyrrhonians didn't write any books themselves, as that would be to claim knowledge, and thereby succumb to dogmatism. Fortunately, we know about their ideas from the descriptions provided by others, most notably Sextus Empiricus and his *Outlines of Pyrrhonism*.

The Pyrrhonians recognized that their radical doubt has its limitations, however, given that it is meant to be a way of life (though it's unclear exactly where they thought the line should be drawn in this regard). If one is to live, then one has to act, and that means that some beliefs are required. Could one act if one doubts whether one exists, for example? The application of the skeptical modes was clearly thought of as a social endeavor too, so that rules out doubt about the existence of others as well. Moreover, some of our beliefs don't seem susceptible to the skeptical modes anyway—how would one go about doubting that one is in excruciating pain, for example? In any case, while the Pyrrhonians had all the materials available to them to doubt the existence of an external world,

they never did so, and the likely explanation is simply that this would take their doubt too far for it to serve its ethical purpose.

Crucially, however, while Pyrrhonian skepticism enjoyed a resurgence in the early modern period as the ancient skeptical texts were rediscovered—and so influenced philosophical figures as diverse as Montaigne, Hume, Bacon, and Hegel—in the transition from antiquity to modernity a new form of skepticism emerged which had very different concerns. Rather than being advanced as a way of life, it was instead understood as a kind of methodological challenge. The point is not to put forward skeptical ideas in order to embrace them, but rather to test one's epistemological claims to make sure that they are as sound as one supposes.

This is where the *radical* skepticism that Descartes articulates in the *Meditations*, including his doubt about our knowledge of the external world, comes into play. To stress, the doubt is radical because it challenges the possibility that we may know the very existence of *all* physical objects in our surroundings, including the existence of our own bodies. Furthermore, by raising the possibility that we could be not just victims of lucid and persistent dreams—dreams, that is, that are subjectively indistinguishable from veridical experiences—but also of an evil demon, Descartes calls into question the reliability of our own reasoning capacities. Thus, Cartesian skepticism ends up calling into question both the knowledge we ordinarily take ourselves to have of objects in our surroundings and, even more disturbingly, the reliability of our cognitive faculties. Descartes is not a skeptic, nor does he suppose that skeptical ideas are important to living a good life as the Pyrrhonian skeptics did. His concern is rather to satisfy himself that his account of knowledge, such as his account of how we come to have knowledge of the external world, is skeptic-proof. This has an important consequence, which is that it dramatically extends the limits of how much it is possible to doubt, given that one is not attempting to embrace such doubt within a way of life. (Notice an important ambiguity in the notion of doubt in play here. For the Pyrrhonians this is actual doubt, at least in the sense that the propositions at issue are no longer believed, as belief is suspended. When it comes to Cartesian skepticism, however, it is not being claimed that we are actually able to doubt such claims as that there is an external world. To 'doubt' in this sense rather means something more akin to entertaining the falsity of a claim, to taking a challenge to its truth seriously, neither of which implies actual doubt or suspension of belief.) It is this form of radical skepticism

that we derive from Descartes, and which is shorn of the ethical concerns of its ancient predecessors, that will be our primary concern in this book.

The skeptical paradox

The fact that it is highly unlikely that anyone could actually endorse skepticism about the external world as a coherent philosophical stance does not make it any less disturbing. This is because this kind of radical skepticism purports to be a *paradox*. A paradox is a set of inconsistent claims where the truth of each claim, taken in isolation, seems to be compelling. Philosophy abounds in paradoxes, and what they seem to show is that there is a deep tension in our pre-theoretical concepts within the target domain, such as our views about time, free will, personal identity, and so on (see the Introduction). Cartesian skepticism seems to offer us a paradox regarding our knowledge, specifically our knowledge of the external world (we will examine the claims that make up this paradox in a moment).

Notice how very different skepticism construed as a paradox is from skepticism construed as an argument for a skeptical position. The former is not advancing any specific claim at all, but rather simply highlighting a deep tension in our own epistemological concepts. It is then up to us to work out what we will do in response to discovering that tension. One option might be to try to ignore it and continue as before, particularly if we think that the propositions in doubt—in this case, concerning the external world—are such that we can't help believing them anyway. Even so, ignoring the tension is not to resolve it, and it is hard to see how an anxiety about our epistemic position wouldn't result from responding to the paradox in this way. Far better, if we can, to actually 'scratch the itch' and work out what we need to do to solve the paradox, such as by working out which of the claims that make up the puzzle should be rejected.

There is more than one way to resolve a paradox, and some ways of 'scratching the itch' will be more satisfying than others. A useful distinction in this regard is between *undercutting* and *overriding* responses to a paradox. In essence, undercutting responses to a paradox maintain that there isn't really a genuine paradox to be responded to, but merely the appearance of one, in that, properly understood, our concepts in that domain are not fundamentally in tension as alleged. For example, one might argue that a claim that makes up the paradox has been mischaracterized and that, properly characterized, such that it genuinely captures what we are committed to in this

domain, there is no longer anything paradoxical about the propositions in play (i.e., there is no longer any fundamental tension in our commitments in that domain). As we will see later on, one of the most famous undercutting responses to radical skepticism is offered by Wittgenstein; indeed, it is a general theme of Wittgenstein's work that philosophical paradoxes tend to trade on misunderstandings, and hence are illusory.

In contrast, overriding responses to a paradox concede that the paradox is genuine—i.e., that there really are deep conflicts in our commitments in the relevant domain—but maintain nonetheless that we can resist the paradox by opting for a revisionary account of some aspect of that domain. As we will see in Chapter 2, a certain way of understanding responses to radical skepticism that appeal to epistemic externalism conceives of them as overriding, in that they involve a revisionary account of knowledge, one that is very different from our pre-theoretical conception. All other things being equal, undercutting responses to paradoxes are to be preferred to overriding responses. This is because the former case involves removing the puzzle altogether, and so that is bound to bring more intellectual comfort. This is especially so when it comes to a paradox like that posed by radical skepticism, which, as we noted above, has existential force. If the paradox is merely a theoretical puzzle, then whether the best answer available to it is undercutting or overriding might not matter that much to us, but if it is posing a concern of vital importance for our lives, then we are bound to find undercutting responses of greater intellectual appeal. That said, however, perhaps an overriding response is the best that is available, and of course it is better to have an adequate response to a paradox of any kind than no response at all.

The skeptical paradox about our knowledge of the external world can be formulated in many ways. Here we will mainly deal with two particularly influential versions of it—one which is broadly Cartesian in nature, and one which is an instance of a more general skeptical worry about epistemic circularity that is normally attributed to David Hume. Note, however, that we will not be making historical claims in this regard, concerning what Descartes and Hume were trying to achieve with these arguments, nor will we be advancing the skeptical paradox in the specific way that they formulated it. (For example, Descartes' radical skepticism was structured around a demand for certainty, and yet we will be setting this aspect of Cartesian skepticism to one side for our purposes, as it's inessential to the contemporary formulation of the Cartesian skeptical challenge). Our aim is

rather to highlight the structure of the skeptical paradox as it is these days understood, and in particular how it can be sustained and what strategies of response can be proposed.

1.3 The Cartesian skeptical paradox

The Cartesian skeptical paradox appeals to what are known as *radical skeptical hypotheses*. These are scenarios that are completely indistinguishable from normal life but where we are radically in error. Descartes described scenarios where we are misled by a vivid dream or deceived by an evil demon, but we will consider a more contemporary radical skeptical hypothesis that is more akin to the scenario described in *The Matrix*. This is the so-called brain-in-a-vat (BIV) radical skeptical hypothesis. The idea is that all our experiences are being 'fed' to us by a supercomputer so that while we think that we are interacting with physical objects—such as tables, chairs, other human beings, and so on—in fact, we are simply a brain floating in a vat of nutrients that is hooked up to the supercomputer. This scenario may seem far-fetched, but it is at least a logical and metaphysical possibility in that it doesn't conflict in any fundamental way with how things could be. (For example, it is not a metaphysical possibility that something could be both green and red all over, but there is nothing in the BIV hypothesis that is obviously metaphysically impossible in this sense.)

What is distinctive about radical skeptical hypotheses—and which sets them apart from other far-fetched, though logically possible, scenarios—is that we seem to have no method available to us to exclude them. For example, how would one determine that one is not a BIV? Right now I have the experience of sitting in my chair in my office and working on my laptop. How am I to know that I am having a veridical experience of the world around me, as opposed to just seeming to do so from within the vat? Any test that I might devise to determine that I'm not a BIV, such as feeling my legs (BIVs don't have legs, of course), could simply be the product of the experiences generated by the supercomputer. The worry is that if I can't know that I'm not a BIV, then I can't know anything about the world around me—everything I think I know about the world around me presupposes that I'm not a BIV, after all.

Let's try to unpack this reasoning a bit. Consider a claim about the world around you that you are as sure of as anything. For example, that you have hands. If you are wrong about the fact that you have hands right now, then

all bets are off, epistemically speaking. If you are wrong about this, then what could you possibly be right about? Claims of this kind, however, are inconsistent with being a BIV. A BIV, for example, doesn't have hands, since it's just a disembodied brain floating in a vat of nutrients. Accordingly, if you did know that you had hands, then wouldn't it follow that you must be able to know that you are not a BIV (given that you know that BIVs don't have hands)? Conversely, if one is already convinced that one simply cannot know that one is not a BIV, then doesn't it follow that one can't know that one has hands?

The radical skeptical reasoning in play here seems to go as follows:

Cartesian radical skeptical argument (I)

(1) I can't know that I'm not a BIV.

(2) If I know that I have hands, then I know that I'm not a BIV.

(C) I don't know that I have hands.

The inference is clearly valid. Moreover, notice that the conclusion is generalizable. If one doesn't know that one has hands, then one doesn't know anything much. Indeed, we could repeat this argument with any number of claims in the place of the claim about hands and generate the same result. For example, that I am sitting down is incompatible with being a BIV (as BIVs don't sit anywhere). By the same reasoning, we could thus generate the conclusion that I don't know that I'm sitting down.

We've already considered why the first premise of this inference is true, at least *prima facie*. Notice that the claim in play here is not that we are BIVs, or even that we have any reason to think that we are. The point is just that we don't seem to have any way of knowing that we are not BIVs, and that's all that's needed for (1) to be true. This point is important: the conclusion of the skeptical argument is entirely compatible with our not being BIVs; indeed, with everything we experience being just as we take it to be. (Note that it would suffice for the validity of the argument that one *doesn't* know that one is a BIV, as opposed to it being the case that one *can't* know this. But since the radical skeptic is explicitly motivating the former claim by appeal to the latter claim, it will be useful to have the stronger formulation of (1) on the table.)

The second premise also looks very plausible. One knows that having hands is inconsistent with being a BIV. Hence, insofar as one did know that

one had hands, then shouldn't one be able to infer, and thereby know, that one is not a BIV? The general principle in play here is called the *Closure principle*, and it can be roughly stated as follows:

> *Closure principle*
> If one knows that *p*, and one knows that *p* entails *q*, then one knows that *q*.

This principle certainly seems compelling at first blush. For example, if one knows that one is sitting down, and one knows that if one is sitting down, then one is not standing up, then one knows that one is not standing up. But if the principle is granted, then (2) looks to be true, in that if one knew that one has hands, then one would know that one is not a BIV.

We should note that the formulation of the Closure principle that we just offered is susceptible to a particular kind of objection, albeit one that is easily dealt with. Here is the worry: suppose one knows that p, and knows that p entails q, but one forms one's belief that q in a way that has nothing to do with one's knowledge of p and the entailment to q. For example, perhaps one is a detective and one knows that the murderer is female, and hence that the murderer is not the Duke, but that one's belief that the murderer is not the Duke is in fact not based upon what one knows but purely the result of one's natural deference to aristocrats. In that case, one could know the entailing proposition and the entailment without thereby knowing the entailed proposition.

We can easily avoid this concern about Closure, however, by simply casting the principle diachronically rather than synchronically. This means that we think of the principle as capturing a particular kind of inference that the subject undertakes. Here is the formulation, which is now the standard way of thinking about the Closure principle in the literature:

> *Closure principle (diachronic)*
> If one knows that *p*, and one competently deduces *q* from *p*, thereby coming to believe that *q* on this basis while retaining one's knowledge that *p* throughout, then one knows that *q*.

A few remarks about this formulation are in order. It is meant to be built into the idea of a competent deduction that one knows that the relevant

entailment holds. Since competent deductions take place over time, we also need to stipulate that the subject retains her knowledge of the entailing proposition throughout, as of course if this is lost, then it can undermine the epistemic standing of the belief that results from the inference. But most importantly it is key to this principle that the subject bases her belief in the entailed proposition on this competent deduction from her knowledge of the entailing proposition. It is this aspect of the diachronic formulation of the Closure principle that ensures that it doesn't face the objection just outlined, as in that case the subject's belief in the entailed proposition is based on something other than the competent deduction in play. With closure so understood, it is very hard to see how there could be counterexamples to it (though we will consider the case against it in later chapters). Crucially, the new formulation is just as effective at supporting premise (2) as the old formulation.

Given that the argument is valid, however, if (1) and (2) are both true, then the conclusion must follow. And yet that seems to be absurd. How can it be that one doesn't know that one has hands? Indeed, while we've expressed Cartesian skepticism as an argument to a paradoxical conclusion, we could just as easily cast it directly as a paradox as follows:

Cartesian radical skeptical paradox (I)

(I) One can't know the denials of radical skeptical hypotheses.

(II) The Closure principle.

(III) One knows lots of everyday claims about the external world.

All three of these claims look independently plausible, and yet it seems that they cannot all be true. It follows that at least one of them must be rejected. But which one?

Underdetermination-based Cartesian skepticism

While it is common to understand Cartesian radical skepticism in terms of the Closure principle, there is another epistemic principle available that could also be appealed to in this respect. This is known as the *Underdetermination principle*, and it can be stated as follows:

Underdetermination principle
If one knows that *p* and *q* are incompatible propositions, and one knows that *p*, then one has an epistemic basis for believing that *p* that favors *p* over *q*.

For two propositions to be incompatible means that they can't both be true. Accordingly, if one of them is true, then the other must be false. With this in mind, the Underdetermination principle sounds fairly uncontentious. Consider the simplest case of two incompatible propositions, *p* and not-*p*. This is an incompatibility that surely everyone knows. Given the Underdetermination principle, if one knows that *p*, then one must have an epistemic basis for believing that *p* which favors this over not-*p*. 'Favoring' here just means that one has a better epistemic basis for the truth of the one proposition over the other. So, for example, if one knows that Paris is the capital of France, then one must have an epistemic basis that favors this proposition over the alternative that it's not the case that Paris is the capital of France. This might mean, for instance, that one has more reason to think that Paris is the capital of France than for thinking it's not the case that Paris is the capital of France. In particular, it follows from the Underdetermination principle that if one lacks an epistemic basis that favors the target proposition over what are known to be incompatible alternatives, then one can't have knowledge of it. If one really does have no more reason to think that Paris is the capital of France as to think that some other city is the capital of France—if the matter is genuinely *underdetermined*—then one can't know that Paris is the capital of France. And that seems right. Knowledge demands at least favoring reasons in support of the proposition believed.

With the Underdetermination principle in play, however, we can formulate the Cartesian skeptical paradox. The crux of the matter is that we seem to lack any epistemic basis that favors our everyday beliefs, such as that one has hands, over radical skeptical alternatives, even though we know that they can't both be true. But given the Underdetermination principle, this means that we can't have knowledge of these everyday claims, for example, that one can't know that one has hands.

Here is the puzzle stated as an argument:

Cartesian radical skeptical argument (II)

(1*) I lack an epistemic basis that favors my belief that I have hands over the BIV skeptical alternative.

(2*) If I lack an epistemic basis that favors my belief that I have hands over the BIV skeptical alternative, then I don't know that I have hands.

(C*) I don't know that I have hands.

The argument is clearly valid. We have already seen the rationale for (2*), as this is just an instance of the Underdetermination principle. And the rationale for (1*) should now be familiar, as it's essentially the same as the point we made above when discussing the closure-based formulation of Cartesian skepticism. As we saw, the distinctive feature of radical skeptical hypotheses like the BIV scenario is that they are indistinguishable from everyday life. This is why we have no way of knowing that they don't obtain. It's also why we don't have any epistemic basis for thinking that what we are experiencing is normal veridical perception of the world around us (and thus of hands, etc.), as opposed to being a BIV who merely thinks that she is perceiving the world around her. In a nutshell, we lack favoring epistemic support of the kind at issue in (1*). Given that (1*) and (2*) are true, however, then the skeptical conclusion, (C*), immediately follows.

As with the closure-based formulation of Cartesian skepticism, we can also cast this underdetermination-based rendering of the problem as a paradox. Here it is:

Cartesian radical skeptical paradox (II)

(I*) One can't have an epistemic basis that favors one's everyday beliefs over radical skeptical alternatives.

(II*) The Underdetermination principle.

(III*) One knows lots of everyday claims about the external world.

Each claim that makes up this paradox seems antecedently plausible, and yet we have seen that they appear to be collectively inconsistent. So which of them is to go?

It is an interesting question just how Closure-based and Underdetermination-based formulations of Cartesian skepticism relate to one another. Are they just two different ways of codifying essentially the same underlying skeptical concern, or are they instead offering us distinct skeptical problems? This is one issue that we will return to later in this book, for, as we will see, it is at least arguable that some anti-skeptical proposals are only effective against one of these formulations of the Cartesian skeptical challenge, which suggests that there are important differences between them.

1.4 Humean skepticism and epistemic circularity

In addition to the Cartesian form of radical skepticism just formulated, we can also develop a Humean variation on this skeptical line. In essence, this amounts to a kind of generalization of the skeptical worry at issue in Cartesian skepticism that highlights the *epistemic circularity* that is in play.

Consider the closure-based formulation of the Cartesian skeptical argument that we just examined:

Cartesian radical skeptical argument (I)

(1) I can't know that I'm not a BIV.

(2) If I know that I have hands, then I know that I'm not a BIV.

(C) I don't know that I have hands.

We can highlight the nature of the skeptical problem in play here by imagining someone who attempted to run this argument 'in reverse' and thereby come to know the denial of the BIV radical skeptical hypothesis by appealing to her knowledge of her hands. This is what the argument would look like:

An anti-skeptical argument

(1') I know that I have hands.

(2') If I know that I have hands, then I know that I'm not a BIV.

(C') I know that I'm not a BIV.

As with the skeptical argument that it mimics, this anti-skeptical argument is valid, and hence if the premises are true, then the conclusion will follow. Accordingly, what exactly is wrong with arguing against skeptical hypotheses in just this fashion? Indeed, G. E. Moore (1939) famously argues in just this fashion as part of his 'proof' that there is an external world. (Having hands entails that there is an external world, after all, and so one can just as much draw the conclusion that such a world exists from one's knowledge of one's hands as that one is not a BIV.)

There does seem to be something very fishy with this line of reasoning, and this is where the Humean worry comes in. In particular, it seems like there is a kind of epistemic circularity in play here, and this is what is being highlighted by the radical skeptic. After all, think about what reasons one might offer for thinking that one knows that one has hands, such as that one can see one's hands before one right now (and feel them, etc.). The problem is that these only seem to be good reasons for thinking that one has hands if one has already excluded the possibility that one is a BIV, for if one were a BIV then one would have indistinguishable experiences. This is the sense in which such an anti-skeptical argument seems epistemically circular, in that it appears that one already needs to have a reason for thinking that one is not a BIV before one can know that one has hands. If that's right, then clearly it would be epistemically inappropriate to conclude from one's epistemic basis (such as it is) that one has hands that one is not a BIV.

The Humean point is that such epistemic circularity is widespread in our epistemic practices. Famously Hume exposes that our inductive beliefs are also subject to such epistemic circularity. This is because our inductive inferences assume a commitment to the idea that nature is uniform, but that assumption could only be itself justified by induction; hence the epistemic circularity. Epistemic circularity then compromises our ability to justify induction independently of its very application. Moreover, if it is epistemically legitimate to rely on it only if it can be independently justified, then our beliefs based on induction cannot amount to knowledge after all. The application of the Humean idea doesn't end there though. Think about our perceptual knowledge. For such knowledge to be possible, one would need a justification for thinking that one's perceptual beliefs are formed in generally reliable ways. But how would one go about acquiring such a justification except via the use of one's perception (i.e., by forming perceptual beliefs and checking them against the world)? If so,

then there is no non-circular way of justifying the reliability of our senses. If, in turn, our perceptual beliefs can amount to knowledge only if an independent justification for the reliability of our senses can be provided, then those beliefs cannot amount to knowledge. Arguments with a similar structure can be run against other ways of forming beliefs, such as via memory or testimony. It seems, therefore, that our knowledge in general is under threat, as epistemic circularity is endemic.

1.5 Agrippa's trilemma

It will be useful to close our overview of the radical skeptical problematic (at least as we are conceiving of it) by considering a famous trilemma. This is called *Agrippa's trilemma*, and it's named after an ancient Greek philosopher called Agrippa who belonged to the Pyrrhonian skeptical school that we looked at earlier. In essence, the trilemma represents three of the skeptical techniques, or modes, that the Pyrrhonians employed to provoke doubt.

The trilemma can be summarized as follows. Imagine that you have a belief that is challenged. You might well respond by offering reasons in support of it. But of course those reasons can themselves be challenged, in which case one would be required to offer further reasons in support of those reasons, and so on *ad infinitum*. Alternatively, one might eventually offer reasons in support of one's claims that were previously offered—that is, offer a reason in support of itself (rather than merely repeating some reasons which were already mentioned). Finally, the third option would be to at some point simply decline to offer further reasons at all. Here are the three options that make up the trilemma:

Agrippa's trilemma

(i) Offer an infinite chain of reasons (*infinite regress*).

(ii) Offer a reason in support of itself (*circularity*).

(iii) Offer no reasons at all (*groundlessness*).

The point is that all three options look unpalatable. How could an infinite chain of reasons offer support for one's beliefs, especially given that such a chain could never actually be presented (at least not by finite creatures like us anyway)? Equally, however, how could offering circular reasons be any

use when it comes to justifying our beliefs? We've already noted in our discussion of Humean skepticism how epistemic circularity seems to be inimical to knowledge. Finally, if one is unable to offer any reasons at all, then isn't that admitting that one's beliefs are ultimately groundless? Thus, on all three options, it seems that one's beliefs have a dubious epistemic status, and so fail to amount to knowledge. And yet there doesn't appear to be any other option available, which is why this is posing a trilemma.

Agrippa's trilemma is a useful way of mapping out different ways that one might respond to the skeptical challenge. Historically, the focus has been on options (ii) and (iii) (though there is also a contemporary proposal, known as *infinitism*, that aims to challenge (i)). The view known as *coherentism* argues that there can be circular chains of reasons, at least in certain conditions, and we will be looking later on at the prospects for circular reasons.

A more popular view in the history of philosophy has been some form of *foundationalism*, which directly responds to (iii). Perhaps, for example, there are some beliefs that do not require further reasons because they are in a sense self-justifying. Descartes thought that one couldn't doubt one's own existence, for example, and one could argue that an indubitable belief is self-justifying in this sense. If that were so, then such a belief would be rationally supported without being thereby rationally supported by further reasons. Or perhaps there are some beliefs that simply do not require rational support in order to be properly held from an epistemic point of view. If that's right, then the demand that all our beliefs should enjoy rational support is illegitimate. This seems to be Hume's view, as it happens, and we will encounter versions of it below.

The important point is that, in responding to the Cartesian skeptical challenge, we need to keep in mind how the anti-skeptical proposal in play bears on Agrippa's trilemma, and in particular which of the horns of the trilemma it is rejecting (and why).

Further reading

For a general introduction to the problem of radical skepticism, see Pritchard (2019a). For a recent overview of Pyrrhonian skepticism, see Bett (2010). For the classical defense of the innovative nature of Descartes' radical skepticism when compared with Pyrrhonism, see Burnyeat (1982); for a more recent treatment of this issue, see Williams (2010). For Descartes' debt

to Teresa of Ávila for the evil demon hypothesis, see Mercer (2017). For two recent accounts of the relationship between Pyrrhonian skepticism and Eastern mysticism, see Kuzminski (2010) and Beckwith (2015). For an overview of Cartesian skepticism, see Luper (2010); see also Gaukroger (2010) for a discussion of Descartes' own views in this regard, specifically. For further discussion of the idea of radical skepticism as a paradox, see Stroud (1984, *passim*), Wright (1991), and Beebee (2011), and Pritchard (2014a; 2015, Part one). Pritchard (2015, Part one) discusses the distinction between undercutting and overriding responses to paradoxes. See also Williams (1991, Chapter 1) and Cassam (2007, Chapter 1). The Closure principle is formulated diachronically in terms of competent deduction in Williamson (2001a, Chapter 5) and Hawthorne (2005). For some key comparative discussions of the formulations of Cartesian skepticism in terms of either the Closure principle or the Underdetermination principle, see Brueckner (1994), Cohen (1998), and Pritchard (2005c). See Coliva (2010a) for a historical discussion of Moore's (1939) proof of an external world. See Wright (1985; 2002; 2004) for an influential discussion of Humean skepticism along the general lines that it is understood here. For a recent overview of Hume's own treatment of skepticism, see Beebee (2010). For the suggestion that Hume's skepticism was influenced by Buddhism, see Gopnik (2009). Williams (2015) offers a contemporary discussion of Agrippa's trilemma. For a recent overview of the literature on infinitism, coherentism, and foundationalism, see, respectively, Klein (2010), Olsson (2010), and DePaul (2010). See Vogel (1990a) for an influential response to radical skepticism in terms of inference to the best explanation. For further discussion of the history of skepticism, see Popkin's (2003) seminal treatment. See also Machuca and Reed (2018).

2

CONTENT AND EPISTEMIC EXTERNALISM

2.1 Introduction

In this chapter we will be examining responses to the Cartesian skeptical problem that trade on an appeal to philosophical externalism, whether of a content or an epistemic variety. The guiding idea behind proposals of this kind is that Cartesian skepticism presupposes a problematic form of philosophical internalism, and hence that rejecting this internalism is key to resolving the skeptical puzzle. In particular, both proposals aim to block Cartesian skepticism by maintaining that endorsing the relevant brand of externalism enables us to challenge the first premise of the Cartesian skeptical paradox:

(1) I can't know that I'm not a BIV.

DOI: 10.4324/9780429057946-3

(2) If I know I have hands, then I know that I'm not a BIV.

(C) I don't know that I have hands.

Hence, they both challenge the skeptical claim that we are unable to know the denials of radical skeptical hypotheses.

2.2 Content externalism

In general terms, *content externalism* is the idea that the content of one's mental states (one's thoughts, beliefs, and so forth) can be dependent on one's relationship to an external environment (including one's social environment). As such, it stands in opposition to a *content internalism* that insists that the content of our thoughts can be determined completely independently of our interactions with an external environment. We will consider some specific versions of content externalism in a moment, but one can get a sense of the general idea behind the distinction by considering pairs of identical agents, albeit in different environments. According to content internalism, that these agents are identical ensures that the content of their mental states are identical as well, such that if one agent is thinking a thought about, say, articulated lorries, then the other agent will be thinking a thought with the very same content too, regardless of the fact that they occupy different environments. Content externalism, in contrast, resists this claim, and maintains that one's relationship to the environment can have a bearing on the content of one's mental states.

Descartes is usually understood as being committed to content internalism (though whether this is actually his view is a subject of exegetical controversy), and one can see why. After all, the kinds of radical skeptical hypothesis that Descartes presented seem to presuppose that the content of one's thoughts can be determined completely independently of any engagement one might have with one's external environment. Consider, for instance, his famous example of an "evil demon" who systematically deceives you about (almost) everything that you believe (one key exception, according to Descartes, is that he can't deceive you when it comes to your belief that you exist). The idea seems to be that one's thoughts and beliefs could have the content that one takes them to have even if the vast majority of those beliefs are false. Crucially, however, if one's beliefs are radically false in this way, then it follows that one hasn't been interacting with the

world around you as you supposed; indeed, perhaps you haven't interacted with it at all. One can thus see why one might think that Descartes is committed to a kind of content internalism, in the sense that he presupposes that the content of one's mental states could be determined completely independently of environmental factors. In contrast, if one opts for a form of content externalism, then one might be able to block Cartesian skepticism at its source by rejecting the content internalism that it (allegedly) presupposes.

Content externalism faces some challenges, as does any philosophical view. For example, some commentators have argued that it is in tension with the special epistemic authority we are typically held to have regarding our own mental states (Boghossian 1989; McKinsey 1991; Nuccetelli 2003). But in what follows we will set aside these general difficulties facing content externalism and consider whether, if true, the view has application to the problem of radical skepticism.

Putnam's content externalism and the BIV skeptical hypothesis

In Chapter 1, we encountered the 'brain-in-a-vat' (BIV) radical skeptical hypothesis, which we saw was a variation on the kind of radical skeptical hypotheses advanced by Descartes in his *Meditations* (and also the radical skeptical hypotheses found in some movies, such as *The Matrix*). This skeptical hypothesis was brought back to the fore of contemporary philosophical discussion by Hilary Putnam (1981, Chapter 1) in his defense of a distinctive kind of content externalism, albeit with an important twist, as we will see. In particular, building on his earlier influential work in the philosophy of language (especially Putnam 1975), Putnam argues that there is an important sense in which the thought that one is a BIV cannot be true.

Putnam defends this claim by arguing that the content of one's mental states depends upon one standing in an appropriate relationship to the worldly items at issue, whether directly in terms of causal content with those items, or indirectly in virtue of being part of a linguistic community which fixes the meanings of the terms involved (where some members of that linguistic community, at least historically, have been in causal contact with those items). So take a term like 'tree,' for example. Imagine now that no-one has ever actually been in causal contact with a tree. In that case, argues Putnam, then one simply cannot think thoughts that are genuinely about trees. The actual content of one's thoughts in such cases—supposing

that it succeeds in picking out a determinate content at all—must be concerned with something other than trees. Content externalism of this variety certainly seems plausible. After all, as Putnam points out, to suppose otherwise is to endorse a 'magical' theory of meaning, for if there is no causal interaction of the relevant kind, then how else is the content of one's mental states to be determined?

Crucially, however, such a conclusion seems to have some interesting consequences for the BIV skeptical hypothesis. Putnam's own version of this hypothesis is quite specific, and the details are important, as we will see. According to the Putnam version of the BIV hypothesis, we have always been BIVs, with no human being ever existing outside of the vats. The vats themselves are run by supercomputers, who generate the experiences of those in the vats. Almost everything that is represented in the experiences of those in the vats is nonexistent, so there are no trees, no animals, no towns and cities, and so on. Nonetheless, the BIVs will (it seems) think that these things exist, on account of the experiences that are generated for them by the supercomputers.

The crux of the matter concerns the content of the BIVs' thoughts. According to content internalism, there is no essential barrier to supposing that these thoughts pick out the content that the BIVs suppose them to pick out. According to Putnam-style content externalism, however, this simply cannot be the case. Imagine, for example, that the BIV thinks a thought that seems to be about trees (which uses the word 'tree,' say). But no-one has ever causally interacted with a tree on this scenario, so whatever this thought is about, according to Putnam, it cannot be about trees. At most, it will be about something very different, such as 'tree-images-in-the-vat,' or something like that, but of course thinking thoughts about 'tree-images-in-the-vat' is not the same as thinking thoughts about trees.

Interestingly, we can apply the same kind of reasoning to the BIV's attempt to think thoughts about being a BIV. After all, the BIV has never come into contact with the items in question, such as vats or brains; indeed, *ex hypothesi*, no human being has. It follows that, on Putnam's proposal., a BIV simply can't think contentful thoughts about being a BIV at all. As before, at most she can think thoughts with a very different content, of a kind compatible with being a BIV, but not with the content that they are supposed to have. With this point in mind, we are in a position to understand why Putnam claims that there is a sense in which the thought that one is a BIV cannot be true. Either one is a BIV or one isn't. Obviously,

if someone who isn't a BIV has a thought with the content that she is a BIV, then it is clearly false, so the interesting case is that of someone thinking this thought who is a BIV. Putnam's point, however, is that it simply isn't possible for a BIV to even think a thought with this content, and hence there is no way for their thought to have this content and be true. Either way, then, it seems one cannot truly think the thought that one is a BIV. If that's right, however, then doesn't that mean that we can dismiss the BIV skeptical hypothesis out of hand? More specifically, can we now respond to radical skepticism by insisting that we can know that we are not BIVs?

Unfortunately, matters are not quite as straightforward. One issue here is that it is not clear exactly what Putnam's argument shows. Since the argument seems to demonstrate that one cannot think the thought that one is a BIV and think truly in doing so, one might conclude that it is necessarily false that one is a BIV. That certainly would be a dramatic anti-skeptical conclusion, but a moment's reflection reveals that this can't be what Putnam's argument has demonstrated. After all, Putnam isn't disputing that one could be a BIV, as clearly this is possible. His point is rather about our thoughts in this regard, and in particular what contentful thoughts one can think about when one is a BIV. The crux of the matter is that either we are BIVs, in which case we can't think contentful thoughts about being a BIV, and hence can't think true thoughts in this regard either, or else we are not BIVs, in which case we also can't think contentful true thoughts about being a BIV, for obvious reasons. If we are BIVs, then there is still a sense in which we are radically in error; it's just that we can't give expression in thought to this error since our thoughts are not about what we think they are about (and hence don't have the content that we take them to have).

Moreover, remember that Putnam isn't disputing that it will be indistinguishable to one whether one is a BIV. Accordingly, not only is there still a sense in which one is radically in error if one is a BIV, but one also has no way of knowing whether one is in the 'good' case of not being a BIV (and therefore thinking contentful thoughts in the usual way) or one is in the 'bad' case of being a BIV (and hence not thinking the contentful thoughts that one takes oneself to be thinking). But that doesn't look all that different from the skeptical problem. In particular, it doesn't seem that Putnam's content externalism offers us a way of knowing that we are not BIVs.

Another issue in this regard is that Putnam's response to the BIV-based skeptical puzzle has very narrow application. The reason for this relates to

the very specific way in which he understands this skeptical hypothesis. We noted above that, for Putnam, it is important that everyone has always been a BIV; indeed, there is no thinking creature described as part of this scenario who is, or ever has been, outside of a vat (remember that the deception is being conducted by non-sentient supercomputers). This is crucial to the argument that Putnam offers, since without it, his conclusion doesn't go through. Ordinarily, however, the BIV skeptical hypothesis is not formulated in such a restrictive fashion. For example, one natural formulation of this skeptical hypothesis imagines that one has led a normal life up to the point at which one's brain gets undetectably 'harvested' to go into the vat. Moreover, it is natural on this version to suppose that the deception is being undertaken by superscientists rather than supercomputers. Adding these details to the case severely undermines the supposed anti-skeptical import of Putnam's content externalism, however. After all, one will have had the usual causal interactions with the world around one prior to going into the vat, and presumably the superscientists who are engineering the deception will continue to have such causal interactions. There is thus nothing in principle within Putnam's content externalism that would prevent one from thinking thoughts with the usual content while in the vat.

Moreover, even if one hasn't had the requisite causal interactions prior to entering the vat, this still need not be a barrier to thinking contentful thoughts while in the vat. Although content externalism focusses on the importance of having the relevant causal connections to the referents of one's thoughts, it allows that such a causal connection could be attained indirectly, in virtue of being part of a linguistic community (some of whose members have had the requisite causal relationships directly). This claim is needed to make content externalism plausible, since otherwise it would place an implausible causal constraint on the content of our mental states. If trees disappeared from the planet tomorrow, such that our causal connections with them was severed, we could surely still teach our children about trees, and thereby enable them to think contentful thoughts about trees, even though they couldn't causally interact with trees directly. If that's right, however, then even if one has always been a BIV, it needn't follow that one is unable to think thoughts with the same content as those who are not BIVs, as in principle it is possible that whoever is conducting the deception (the superscientists, say) could be appropriately inducting the BIV into their linguistic community, just as we are able to teach

children in a newly treeless world about trees so that they are able to think thoughts about trees.

The problem, however, is that the radical skeptic only needs one plausible radical skeptical hypothesis in order to motivate their skepticism. So if there are versions of the BIV skeptical hypothesis that are consistent with content externalism, then they can simply run their argument using those formulations. The effect of such an argument, if successful, would be no less devastating on the proviso that there are other formulations of this skeptical hypothesis that would be ruled out by content externalism. It follows that the prospects of employing Putnam-style content externalism to respond to the radical skeptical problem are dim indeed.

Davidson's holistic content externalism

A different kind of content externalism, one that also seems to have anti-skeptical implications, is offered by Donald Davidson. In particular, Davidson (1983, 314) famously motivates a form of content externalism that implies that "belief is in its nature veridical." If this were true, then it would have obvious ramifications for radical skepticism, not least in terms of radical skeptical hypotheses, like the BIV hypothesis, which involve widespread falsity in one's beliefs. In particular, if we could be assured that our beliefs cannot be massively false, then that ought to provide us with a basis to know that such radical skeptical hypotheses are false.

The starting-point for Davidson's content externalism is the notion of *radical interpretation*, which is when one interprets a speaker without reliance on any prior knowledge of either the speaker's beliefs or the meanings of the speaker's utterances. (Think, for example, of what is involved in an anthropologist trying to interpret the utterances of people who are being encountered for the first time.) The importance of radical interpretation for Davidson's work is that he holds that it is a necessary truth that any content-bearing sentence is interpretable under these epistemological conditions. Radical interpretation faces a problem, however, which is that one cannot assign a meaning to a speaker's utterances without knowing what the speaker believes, and one cannot identify the speaker's beliefs without knowing what her utterances mean. So, for example, if one does not already have a purchase on a speaker's beliefs, then how is one to interpret an utterance of a sentence (in the vicinity of a rabbit, say) as being a sentence with a particular content (such as, 'that's a rabbit')? But if

one cannot assign meanings to the speaker's utterances, then how is one to determine, in the conditions of radical interpretation, what the speaker's beliefs are? One is thus apparently stuck within an interpretative circle.

Davidson's resolution of this problem is to appeal to the *principle of charity*. While this is never given a particularly precise rendering, it roughly instructs the interpreter to treat the speaker as having mostly true beliefs (by the interpreter's lights anyway). It's clear from how Davidson applies this principle that there are two key dimensions to applications of charity. On the one hand, interpreting a speaker as charity demands will tend to lead to an interpretation which treats the speaker as having a coherent—or, at least, *consistent* anyway—set of beliefs. On the other hand, charity will tend to lead to an interpretation which treats the speaker's beliefs as being correctly about the objects in the speaker's immediate environment which the speaker is interacting with.

By using the principle of charity in this fashion, we have a way out of the interpretative circle just noted. The radical interpreter is using her beliefs as a way of assigning beliefs to the speaker, and in doing so is able to assign meanings to her utterances. So, for example, if the speaker is observed uttering sentences in the vicinity of a certain object, such as a rabbit, then by treating the speaker as having mostly true beliefs (by the interpreter's lights, anyway), one will have a way of ascribing a meaning to the speaker's utterances, such as 'that's a rabbit.' Of course, the ascription will be defeasible, and in time the radical interpreter may settle on a very different interpretation of the speaker's utterances. But at least this application of the principle of charity enables the radical interpreter to coherently begin the project of radical interpretation.

On the face of it, the principle of charity offers us no particular respite from the radical skeptical challenge. That it may be methodologically necessary for the project of radical interpretation to treat a speaker's beliefs as mostly true by our lights does not in itself give us any more reason to think that the speaker's beliefs are in fact true than it does for thinking that our beliefs are true. So how does Davidson get from the principle of charity to the idea that belief is in its nature veridical?

One key dialectical move that Davidson (1975; 1977; 1983) makes in this regard is to appeal to an *omniscient interpreter*. This is "an interpreter who is omniscient about the world, and about what does and would cause a speaker to assent to any sentence in his (potentially unlimited) repertoire" (Davidson 1983, 317). According to Davidson, the omniscient interpreter

would be obliged to employ the principle of charity too, and hence in interpreting our utterances would treat most of our beliefs as true. But, being omniscient, our interpreter's beliefs are by definition true. Hence, it follows that our beliefs, which are in accordance with the omniscient interpreter's beliefs, must be true as well, and hence we get the conclusion that belief is in its nature veridical.

This argument is highly contentious. One worry is that the very idea of an omniscient interpreter being bound to use the principle of charity in making sense of our utterances is itself problematic. For although we can understand why non-omniscient creatures such as ourselves might need to employ the principle of charity in interpreting the utterances of others, why would an omniscient creature be so restricted? In particular, why would a creature who is omniscient about 'what does and would cause a speaker to assent to any sentence in his (potentially unlimited) repertoire' need to rely on a methodological principle like the principle of charity in making sense of a speaker's utterances? After all, a core part of the reason why we need to appeal to this principle is our supposed lack of epistemic access, in the context of radical interpretation anyway, to what is causing the speaker's utterances.

Even setting aside this concern, there is a deeper worry in play here, which is that even if we grant that the omniscient interpreter will ascribe mostly true beliefs to us, this still seems consistent with there being a significant mismatch in how we conceive of our own situation and how the omniscient interpreter, from its epistemically elevated viewpoint, conceives of it. In short, that our beliefs are mostly true does not in itself guarantee that we are not in some fundamental respect in error. For example, suppose that we are BIVs, and being interpreted by an omniscient interpreter outside of the vat. The omniscient interpreter could treat us as having mostly true beliefs by taking those beliefs to be about events that are internal to our (artificially-generated) experience of the world, rather than being caused by events in the world outside the vat (which is what we take our beliefs to be about). The point is that there is a way of satisfying the principle of charity in this interpretative enterprise involving the omniscient interpreter that is entirely consistent with us having beliefs that fundamentally fail to hook up with the world in the manner that we suppose. But if that's right, then what possible anti-skeptical comfort could be provided by the appeal to this thought-experiment? (This should remind us of how Putnam's brand of content externalism seemed to also be

consistent with the idea that the BIV is subject to massive error, even if that is not something that the BIV herself can formulate in her thoughts.)

Interestingly, while critics of Davidson treat the omniscient interpreter argument as central to his view, he himself has distanced himself from it. But if that's right, then how does Davidson draw anti-skeptical conclusions from his content externalism? In order to understand this, we need to look again at how he motivates his position, and in particular the notion of *triangulation* that he appeals to.

According to Davidson, triangulation involves a causal nexus involving two subjects and an object in a common physical environment. Triangulation occurs when both creatures react to that object and then react in turn to each other's reactions. Crucially, according to Davidson (e.g., 1991), triangulation is essential to there being mental content in the first place, in the sense that it is a metaphysically necessary condition for the acquisition of (contentful) thought. This commitment is core to his variety of content externalism, since it makes causal relationships to matters external to the subject necessary for contentful thought.

In order to see how triangulation might be related to Davidson's views on radical interpretation and the principle of charity, consider this passage, which comes just after a discussion of how the principle of charity blocks radical skepticism:

> What stands in the way of global skepticism of the senses is ... the fact that we must, in the plainest and methodologically most basic cases, take the objects of a belief to be the causes of that belief. And what we, as interpreters, must take them to be is what they in fact are. Communication begins where causes converge: your belief means what mine does if belief in its truth is systematically caused by the same events and objects.
>
> (Davidson 1983, 317–318)

Here we can see how the principle of charity, *qua* an indispensable ingredient in radical interpretation, is guiding an implicit commitment to triangulation. The connecting thought is that those cases where triangulation occurs are precisely the kind of "plainest and methodologically most basic cases" in which we are required, following the principle of charity, to interpret a speaker so that she ends up speaking truly. Moreover, notice the remark that "what we, as interpreters, must take them [i.e., the objects

of belief] to be is what they in fact are." Davidson's point is that these most basic cases effectively determine the contents of the relevant beliefs, so that there is no logical gap between what we as interpreters take the objects of a belief to be and the causes of that belief which could allow for the possibility of massive falsehood in one's beliefs. This is why the possibility of massive error in one's beliefs is impossible, *contra* the skeptic, and thus why "belief is in its nature veridical." For to suppose that this is possible is to suppose that the belief ascriptions offered in the "plainest and methodologically most basic" cases of triangulation could be systematically false, and that claim is incompatible with Davidson's content externalism.

When Davidson's argument is put this way, it should be clear that it represents a broadly transcendental argument against skepticism. The idea is to demonstrate on purely *a priori* grounds that there is a necessary condition for one even thinking a contentful thought—*viz.*, that most of one's beliefs must be true. Since even the skeptic is committed to the possibility of there being contentful thought (as otherwise we could not even make sense of the skeptical enterprise, still less the specific skeptical appeal to radical error-possibilities, such as the scenario that one might be a BIV), so even the skeptic must accept the anti-skeptical consequences of this transcendental argument if it is sound. Indeed, if this transcendental argument is sound, then one can never even coherently expound radical skepticism.

The transcendental nature of Davidson's anti-skepticism does, however, leave it open to the familiar criticisms that are leveled against transcendental arguments. In particular, there is the influential criticism of this style of argument that is due to Barry Stroud (1968). Very roughly, Stroud's charge is that transcendental arguments demonstrate only what we are committed to thinking is the case, but fall short of actually showing what is the case. In terms of Davidson's transcendental treatment of radical skepticism, the point would be that all that follows is that we must think of our beliefs as being in general veridical, but that it need not follow from this that our beliefs are in general veridical. In particular, for all Davidson has argued, our beliefs could still be massively false.

Still, one might think that it is philosophical progress of a certain kind to discover that we are all, including any actual radical skeptics who might be out there, committed to treating our beliefs as mostly true. Recall, however, our point above about how even if it is true that our beliefs are in general veridical, this could nonetheless be compatible with our relationship to the world being radically different from how we suppose it to be (including

such that we are in fact BIVs). We made that claim above with regard to how we might be interpreted by an omniscient interpreter. But now the omniscient interpreter has left the scene, the point is even more apt. If all Davidson has shown is only that we are obliged to treat our beliefs as mostly true, but not that they are mostly true, then doesn't that leave the intellectual anxiety that is prompted by radical skepticism entirely intact? In particular, are we any closer to being able to recover our knowledge of such propositions from the radical skeptical threat by becoming aware that we must treat our beliefs as mostly true regardless of how we are in fact related to the world around us?

Chalmers: The Matrix as metaphysics

In a fairly recent paper, David Chalmers (2005) offers an intriguing twist on standard content externalist responses to radical skepticism. Focusing on the BIV scenario, Chalmers argues that, properly understood, this is not a radical skeptical hypothesis at all (i.e., one that undermines our knowledge of the world around us), but rather a *metaphysical hypothesis*, in that it is offering an account of the underlying nature of reality. If that's right, then it has some important consequences, since it means even if the BIV hypothesis were true, this needn't imply that most of our everyday beliefs about the world are false. Chalmers' anti-skepticism is thus different to Putnam's and Davidson's, since he isn't challenging the truth of the BIV hypothesis itself, but rather its putative radical skeptical implications.

Chalmers notes that there are all sorts of competing metaphysical hypotheses about the underlying nature of reality. He considers three examples. The first is the *computational hypothesis* that underlying even the most fundamental physical processes is a further level of computational processes that underpin everything. The second is the *creation hypothesis* that physical space-time was created by beings outside of physical space-time. The third is the *mind-body hypothesis* which states that our minds are constituted by processes outside physical space-time. Chalmers notes that all three metaphysical hypotheses have had their proponents. There are some scientists today who take the computational hypothesis seriously, and of course there are those with theistic religious beliefs who would be willing to endorse the creation hypothesis. Some of those religious believers might also endorse the mind-body hypothesis too. Historically this also has had some prominent adherents, such as Descartes.

What is important is that in each case the truth of the metaphysical hypothesis doesn't seem to entail that most of our everyday beliefs are false. Of course, it entails that *some* of our beliefs are false. For example, if one believes that there is nothing outside of physical space-time, and certainly no 'creator,' then, of course, the truth of the creation hypothesis entails that this belief is false. But why should it follow from this that one's everyday beliefs about tables and chairs and such like are false? The same goes for the other two metaphysical hypotheses, in that their truth only seems to imply that a certain sub-set of our beliefs (mostly theoretical in nature) are false, and not our beliefs about the external world in general. If they are true, then we are ignorant about the fundamental nature of things, but that doesn't mean that one's everyday beliefs are false, or that they don't amount to knowledge.

If that's right, and the BIV scenario is also a metaphysical hypothesis, then it holds out the possibility that its truth could also be consistent with the widespread truth of our everyday beliefs. Interestingly, Chalmers is quite specific about the metaphysical hypothesis that the BIV scenario concerns (at least the version of the BIV scenario where we have all always been envatted), as he claims that it is the combination of all three of the metaphysical hypotheses just noted. For the sake of argument, let's grant Chalmers this claim. Does it follow that thinking of the BIV scenario in these metaphysical terms means that its truth can be compatible with the truth of our everyday beliefs?

The idea that our everyday beliefs are true regardless of the truth of the BIV scenario should remind us of Putnam's content externalism. Remember that on this proposal the content of our beliefs when we have always been envatted are simply about different things. When a non-envatted person forms a belief about their hands, for example, then the content of this belief concerns the worldly items that the belief is about. In contrast, when an envatted person forms a corresponding belief, then its content concerns something very different—images of hands in the vat, or something like that. That's why both beliefs can be true, as they are picking out very different contents, with the way that content is fixed determined, at least in part, by factors that are external to the subject thinking the thought, such as their actual relationship to the world or their participation in a certain social and linguistic practice.

Chalmers seems to have a similar idea in mind. If we are envatted, then 'real' hands are made out of the bits that comprise the underlying

computational processes rather than physical objects, but they are no less real as a result. So the beliefs one has about hands in the vat, while having a different content to our beliefs about hands outside of the vat, are no less true as a consequence. In particular, why should it matter whether the fundamental nature of the hands that we are having thoughts about is physical or computational? This is the main difference between Chalmers' approach and Putnam's. Whereas Putnam was concerned to show that there is a sense in which the BIV hypothesis must be false, Chalmers wants to argue that its truth is essentially harmless.

Even if one grants Chalmers' point, however, it nonetheless has limited appeal when it comes to the problem of radical skepticism, and for similar reasons to why Putnam's line also had limited appeal. Although he claims otherwise, Chalmers' argument only works against BIV scenarios where we have always been envatted, since it is only on this supposition that it works as a metaphysical hypothesis. In order to see this, imagine a version of the BIV scenario which is more akin to how it is usually understood, whereby an individual who had previously lived a normal life is captured and placed in the vat where she is 'fed' deceptive experiences. On this interpretation of the BIV scenario, there is no claim that the fundamental nature of reality is any different from how we, outside of the vat, imagine it to be. In particular, it is precisely built into the scenario that the 'reality' presented to the envatted subject is deceptive. Moreover, as we noted above with regard to Putnam's content externalism, a recently envatted person ought to be able to think thoughts with the same content as she previously did outside of the vat, given that she was inculcated into the relevant socio-linguistic practices (with their causal relationships to the items referred to as part of those practices). It follows that the thoughts that she thinks in the vat will have the same contents as corresponding thoughts that one might think outside of the vat, and hence can be false. For example, if she thinks that her hands are in front of her, because that is what she is 'seeing' in the vat, then her thoughts will really be about hands (i.e., and not vat-hands), and hence will be false.

Crucially, however, the radical skeptic only needs one good radical skeptical hypothesis in order to run their skeptical challenge. It follows that even if some putative radical skeptical hypotheses can be dismissed in this regard as metaphysical hypotheses, this doesn't give us any purchase on the problem of radical skepticism unless we can show that this strategy applies to all radical skeptical hypotheses. Since it doesn't, it leaves the problem of radical skepticism essentially untouched.

2.3 Epistemic externalism

A different way of responding to the problem of radical skepticism by appealing to a form of philosophical externalism is to focus on the prospects of a specifically *epistemic* externalism in this regard. The guiding idea is very similar to that we saw being put forward by semantic externalists above, however, which is that radical skepticism is problematic because it illicitly trades on a kind of philosophical internalism, in this case, of an epistemic variety, rather than being concerned with internalism about content.

How one should draw the epistemic externalism/internalism distinction is itself a controversial matter, but we can side-step some of this controversy by focusing on the best way of marking this distinction if one is to argue for an epistemic externalist response to radical skepticism. This would require us to focus our formulation of the distinction at the level of knowledge (i.e., as opposed to another epistemic standing, like justification), given that we have formulated the skeptical problem in terms of the possibility of knowledge regarding an external world. With this in mind, let's define some terminology. An *epistemic condition* is a necessary condition for knowledge, over and above the demand that knowledge entails true belief. An *internalist epistemic condition* is then an epistemic condition that one can come to know that one has satisfied by reflection alone. (This is an *accessibilist* way of understanding epistemic internalism, as opposed to the competing *mentalist* account. In essence, mentalism argues that internalist epistemic conditions supervene on one's mental states. But since defenders of mentalism usually grant that one's mental states are reflectively accessible, we can bracket this distinction for our purposes.) In contrast, an *externalist epistemic condition* is an epistemic condition that one cannot come to know that one has satisfied by reflection alone. *Epistemic internalism about knowledge* is then best understood as the idea that knowledge entails an internalist epistemic condition. *Epistemic externalism about knowledge*, in contrast, is the denial of this claim, and hence maintains that there are no necessary internalist epistemic conditions on knowledge.

Let's try to unpack this distinction between externalism and internalism about knowledge, beginning with what is involved in an internalist and an externalist epistemic condition. The most natural way to understand an internalist epistemic condition is in terms of the justification condition, classically conceived. This is usually cast as demanding that the subject has good reasons for believing the target proposition to be true, where these

are reasons that are reflectively available to the subject concerned, in the sense that the subject is, in principle at least, able to cite them. In particular, such reasons are available to the subject without her having to conduct any further empirical investigation.

It is justification in this sense that was part of the classical tripartite account of knowledge as justified true belief, which was famously attacked by the counterexamples offered by Edmund Gettier (1963). These counterexamples turn on the fact that the traditional notion of justification is non-factive, which means that one can have a justified belief in a false proposition. What Gettier's cases do is take a justified belief that would ordinarily be false and add a twist to the scenario such that the belief is true regardless, albeit in a way that has nothing to do with the subject's justification. The result is that the subject's belief, even though true and justified, is nonetheless only luckily true, and that is held to be incompatible with having knowledge.

For example, consider this famous Gettier-style case, due to Roderick Chisholm (1977, 105). Imagine that a farmer is looking into a field and sees, in good perceptual conditions, what appears to be a sheep there. Accordingly, he forms the belief that there is a sheep in the field. Given how he has formed his belief, it is surely justified—farmers are generally quite good at spotting sheep in good perceptual conditions after all. The twist in the tale is that what the farmer is looking at is not a sheep at all, but rather something that looks like a sheep (such as a big hairy dog). The belief is nonetheless true, however, as there is a genuine sheep hidden from view behind the sheep-shaped object that the farmer is looking at. The farmer's belief is thus justified and true, but is only true as a matter of luck (i.e., the luck that there happened to be a real sheep behind the sheep-shaped object that he is looking at), and hence is not an instance of knowledge.

Gettier-style cases demonstrate the untenability of a certain kind of strong epistemic internalism about knowledge, whereby there is only one epistemic condition on knowledge and it is an internalist epistemic condition (at least where that internalist epistemic condition is understood along classical lines—we will return to this issue when we examine epistemological disjunctivism below). But such cases don't thereby show that epistemic internalism about knowledge in general is untenable, since it is possible that one might supplement the classical tripartite account of knowledge with further epistemic conditions, including, potentially, further externalist epistemic conditions.

The paradigm case of an externalist epistemic condition in the literature is that provided by *reliabilism*, whereby a belief enjoys a positive epistemic status if it is formed via a reliable belief-forming process (i.e., roughly, one that tends to produce true beliefs). Reliabilism as part of a theory of knowledge would thus hold that this is an epistemic condition on knowledge; in its strongest form, it is the view that it is the *only* epistemic condition on knowledge. What makes this an externalist epistemic condition is that whether or not one's belief-forming process is reliable depends on the nature of the world. The reliability of one's perceptual beliefs, for example, will depend on how well one's perceptual faculties are functioning, and the kinds of environments in which they are employed. Accordingly, it won't always be possible to determine whether one has met this condition by reflection alone. Instead, one will typically need to undertake an empirical investigation regarding the reliability of the belief-forming process that one is using.

The idea that knowledge demands the satisfaction of an internalist epistemic condition tends to go hand-in-hand with the thought that good reflectively accessible reasons are important to knowledge, such that even if there is more to knowing than having a true belief supported by such reasons (e.g., because of Gettier-style cases), one must at least have these supporting reasons. One can see the attraction of the idea, since if one's knowledge is not supported by such reasons, then it seems to be in a certain sense opaque to one that one knows. Consider, for example, a subject whose true belief was formed via a reliable belief-forming process, but who has no good reflectively accessible reasons in support of that belief. Perhaps, for example, she just finds herself believing what she does and has never even considered what would be a good rational basis for this belief. On an externalist account of knowledge (at least of reliabilist stripe), such a belief ought to be a good candidate for knowledge. And yet it does seem odd that our agent could have knowledge even while being unable to articulate any good reasons in support of what she believes.

In any case, we don't need to adjudicate between externalism and internalism about knowledge; our concern is rather to work out if this distinction has any bearing on the problem of radical skepticism. Here is one way that it might be relevant. Doesn't the radical skeptical problem essentially trade on the idea that we lack good reflectively accessible reasons in support of our beliefs? For example, radical skepticism appeals to the thesis that since we have no reason to exclude the possibility that we

are BIVs, so we don't have any good reasons in support of our everyday beliefs either. In particular, one's reasons for believing that one has hands—that one seems to be able to feel them and see them before one—don't seem to be good reasons for believing this at all insofar as the BIV radical skeptical hypothesis is in play (since the BIV will also seem to be able to feel and see their hands before them). If the radical skeptical problem does essentially trade on this appeal to reasons, however, then it appears to presuppose a commitment to epistemic internalism about knowledge. In particular, it seems to presuppose that having knowledge means having a justification to believe what one does in the classical internalist sense of that term.

If that's right, then this might provide the conceptual space for epistemic externalists to offer a response to radical skepticism. In a nutshell, the thought will be that all that the radical skeptic has demonstrated is that we lack knowledge by the lights of an epistemic internalist treatment of this notion, and not that we lack knowledge *simpliciter*. In particular, it would be consistent with the radical skeptical argument that our beliefs meet whatever externalist epistemic conditions are necessary for knowledge. Indeed, for all the radical skeptic has shown our beliefs *do* meet externalist epistemic criteria. For all the radical skeptic has argued, for example, our beliefs could well be formed via reliable belief-forming processes, at least insofar as the world really is how we think it is. (Remember that the radical skeptic hasn't offered us any reason for thinking that the world *isn't* the way that we think it is. For instance, they haven't offered us any reason for thinking that we *are* BIVs, but only claimed that we lack good reasons for thinking that we are not BIVs.) Moreover, the epistemic externalist can diagnose the enduring appeal of the radical skeptical problem as arising out of a mistaken commitment to epistemic internalism.

Part of the challenge facing epistemic externalist responses to radical skepticism is to explain why this isn't a capitulation in disguise. Wasn't it knowledge of a specifically internalist kind that we were trying to rescue from the radical skeptic? To be told that our knowledge is secure from skeptical challenge, but only on the understanding that we accept that we have no good reasons in support of our beliefs doesn't sound very *anti*-skeptical at all. Although the epistemic externalist is right that the radical skeptic hasn't given us any basis for thinking that our beliefs fail to meet the target externalist epistemic condition—they haven't given us any grounds for thinking that our beliefs are formed in unreliable ways, for example—the

point remains that if the skeptic is right, then we have no reflectively accessible rational basis for thinking that we have met such an externalist epistemic condition. We are thus left in a kind of intellectual limbo. The skeptic hasn't shown that we necessarily lack knowledge, but the point remains that whether we in fact have knowledge is not something that we have any reflective access to, since it depends on the satisfaction of external epistemic conditions that are beyond our reflective ken.

The concern in play here is sometimes expressed in terms of the *iterativity principle* for knowledge. This holds that when one knows, one knows that one knows. So the thought runs, epistemic internalism about knowledge entails iterativity, but epistemic externalism about knowledge, in contrast, is inconsistent with it. And yet isn't it important not just that we have knowledge, *contra* the skeptic, but also that we know that we do?

While this way of putting matters might be superficially appealing, it is in fact deeply mistaken. This is because it is simply false both that epistemic internalism about knowledge entails iterativity and that epistemic externalism about knowledge is inconsistent with iterativity. Let's take these claims in turn. Epistemic internalism about knowledge entails that when one knows one's belief that one knows meets an internalist epistemic condition, such as the justification condition, classically conceived. But that's not enough to ensure that one has second-order knowledge. Perhaps, for example, one's second-order belief is Gettierized? The confusion here is that while epistemic internalism about knowledge might entail, for example, that when one knows one has a good reason to think that one knows, this is not the same thing as knowing that one knows—one could satisfy the former without thereby satisfying the latter.

If one advanced a pure form of epistemic internalism about knowledge, such that meeting an internalist epistemic condition sufficed for knowledge, then one would be better placed to argue that one should in addition accept iterativity. In being aware that one meets such a condition, and that it suffices for knowledge, wouldn't one at least be in a position to acquire second-order knowledge? We've already noted, however, the implausibility of such an account of knowledge. For example, the idea that knowledge is nothing more than true belief that in addition meets the justification condition, classically conceived, has been undermined by Gettier-style cases. In any case, even if it were true that there are versions of epistemic internalism about knowledge that lead to iterativity, the point

would remain that there isn't a general entailment from epistemic internalism about knowledge to the iterativity principle.

Now consider epistemic externalism about knowledge. Although highly unlikely, there is no principled reason why whatever externalist epistemic conditions that need to be satisfied in order to have first-order knowledge are not also thereby satisfied with regard to second-order knowledge as well. Perhaps, for example, whenever one's beliefs at first-order are reliably formed, one's corresponding beliefs at second-order are also reliably formed? The idea that there is an inconsistency between externalism about knowledge and iterativity seems to turn on the thought that second-order knowledge is by its nature knowledge that satisfies an internalist epistemic condition. Epistemic externalism would certainly be inconsistent with iterativity so understood—why would first-order knowledge that satisfies an externalist epistemic rubric entail second-order knowledge that satisfies an internalist epistemic rubric? Crucially, however, there is no reason why second-order knowledge should demand different epistemic conditions to first-order knowledge.

The status of the iterativity principle is thus completely orthogonal to the externalism/internalism debate about knowledge, and thus the application of this debate to the problem of radical skepticism. What is important here, and which is being mistakenly understood in terms of iterativity, is whether, for example, having knowledge demands having good reasons for thinking that one knows. This is something that epistemic externalism about knowledge is rejecting, in contrast to epistemic internalism about knowledge.

Why might having such reasons be so important? One rationale in this regard concerns the propriety of claims to know. If one claims that one has knowledge, then it seems that one ought to be able to defend one's claim, and that means being able to offer good reasons in its support if called upon to do so. This is the extent to which claiming knowledge demands taking cognitive responsibility for the truth of the target proposition, where discharging that responsibility is naturally understood in epistemic internalist terms.

If that's right, then it seems that opting for epistemic externalism means contending that one can satisfy the conditions for knowledge without thereby satisfying the conditions for being able to appropriately claim that one has knowledge. That's not problematic in itself, as the proponent of epistemic externalism can simply insist that knowing and being in a

position to appropriately claim knowledge are different things. But this issue does loom large with regard to the problem of radical skepticism. It seems that while epistemic externalism maintains that the radical skeptic hasn't shown that we lack knowledge, it must nonetheless concede that one is not in a position to appropriately claim that one has knowledge (since it grants that our anti-skeptical knowledge doesn't satisfy the internalist rubric). And that looks like a rather big concession to make in the face of the problem of radical skepticism.

Nonetheless, the epistemic externalist can push back in this regard, since they can counter that it is simply part of the human epistemic condition that we lack a reflective rational basis that is resistant to radical skeptical challenge. Doesn't this further reinforce the thought that we were mistaken in ever thinking that such a basis must be essential to knowledge? Moreover, epistemic externalists will point out that they have their own ways of capturing notions of cognitive responsibility that are distinct from the intellectualist model provided by epistemic internalism, with its focus on reflectively accessible reasons.

Virtue reliabilism, for example, treats knowledge as arising not from reliable belief-forming processes in general, but rather from those specific skill-like reliable belief-forming processes that are stable and integrated features of a subject's cognitive character, and thus represent the subject's cognitive *virtues*. The idea is that whereas mere reliability in one's belief-forming processes doesn't underwrite any significant notion of cognitive responsibility—the reliability might have nothing to do with one's exercise of cognitive agency, for example—cognitive success that is the result of cognitive virtue does underwrite a substantive kind of cognitive responsibility. In particular, when one is cognitive successful as a result of manifesting cognitive virtue, even when one is lacking in reflectively accessible supporting reasons, that cognitive success is attributable to one's exercise of cognitive agency. Accordingly, while the knowledge that we possess might not be supported by reflectively accessible reasons, it can nonetheless manifest one's cognitive agency. In short, skillfully getting to the truth in the absence of reflectively accessible reasons is not to be equated with merely stumbling across the truth in the dark (even if the absence of such reasons means that one cannot properly claim the knowledge that one possesses).

One question we might ask of an epistemic externalist anti-skeptical proposal is what role it allows for reflectively accessible reasons, given that they are inessential for knowledge. Modest forms of epistemic externalism

usually accommodate some role for reflectively accessible reasons. For example, one might naturally opt for a kind of 'two-tiered' view whereby while knowledge in general doesn't demand the satisfaction of an internalist epistemic condition, nonetheless there is a kind of knowledge that satisfies such a condition, albeit where the rational support in question is not to be understood as being resistant to the radical skeptical challenge. Knowledge in general is thus an externalist notion, but that is compatible with some knowledge meeting the internalist rubric (call this *internalist knowledge*). Moreover, it is open to the epistemic externalist to allow that there is something particularly desirable about internalist knowledge; all that is important for their purposes is that knowledge is fundamentally concerned with the satisfaction of externalist epistemic desiderata—*viz.*, that not all knowledge is internalist knowledge (because some of it is merely externalist knowledge).

The problem, however, is that such moderate epistemic externalism, while independently attractive, doesn't really survive radical skeptical scrutiny. The difficulty lies in the deceptively simple idea that there can be reflectively accessible rational support that is both genuine and yet which doesn't survive radical skeptical challenge. At first blush, the thesis in play can seem quite straightforward: there are reflectively accessible reasons available to us, but only of a limited kind, and not of a sort that would be resistant to radical skepticism. But a moment's reflection reveals that such a position is not sustainable. For recall that if the radical skeptic is right, then we have no good reasons *at all* for believing what we do. That is, we think that we have such reasons, because of what we seem to experience (and remember etc.), but all of these reasons turn out to be illusory once it becomes clear that we cannot rationally exclude radical skeptical hypotheses. It follows that epistemic externalism as an anti-skeptical strategy cannot advance the moderate epistemic externalism about knowledge just described, but rather must embrace a much more radical view: knowledge is only ever of a kind that involves the satisfaction of externalist epistemic conditions, as none of our beliefs (none of our beliefs about the external world at any rate) ever satisfy an internalist epistemic condition.

The nub of the matter is what the epistemic externalist will say about our belief that we are not the victims of radical skeptical hypotheses, like the BIV hypothesis. The moderate externalist line sketched above seemed to suggest that we merely know such a proposition in an externalist sense of the term, but that this is compatible with our being able to know all

manner of everyday propositions about the external world in a manner that would satisfy internalist demands. At any rate, we are able to have reflectively accessible rational support for our everyday beliefs, even if such support isn't of a kind that would enable us to rationally exclude radical skeptical hypotheses.

But this way of understanding moderate epistemic externalist anti-skepticism isn't very plausible on closer inspection. For according to radical skepticism, if one is unable to have a rational basis for excluding radical skeptical hypotheses—and thus one fails to have internalist knowledge that such a scenario doesn't obtain—then one thereby fails to have any rational basis at all for one's everyday beliefs as well. Accordingly, the balancing act just described, whereby there is a kind of internalist knowledge that coexists with our generally having merely externalist knowledge, including of the denials of radical skeptical hypotheses, is simply untenable. The upshot is that if epistemic externalism is to offer us a resolution to the radical skeptical problem, then it needs to embrace a much more radical form of epistemic externalism, one that accepts that we have no rational basis for our everyday beliefs at all (at least insofar as those beliefs concern the external world anyway).

Adopting an austere epistemic externalism of this revisionist sort in response to the radical skeptical paradox would put it at a major dialectical disadvantage. In Chapter 1, we saw that responses to paradoxes can be divided into those that are undercutting and those that are overriding. While the former effectively demonstrates that the paradox in question is illusory, and hence that there isn't the fundamental tension in our concepts that is alleged, the latter grants the force of the paradox but argues on independent grounds that we can resist it. Marshalling an austere epistemic externalism in the service of one's anti-skepticism would fall into the overriding category, as it would be a highly revisionary way of responding to the paradox. The problem facing such anti-skeptical proposals, however, is that they leave us torn between the intuitions that drive the paradox and the theoretical position that aims to resolve that paradox, which is why they struggle to offer us the kind of relief from intellectual anxiety that we seek from a solution to the problem at hand. As we noted in Chapter 1, this is especially so when we are dealing with a paradox that has existential force, like the problem of radical skepticism.

There is a further challenge facing this particular, austere, epistemic externalist anti-skeptical stance, which is that it is actually quite difficult to

understand how it is possible to have knowledge of the denials of radical skeptical hypotheses even by externalist lights. Presumably the idea is that since the knowledge is rationally ungrounded, it follows that the Humean skeptical worry doesn't get a grip, as that turns on the claim that there is no non-circular way of *rationally grounding* one's belief in the target proposition (in this case concerning the denial of a radical skeptical hypothesis). But even so, we are owed a story about what the epistemic basis of this belief is which enables it to be knowledge, given that it is not an epistemic basis cast in terms of reflectively accessible reasons.

The difficulty, however, is that it is unclear what this epistemic basis is supposed to be, precisely because it is unclear what the source of the belief is. This is important, since that is usually crucial to working out whether the belief satisfies an externalist epistemic condition. In order to determine whether a belief satisfies the rubric laid down by either process reliabilism or virtue reliabilism, for example, we first need to know what belief-forming process is in play. But what is the belief-forming process that results in our belief that we are not victims of radical skeptical hypotheses? This doesn't seem to be a belief that we simply acquire in virtue of possessing our ordinary cognitive faculties, in the way that we might acquire our perceptual beliefs. Indeed, prior to engaging with the skeptical problematic, we would surely not regard someone as having *any* beliefs about this matter. But if that's right, then it suggests that this belief is in fact formed as part of an *inference* from one's prior beliefs, such as one's perceptual beliefs. The problem, however, is that it is hard to see why Humean skepticism does not now get a grip after all, for isn't such an inference clearly epistemically circular? More generally, if it's right that our commitment to everyday propositions about the external world already presupposes the falsity of radical skeptical hypotheses, as the Humean skeptical problematic alleges, then how, for example, can it be the manifestation of cognitive virtue that one forms one's belief in the denials of radical skeptical hypotheses by inferring them from one's prior perceptual beliefs about the external world?

There is a further way that we might develop the epistemic externalist response to radical skepticism, which is to use it to motivate a denial of the Closure principle for knowledge. In particular, one could appeal to epistemic externalism to explain why one can know everyday claims about the external world even though one is unable to know the denials of radical skeptical hypotheses. That would avoid many of the problems just noted

by severing the connection between everyday knowledge and knowledge of the denials of radical skeptical hypotheses altogether, though of course it would come at the cost of entailing the rejection of Closure. Some epistemic externalists have been tempted down this route. Indeed, it has been argued—most notably by Ernest Sosa (e.g., 2021)—that virtue reliabilism lends itself to such a proposal.

Recall that virtue reliabilism understands knowledge in terms of the manifestation of cognitively skillful belief-forming processes. In particular, knowledge is, roughly, what results when one's cognitive success (one's true belief) is appropriately attributable to one's manifestation of cognitive skill. Knowledge is thus understood as a specifically cognitive kind of achievement—just as achievements in general are successes that are appropriately attributable to one's skillful performance (rather than, say, to dumb luck), so knowledge is a cognitive success that is appropriately attributable to one's cognitively skillful performance. Indeed, this is why a virtue reliabilist like Sosa holds that knowledge is lacking in Gettier-style cases, since while there is the conjunction of cognitive success and cognitive skill on display, the former is not appropriately attributable to the latter (but is rather attributable to the epistemic luck in play). The subject's cognitive success is thus not a cognitive achievement of theirs, and hence not knowledge.

Interestingly, however, it seems that one's achievements more generally can take place against a backdrop of certain background conditions obtaining, even if one has no rational basis for thinking that they obtain, and that suggests that the same ought to be the case with regard to one's cognitive achievements too, and thus one's knowledge. To use one of Sosa's examples, a baseball pitcher in a nighttime game might skillfully pitch the ball while taking it for granted that the lighting will continue working as he makes his pitch. He couldn't make the pitch were the lighting to fail, so that it is working is important to his successful manifestation of his pitching skill. But that it is working seems like a background condition for his performance, and as such is the kind of thing that the pitcher can legitimately take for granted. In particular, we wouldn't think that it is part of the pitcher's skillful performance that he checks on the lighting before making the pitch, in the way that we would think it is part of his performance that he checks on the strength and direction of the wind in the arena. So long as this background condition obtains, then the pitcher's success can be appropriately attributable to his skill, and hence count as an achievement on his part.

Transplanting this idea back to the cognitive realm, perhaps we can treat our commitment to the denials of radical skeptical hypotheses in the same vein. In coming to know propositions about one's environment, it is necessary to take for granted that one is not, say, a BIV. But one's cognitively skillful performance in this regard does not involve one being able to rationally exclude this possibility. All that matters is that it is *de facto* excluded—i.e., that one is not a BIV. So long as this background (metaphysical) condition obtains, then one can have knowledge of specific claims about one's environment—i.e., one's cognitive success in this regard will be appropriately attributable to one's cognitive skills. In particular, it's not necessary to know that one is not a BIV, or even to have any reason at all for excluding this error-possibility; all that matters is that it is indeed false.

Of course, if one takes this route, then one is committed to denying the Closure principle, since one can know everyday claims that entail the denials of skeptical hypotheses without being able to know that these skeptical hypotheses are false. Given the intuitive nature of the Closure principle, this would thus also be an overriding response to the problem of radical skepticism. Nonetheless, there would be a rationale in support of this revisionary proposal, which is to argue that we should treat knowledge as a cognitive achievement, and thus as akin to other kinds of achievement, in that it can coexist with there being background conditions that need to obtain for that achievement that one is effectively merely assuming obtain (as opposed to knowing that they obtain, or otherwise skillfully determining that they obtain). (For further discussion of Sosa's more recent proposal, see Chapter 6, Section 6.5.)

In any case, denying the Closure principle is a way of responding to the problem of radical skepticism in its own right, one that can be employed in responding to this problem independently of whether one endorses epistemic externalism. We will thus be examining this form of anti-skepticism in Chapter 3. We will also be considering the general idea that our commitment to the denials of radical skeptical hypotheses might be epistemically legitimate even though we can't have knowledge in this regard when we discuss the Wittgensteinian notion of a hinge commitment (in Chapters 4, 6, and 7).

Further reading

See Putnam (1981) for his main defense of content externalism and its relevance for radical skepticism. See also Putnam (1975) for an earlier

statement of his content externalism, and Burge (1993; 2003) for an important defense of a related form of content externalism. Whether Descartes is committed to content internalism is a controversial exegetical issue. For some useful discussion, see Burge (1986), Normore's (2007) commentary, and Burge's (2007) response. For a sample of discussions of the putative tension between content externalism and our first-person authority when it comes to our mental states, see Nuccetelli (2003). See also Burge (1988), Boghossian (1989), and McKinsey (1991). For an overview of Putnam's work on radical skepticism, see Pritchard (forthcoming b); for specific discussion of his response to BIV-style skepticism, see Pritchard and Ranalli (2016). Davidson's work on radical interpretation draws heavily on influential earlier work on radical translation by Quine (1960). See Davidson (1983) for his key discussion of the omniscient interpreter argument. See also Davidson (1975; 1977) and his later qualifying comments in this regard in Davidson (1990). For more on his notion of triangulation and its import for content externalism, see Davidson (2001). For an overview of Davidson's response to radical skepticism, see Pritchard (2013). For some key critical discussions of Davidson's anti-skepticism, see Williams (1988), Ludwig (1992), and Genova (1999). See Stroud (1968) for an influential critique of transcendental arguments. See Chalmers (2017; 2018) for further development of some of the ideas introduced in Chalmers (2005). For further discussion of the epistemic externalism/internalism distinction, see Vahid (2010) and Pappas (2014). Accessibilism with regard to epistemic internalism is defended by Chisholm (1977) and BonJour (1985, Chapter 2), among others. For an influential defense of mentalism, which is the main competing proposal in this regard, see Conee and Feldman (2004). We've here followed the standard post-Gettier narrative regarding how Gettier (1963) definitively undermined a prevailing internalist conception of knowledge. See Dutant (2015) and Le Morvan (2017) for critical discussion of this narrative. Process reliabilism is usually associated with the work of Goldman (e.g., 1979; 1986). For an important development of virtue reliabilism, including its potential application to the problem of radical skepticism, see Greco (1999; 2000; 2010). For two related views, see Plantinga's (1993) proper functionalism and Goldman's (1992) own later reflections on process reliabilism. See also Sosa's (1991; 2007; 2009; 2011; 2015; 2021) influential development of a fully-fledged virtue epistemology that is more receptive to certain insights from epistemic

internalism, and which accordingly includes a kind of 'two-tiered' conception of knowledge (albeit where the second-order tier is not meant to be specifically internalist). For further discussion of 'modest' forms of epistemic externalism about knowledge, see Pritchard (forthcoming a). The particular anti-skeptical line attributed to Sosa here is given its clearest expression in Sosa (2021, Chapter 7). For a prominent defense of a version of virtue epistemology that is cast along broadly epistemic internalist lines, see Zagzebski (1996). For a survey of recent work on virtue epistemology, see Kvanvig (2010) and Alfano, Greco, and Turri (2017). For a survey of epistemic externalist responses to radical skepticism, see Fumerton (2010).

3

THE DENIAL OF THE CLOSURE PRINCIPLE AND CONTEXTUALISM

3.1 Introduction

In this chapter, we will examine responses to the Cartesian skeptical paradox that take issue with its form or scope of application. In order to introduce them, it is useful to look at what exactly they object to in the Cartesian skeptical argument in its simpler form:

(1) I don't know that I am not a BIV.

(2) If I know I have hands, I know I am not a BIV.

(C) I don't know that I have hands.

According to some theorists, the problem with this argument is that it relies on the Closure principle (see Chapter 1, Section 1.3). That is, the

DOI: 10.4324/9780429057946-4

principle that says, roughly, that if one knows that *p* and one knows that *p* entails *q*, then one knows that *q*. For it is only if one adheres to it that, from the fact that one does not know that one is not a BIV, one can derive that one does not know that one has hands. To see this clearly, consider the following instance of the Closure principle:

(1) I know that I have hands.

(2) I know that if I have hands, then I am not a BIV.

(3) I know that I am not a BIV.

If the Closure principle holds, then, by contraposition, if I don't know that I am not a BIV, then I don't know that I have hands either. If, in contrast, the Closure principle doesn't hold, from not knowing that I am not a BIV, it doesn't follow that I don't know that I have hands. We will look at this proposal in more detail in Section 3.2.

Yet, as we saw in Chapter 1, many theorists think that the Closure principle is irrevocable, for reasons we have already explored. Moreover, they point out that renouncing it would entail endorsing "abominable conjunctions," such as 'I know I have hands, but I don't know I am not a BIV,' which, on the face of it, strike us as clearly mistaken. Their strategy, therefore, consists in retaining the Closure principle while objecting to the scope of application of the Cartesian skeptical argument. According to contextualists, the argument goes through only in those typically philosophical contexts in which skeptical hypotheses are relevant. Yet, in ordinary contexts, when such hypotheses are not relevant, the argument does not go through and, given the Closure principle, by knowing that we have hands, we also know that we are not BIVs. We will consider this proposal in more detail in Section 3.3.

3.2 The denial of the Closure principle

An influential attempt to block the skeptical paradox, which has been proposed in the literature, consists in the denial of the Closure principle, which, as we have seen, is fundamental to the derivation of the Cartesian paradox. This has traditionally been proposed as a *damage limitation* strategy for skepticism (Dretske 1970; Nozick 1981). Indeed, it is taken for granted that one cannot know (or have a justification to believe) that one is not dreaming (or

that one isn't a BIV). As we saw, this is the Cartesian skeptical hypothesis that constitutes step (1) of the Cartesian argument. To take it for granted, therefore, is already a big concession to skepticism. However, this strategy tries to avoid the conclusion (C), *viz.*, that we have no knowledge (or no justification for our beliefs) about physical objects in our surroundings by denying that the inferential principle which would allow one to derive it is in fact invalid. If it is invalid, then there is no real Cartesian skeptical paradox to be responded to. In that sense, this strategy allows one to present an undercutting response to the Cartesian skeptical paradox (see Chapter 1, Section 1.3). Yet, as we shall see, it is based on a revisionary understanding of the Closure principle, whose validity is either denied or confined to specific sets of premises and evidential conditions, based on an externalist account of knowledge, which, as we saw in Chapter 2, Section 2.3, is revisionary of some aspects of knowledge itself. Furthermore, as noted, by conceding (1)—that we don't know we are not BIVs—this strategy concedes quite a lot to the Cartesian skeptical paradox. For these reasons, it could also be listed as an overriding response to the Cartesian skeptical paradox.

This strategy of response to the Cartesian skeptical paradox often goes hand in hand with an externalist conception of knowledge (and/or justification). Briefly, the idea is that in order to have knowledge (or justification) for ordinary empirical propositions, it would suffice to exclude all *relevant* alternatives without having to rule out the irrelevant ones. Now, when is an alternative relevant (in a metaphysical and not pragmatic sense)? There are various ways of understanding this idea. One is that, whenever an alternative is being considered, it then becomes relevant.

Another way of understanding the notion of relevant alternative is by considering whether it could *easily* happen. We can clarify this idea by considering possible worlds and their relationship. The relevant alternatives are those that are given in the possible worlds, which are closest to the actual one. For example, a relevant alternative for the actual world in which 'I have hands' is true, is a possible world where I don't have hands because I have been the victim of a car accident. Alas, the possible world in which such a contingency occurs does not appear to be that distant from the real one. That is, not too many things should differ from what obtains in the actual world for that contingency to take place. An irrelevant alternative, by contrast, is the one where I do not have hands because I am a brain in a vat, who is hallucinating having hands. Such a possible world is much further away than the one in which I would have lost my hands,

because many more things would have to differ from what obtains in the actual world for that contingency to obtain.

Now, the notion of relevant alternative is surely relative. The fact that one is only dreaming of seeing one's hands doesn't seem to be at all a remote hypothesis; certainly not as remote as that of being just a brain in a vat, or a victim of the Matrix, or deceived by an evil genius. Yet all these alternatives could become relevant when dealing with skepticism.

Furthermore, there is an issue about whether we can confidently hold that the *actual* world is as we normally think of it, as opposed to being one of these worlds in which everything looks the same to us but we are in fact brains in a vat, or deluded by powerful machines or by an evil demon. That is to say, it seems to be begging the question against Cartesian skeptical scenarios to hold that the actual world is as we normally think of it. Let us leave these preoccupations aside for the time being, to see how embracing such an account of knowledge would allow one to call the Closure principle into question.

According to Dretske and Nozick, the subject would know that she has hands because in the possible world in which a relevant alternative obtains, she would not believe that she had them. By contrast, she wouldn't know that she is not dreaming or that she is a victim of the Matrix because, in the possible world in which these alternatives obtain, she would not believe that she is dreaming, or that she is a victim of the Matrix.

Following Nozick, the belief that I have two hands 'tracks the truth,' because *the same methods* used to form that belief would return an opposite verdict in those nearby possible worlds in which that belief were false—that is, in those possible worlds in which I wouldn't have hands due to a car accident. By contrast, the belief that one is not dreaming or that one is not a victim of the Matrix does not track the truth, because one would have it even if it were false. As we saw in Chapter 1, Section 1.3, any method we may use to form the belief that we are not dreaming, or that we are not BIVs, or aren't deceived by the Matrix or by an evil demon, would be compatible with the possibility that we might just be thinking of applying it and could not be used to determine the falsity of these scenarios. For instance, my sensory impressions and my memories could just be the same as the ones I have in the actual world, in which I am not dreaming or I am not a victim of the Matrix, and make me think that I am not currently dreaming or the victim of the Matrix, while I am. Therefore, the belief 'I have hands' is 'sensitive'—as it is customary to call it in the

literature—with respect to these different scenarios, while the belief 'I am not dreaming/a BIV' is not. For that reason, only the belief 'I have hands' amounts to knowledge.

More specifically, Nozick proposes the following conditions for knowledge:

(1) P is true.

(2) S believes that P.

(3) Sensitivity condition: If it were the case that not-P, S would not believe that P (by application of the same method).

(4) If it were the case that P, S would believe that P (by application of the same method).

Given what we have just seen, while 'I have hands' would satisfy these conditions, 'I am not a BIV' wouldn't. In particular, it would flout condition (3).[1] Thus, the following is a counterexample to Closure:

(1) I know that P ('I have hands').

(2) I know that if P ('I have hands'), then Q ('I am not a BIV').

(C) I know Q ('I am not a BIV').

For, while we do know that we have hands, and that having hands entails that one is not a BIV, we don't know that we are not BIVs.

The underpinning conception of knowledge has been subject to many objections, which we won't explore here. We will simply mention the fact that, nowadays, most modal accounts of knowledge replace the Sensitivity condition with a Safety condition:

Safety condition: In all nearby worlds where *S* believes that *p*, *p* is not false.

Be as it may, the criticism of the Closure principle does not necessarily depend on an externalist conception of knowledge or justification. Indeed, we will now analyze its internalist variant in more in detail.

To that end, we will consider the Closure principle for justification internalistically conceived. If such a notion of justification is taken to be a

necessary component of knowledge, then what we will be saying will carry over to knowledge as well. Yet, in principle, we may divorce these two notions—that is, we might retain an externalist conception of knowledge while holding on to an internalist notion of justification, which in turn would not be necessary for knowledge. Hence, for expository reasons, it is better to focus just on the denial of Closure for justification internalistically construed. Now, suppose you go to a zoo and see an animal in a cage that looks like a zebra. It seems that you are justified in believing *p* 'Here is a zebra,' based on your current perception. Furthermore, it seems that you are *a priori* justified in holding that if this is a zebra (*p*), then this isn't a cleverly disguised mule (*q*). Yet, by application of Closure (understood diachronically, cf. Chapter 1, Section 1.3), you should now also possess a justification for (*q*) 'This isn't a cleverly disguised mule.' Yet, it seems weird to hold that your perception, together with a bit of *a priori* reasoning, could give you a justification, which you could deploy to support your belief *q*. After all, it seems entirely compatible with your subjectively available evidence at that time that you might be seeing a cleverly disguised mule. And if that is the status of your evidence, you cannot adduce it to support your belief in *p* over *q*.

Notice that this counterexample has a specific structure. Namely, not-*q* is incompatible with *p* and yet sustainable on the basis of the same evidence one may have at one's disposal. Thus, one may take this to show not that the Closure principle never holds, but that it doesn't hold only when these more specific conditions are further met by *p*, not-*q* and the available evidence.

Interestingly, skeptical scenarios would fulfill the conditions for Closure failure. As we saw in Chapter 1, Section 1.3, the Underdetermination-based version of the Cartesian skeptical paradox shows exactly that. For one's available evidence as of a hand here would be compatible both with *p* 'Here is my hand' and with not-*q* 'I am a BIV having a hand-like experience,' where clearly *p* and not-*q* are incompatible.

If this is the case, and Closure does not hold unconditionally, one may conclude that Descartes was right to believe that we cannot have a justification to believe that we are not dreaming. Yet, he was wrong to maintain that this prevented us from having perfectly justified beliefs about physical objects around us.

By contrast, arguments such as the following one:

(1) I know there is a 200-page book here.

(2) I know that if this is a 200-page book, then it has more than 100 pages.

(3) I know that this book has more than 100 pages.

Would not meet the conditions for Closure failure. Surely *p*—'This book has 200 pages'—is incompatible with not-*q*—'This book doesn't have more than 100 pages.' Yet, my current evidence—that is reading the number on the last page of the book after reading the whole of it and having checked page numbers as I went along—is not compatible with not-*q*. If so, one may still retain the Closure principle but reduce its scope of application. In other words, even if one is convinced by the zebra example, one is not forced to deny the validity of the Closure principle *absolutely*. For, as we have just seen with the book example, that principle would still be valid in specific, determinable cases. Thus, the moral to draw is simply that the Closure principle (for justification internalistically understood, at least) is not *unconditionally* valid.

In effect, what is being proposed here is an instance of the principle that rules (like the Closure principle) admit of exceptions (when certain other conditions obtain, as we have just seen). This is a very important principle because limiting the range of application of a rule (in a principled way) allows us to preserve the rule in all other cases. Thus, it is important to note that overstating the claim being made by detractors of Closure by saying that they deny its validity *tout court* is both a mistake—for they don't, as they merely deny its unconditional validity—and a problem in its own right. For now a supporter of Closure will have to hold, implausibly, that Closure holds everywhere, even when certain special conditions relating *p*, *q* and a subject's available evidence are met.

Closure and Transmission

The distinction between the two formulations of the Closure principle introduced in Chapter 1, Section 1.3 (with the substitution of justification for knowledge):

> *Closure principle:* If one is justified in believing that *p*, and one knows (or is justified in believing) that *p* entails *q*, then one is justified in believing that *q*.

and

> *Closure principle (diachronic):* If one is justified in believing that *p*, and one competently deduces *q* from *p*, thereby coming to believe that *q* on this basis while retaining one's justification for the belief that *p* throughout, then one is justified in believing that *q*.

Has not always been clear, as it became standard only recently. Before it did, it was pointed out (Wright 1985) that the zebra example can be reformulated as a case of Transmission failure of justification from the premises to the conclusion. Namely, as a case in which occurs what Humean skepticism believes happens with a proof for the existence of the external world, like Moore's, which starts from a premise such as 'Here is a hand' (cf. Chapter 1, Section 1.4). Accordingly, the zebra example does not show that the Closure principle (understood synchronically) is not valid, but only that not every logically valid argument transmits justification from the premises to its conclusion.

> *Transmission of justification principle*: a logically valid argument transmits justification from the premises to the conclusion if (and only if), by running it, it is possible to acquire a first justification to believe the conclusion.

Equivalently, an argument transmits a justification from its premises to its conclusion—and thus allows one to acquire a first justification for the conclusion—if (and only if) having a justification for the premises does not presuppose that there is already a justification for its conclusion. According to Wright, the zebra example and Moore's proof would then be cases where in order to have a justification for the respective premises 'Here is a zebra' and 'Here is a hand,' one must already have a justification for 'This is not a cleverly disguised mule' and for 'There is an external world,' or for 'I'm not dreaming,' respectively. For one's own experiences as of a striped animal and as of a five-fingered pinkish expanse could not give one a justification for 'Here's a zebra' or 'Here's a hand' just by themselves. As we noted, those experiences would be compatible with 'Here is a cleverly disguised mule' and 'I am a BIV hallucinating having a hand,' respectively. Therefore, the zebra example and Moore's proof are epistemically defective but do not prove that the Closure principle as applied to justification (understood synchronically) does not apply, but only that the Transmission of justification fails.

It is important to note that the failure of Transmission understood along these lines *entails* the validity of the Closure principle understood synchronically.[2] For it is only when there must already be a justification for the conclusion in order to have a justification for the premise(s), that the argument cannot give one a first justification to believe the conclusion. Therefore, it is only when we have the following argumentative structure that it may be a case of justification's Transmission failure:

(1) One has a justification for *p*.

(2) One has a justification for *p* entails *q*.

---

(3) One has a justification for *q*.

Where having a justification for q is a necessary condition for having a justification for p. But this pattern, as we have just seen, is just the Closure principle for justification understood synchronically. That is, as a principle that only tells us how the epistemic operator—i.e., let it be knowledge or justification—is distributed with respect to the premises and conclusion of a given argument. On the contrary, the Transmission principle should be understood as a generative principle, in the sense that it considers the conditions under which justification (or knowledge) of a conclusion can be obtained by running an argument. Thus, a way to represent the difference between the failure of Closure (synchronically understood) and the failure of Transmission is as follows:

Failure of the Closure principle

(1) One has a justification for *p*.

(2) One has a justification for *p* entails *q*.

---

(3) One does not have a justification for *q*.

Failure of Transmission of justification principle

(1) One has a justification for *p*.

(2) One has a justification for *p* entails *q*.

---

(3) One has a justification for *q*.

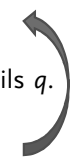

We will have to go back to the relationship between these two principles in Chapter 7. In fact, it seems that, by more contemporary lights, the principle of Transmission is very similar to the principle of Closure *diachronically* understood. If so, distinguishing between Closure and Transmission to claim that in the zebra case the latter fails while the former doesn't would be quite pointless, not because Closure is now typically understood diachronically instead of synchronically, but, far more interestingly, because the two principles—Closure diachronically understood and Transmission—would be identical (or at least sufficiently similar). Hence, it is only to be expected that if one failed, the other would do too. In fact, as we saw, the zebra example is a counterexample to Closure for justification *diachronically* understood, which hinges on the fact that the same evidence at one's disposal would support both *p* and not-*q*, where *p* and not-*q* are incompatible. Thus, an argument exhibiting that pattern could not provide one with a first justification to believe its conclusion. For one should already be justified in holding it, and thereby be justified in discounting not-*q* (or other hypotheses incompatible with *q*), in order to be justified in holding its premises in the first place.

To conclude, if one understands Closure synchronically, then the zebra argument is not a counterexample to it, but to the principle of Transmission. If, by contrast, Closure is understood diachronically, then the zebra argument is both a counterexample to Closure and to Transmission (understood along the lines presented so far—that is, as entailing that Transmission would fail only if the justification one has for the premises of a given argument depended on already possessing a justification for its conclusion).

"Abominable conjunctions"

Another more recent objection to the denial of the Closure principle is that it would give rise to "abominable conjunctions" (DeRose 1995). One could in fact say (or judge) 'I know (or I have a justification to believe) that I have hands, but I don't know (or I have no justification to believe) that I'm not a BIV.' Yet, on the face of it, how could it be correct to affirm the first conjunct—'I know I have hands'—and yet be correct to affirm the second—'I don't know I am not a BIV'—that leaves it open the possibility that I might be a handless brain floating in a vat?

Now, one might argue, with Sosa (2000), that it is one thing to correctly describe our epistemic situation, and another is to explain the legitimacy of certain linguistic assertions. To put it differently: it is one thing

to clarify the conditions under which knowledge and justification obtain, and another is to clarify under which conditions it is appropriate to make epistemic assertions. Thus, it may be the case that one does not possess knowledge or justification for 'I am not a BIV' and yet have them for 'I have hands,' while it may be weird to assert that much. Epistemology, after all, is interested only in the conditions in which knowledge and justification obtain, while the conditions under which assertions are felicitous or not, if anything, pertain to the philosophy of language.

Now, it is evident that this response strategy must clarify why it sounds strange to make assertions that according to one's theory are true. Moreover, we should note that it would be a little awkward to be in a situation in which one is in possession of a truth that could not be properly asserted.

Pritchard (2005b) has argued that the assertion sounds strange because it violates the Gricean maxim of quality, which prescribes that one should only assert what one knows or has a justification for. Now, if it is true that we have no justification to believe or knowledge of the fact that we are not BIVs, then it is obvious that we commit a pragmatically infelicitous speech-act by asserting it.

However, it seems that the problem is a little thornier. In fact, the assertion that sounds weird is not 'I am not a BIV'; rather, it is 'I know (or I have justification to believe) that I have hands and I don't know (or have no justification to believe) that I'm not a BIV.' This is a true assertion, since according to the detractors of the Closure principle both conjuncts are true. Furthermore, in their opinion, one would also have very good justifications for both conjuncts. Thus, it is not clear how the appeal to the Gricean maxim of quality can explain why the assertion sounds strange. After all, that maxim seems to be respected (by the lights of detractors of Closure).

Moreover, even just *judging* the conjunction is very counterintuitive. This impression, however, cannot be explained by reference to conversational maxims that regulate the pragmatic propriety of the speech-act of assertion. Hence, it is not clear how one can respond to the non-linguistic version of the objection.

We will return to the problem of abominable conjunctions once in possession of further notions and distinctions (which will be introduced in Chapter 7, Section 7.5) that will help us clarify how one might correctly assert these *prima facie* abominable contents, and why, therefore, this objection is not fatal after all.

3.3 The contextualist response: Cohen and DeRose

Contextualists agree with Dretske and Nozick that we cannot know that we are not BIVs, at least in some circumstances, as we shall see. However, they do not reject the Closure principle (DeRose 1999; Cohen 2005a; 2005b). Yet, they don't want to conclude that we do not know that we have hands. In order to do so, they claim that the truth conditions of 'S knows that *p*'—i.e. the proposition expressed by that statement—vary depending on the context of the subject who asserts it. It is important to note that the subject who asserts the statement might be different from the subject to whom knowledge is attributed. If the subject who makes the attribution of knowledge—'S knows that *p*'—is in the ordinary context, the standards are low, and the subject S knows both that she has hands and that she is not a BIV. Conversely, if the subject who makes the knowledge attribution—'S knows that *p*'—is in the skeptical context, the standards are much higher, and S doesn't know either that she is not a BIV, or that she has hands. Consequently, the corresponding knowledge ascription would be true in the former case and false in the latter, while nothing has changed in S's overall epistemic conditions.

That said, it is clear that, as in the case of the denial of the Closure principle, we are dealing with a damage limitation strategy: contextualism doesn't disprove skepticism but it tries to prevent it from having any consequences on our ordinary epistemic practices, which enable us to have knowledge of physical objects in our surroundings. Specifically, contextualism proposes an overriding solution to the Cartesian skeptical paradox because, by proposing a revision of our ordinary understanding of knowledge as invariant, it contains the effects of the paradox: contrary to the moral typically drawn from the Cartesian paradox, for contextualists, knowledge would still be possible, at least in ordinary contexts, while it wouldn't in relatively far-fetched contexts.

The contextual sensitivity of 'know'

In order to illustrate the contextual variability of knowledge attributions, suppose that Lucy wants to deposit a check in the bank and that Mary tells her that the bank is open on Friday afternoon since she looked at the opening hours. Intuitively, we would say that Mary knows that the bank is open on Friday afternoon. Let us now suppose that Lucy has an absolute need to deposit that check, because if she does not, she is going to get cut

off. All things being equal, according to contextualists, Lucy's different practical interests change the truth conditions of 'Mary knows that the bank is open on Friday afternoon.' In this sense, the standards to be met, in order to know that the bank is open on Friday afternoon, have raised. Therefore, it is now false that Mary knows that the bank is open on Friday afternoons.

This is just an illustration of the contextualist take on knowledge ascriptions. When applied to Cartesian skepticism, however, practical interests are quite irrelevant. Rather, we are now dealing with error possibilities raised by engaging with skeptical scenarios. To clarify the structure of the contextualist strategy in this connection, it is useful to point out that, in the derivation of the Cartesian paradox, we have the following premises:

(1) One can't know 'I am not a BIV.'

(2) One knows 'I have hands.'

(3) If one knows 'I have hands,' then one knows 'I'm not a BIV.'

Now, according to the contextualist, when the lower standards are in place, because skeptical error possibilities aren't salient, (2) is true, while (1) isn't; when the higher standards are in place, in contrast, because skeptical error possibilities are being considered by engaging with skepticism, (1) is true, while (2) is false, whereas (3) is true in all cases. The paradox from which one would get that we *never* know that we have hands (and that we are not BIVs) is only apparent, because it depends on a conflation between low- and high-standard contexts. In fact, it is only in the skeptical context that (1) is true, and, given (3), we can derive the denial of (2). Conversely, in the ordinary context we know both that we have hands (2), and the denial of (1)—that is, that we know we are not BIVs.

Contextualism has been the object of criticism because it is not clear that 'know' is as context-sensitive as clearly are 'I,' 'tomorrow,' 'flat,' 'tall,' 'sharp,' etc. The linguistic data and the intuitions of the speakers do not seem to clearly confirm this hypothesis about 'know.' For instance, Earl Conee (2005) has argued that the variability in knowledge attributions sometimes found among speakers may depend on the passage from a loose-speaking context to a strict-speaking context. Thus, the truth conditions of 'S knows that p' would not really be context-sensitive. When we

speak properly, 'S knows that p' is absolutely true (or false). However, for various pragmatic reasons, sometimes we don't speak properly. One example is when we say that France is hexagonal. There is no context in which this is true. Yet, sometimes, if our aim is just to give an idea of the shape of the country we act as if it were true. Conee also rightly points out that in philosophy we are not interested in the conditions for knowledge 'roughly speaking,' but, rather, in the conditions that need to obtain to have knowledge properly so regarded. Thus, in philosophy, the truth conditions of 'S knows that p' would be absolute and fixed according to stringent conditions. (Of course, for Conee, it is still an open issue if these are the conditions posed by Cartesian skepticism, or less rigorous ones.)

Another related objection has been raised by Ernest Sosa (2000). It consists in saying that contextualism confuses the conditions under which it is (or is not) appropriate to make assertions of knowledge attributions, with the conditions under which knowledge obtains. Only the latter are relevant to epistemology, while the former, at most, may concern the philosophy of language.

Cohen (2005) has responded by pointing out that speakers' intuitions are not always conclusive in deciding whether a term is context-sensitive. Think, for instance, of the case of 'simultaneous' before and after the advent of relativity theory. In addition, Cohen has also highlighted how in philosophy we oscillate between different standards for the attribution of knowledge. When we consider Gettier cases, for instance, we are willing to admit that we know many things; not so when we consider Cartesian skepticism.

Finally, epistemic contextualism does not deal with *assertibility* conditions of 'S knows that p,' but with its *truth* conditions. As we have seen, contextualism claims that the proposition expressed by that statement is context-sensitive. Thus, according to contextualists, 'knowledge' itself is context-sensitive, and not only the pragmatic propriety of asserting its attribution to a given subject.

The problem of disagreement, retraction and MacFarlane's epistemic relativism

From the point of view of the philosophy of language, it is possible to raise other objections to contextualism. The first is that if the truth conditions of 'S knows that p' vary based on the context, then it is not clear how there

can be genuine disagreement between a subject who were to assert that we don't know that we are not BIVs or that we have hands and another subject who denied that. For the two of them would employ different standards and would express different propositions with their respective statements. To see the point, consider that, after all, we typically don't think that A and B disagree with each other when the former says, 'I was born in Milan,' and the latter denies it, since they are talking about different people.

The second objection is that if, for contextualists, the relevant epistemic standards are those of the subject who makes the knowledge attribution, then it is not clear how the attributor could retract her statement later. After all, when she made it, different standards were in place from those that applied to her at the time of retraction. However, contextualism implies that the standards applied at the time the knowledge attribution was made are the right ones to semantically assess 'S knows that p.' For this reason, they remain the only ones valid even when, later on, it would seem they are no longer operative.

To both these objections one can respond by introducing a different form of contextualism—known in the literature as "alethic relativism" (or "assessment sensitivity")—proposed by John MacFarlane (2005; 2014, Chapter 8). The stratagem is to say that the standards that determine the truth of 'S knows that p' are neither those of the one who has asserted 'p,' nor those of the one who asserted 'S knows that p,' but the standards of the one who *assesses* the attribution of knowledge, which in principle could be a different subject from the other two (but of course it could also be identical). By these means, the problem of retraction could easily be solved, since the assessor's standards may vary over time and therefore the same person may retract one's previous statements 'S knows that p,' if her standards happen to change over time.

The problem of disagreement, however, is more difficult to solve. Certainly, the truth conditions of 'S knows that p' would not vary with the change of standards of those who make or deny the attribution. Yet, there is no doubt that the proposition would be evaluated semantically by means of different standards. Not being able to establish that only one of them is the correct one, the problem of disagreement would rise again, although moved to a different level.

To address this worry, MacFarlane proposes to understand the disagreement as relative to the assessor's standards. That is to say, with respect to the standards of the assessor A, who considers 'S knows that p' true, B's

denial of 'S knows that *p*' generates a contradiction, and that is enough to account for the idea that A and B disagree with each other.

The objection that can be made to this account of the disagreement between A and B is that it depends on not taking into consideration the fact that B's statement '"S knows that *p*" is false' is made according to B's standards, which, *ex hypothesi*, differ from those of A. It is certainly doubtful that this extrapolation of B's claim from the standards within which it has been made is either legitimate, or the proper way of accounting for disputes between advocates of different positions regarding 'S knows that *p*.'

Contextualism and philosophy

Whatever the fate of these objections to contextualism may be, it seems appropriate to turn our attention to another kind of perplexity, which admits *pro tempore* that contextualism is plausible compared to the linguistic data we have, to the oscillation between different standards in epistemology, and finally as a thesis about knowledge, and not only about the pragmatic propriety of knowledge attributions.

Let us start by asking what determines the transition from one context to another. According to DeRose (1995), it is enough to assert that S knows that *p* (or its negation). Now, if the belief in question is 'insensitive'—that is, such that we would believe it to be true even in the closest possible world where it were false, like 'I'm not a BIV'—then that would make the standard automatically higher. Therefore, in this case, S would not know either that she is not a BIV, or that she has hands. However, if the belief in question is 'sensitive'—that is, such that we wouldn't believe it to be true in the closest possible world in which it was false, like 'I have hands'—then standards would remain low. We would then know both that we have hands and that we are not BIVs. Therefore, standards rise in such a way as to deprive us of knowledge, even of propositions we commonly think we know, only when skeptical hypotheses like the Cartesian ones are introduced.

Now, the problem arises when the person making the attribution of knowledge and the subject asserting that *p* is the same person. Let us consider a lay person who has just been exposed to philosophy. Would she know that she has hands? Would she know when she is not considering the issue, and stop knowing it as soon as she starts considering skepticism? What if she puts philosophy aside now, then does she know it again? Her epistemic situation would be unstable.

To stabilize it, one might think that once we have been exposed to philosophy, we can no longer use ordinary standards. Philosophers—or, at any rate, people exposed to it and to skeptical worries in particular—would no longer know anything, unlike those who ignore philosophy. Perhaps it could be argued that philosophers know that they do not know. Yet, this would be small consolation. According to Socrates, in fact, knowing that you don't know is a consolation because ignorance of how things stand is common to philosophers as well as to lay people, who are untouched by philosophy. However, only philosophers have knowledge at least of one thing: their ignorance. Instead, according to contextualism, people never exposed to philosophy would know many things that philosophers would not know. Thus, philosophy would destroy knowledge instead of simply making us aware of our own ignorance, and making us wiser by giving us a piece of knowledge ordinary people typically lack.

Obviously the contextualist could bite the bullet and claim that that's just the way it is, too bad for those of us who care about skeptical hypotheses. This, however, would not be a very philosophically satisfactory response. After all, it turns out that the best antidote against skepticism would be simply to ignore it!

Another objection against contextualism (see Williamson 2001b; Wright 2005) can be raised by considering that 'know' is a factive verb. That is, from the truth of 'S knows that *p*' it follows that *p* is true. For instance, if 'I know I have hands' is true, it follows that 'I have hands' is true too. Alternatively, if I know that I have hands, it can't be false that I have hands. Now, according to contextualists, in the ordinary context (c_1) I do know *p* 'I have hands,' and I know that if I know that *p*, then *p* is true (due to the factivity of 'know'). In fact, the truth of this conditional is known in any context (c_*). Similarly, it is known in any context that in the ordinary context I know that I have hands. From this, together with the Closure principle, it follows that 'I have hands' is known in *any context*, including in the skeptical context! Contextualism would thus collapse into invariantism about knowledge. In symbols:

(1)	$K_{c_1}p \rightarrow p$	Factivity of 'to know'
(2)	$K_{c*}(K_{c_1}p \rightarrow p)$	Assumption
(3)	$K_{c_1}p$	Contextualist Assumption

(4) $K_c*(K_{c_1}p)$ Contextualist Assumption

(5) $K_c*p2, 4$ Closure

A contextualist might deny (4)—*viz.*, that in any context one would know that in the ordinary context one knows that one has hands. For if one is in the skeptical context, one cannot admit that in the ordinary context one knows that one has hands. The problem with this answer, however, is that the main attraction of contextualism consists in allowing for knowledge in the ordinary context. However, if from a skeptic's point of view that piece of knowledge cannot be guaranteed, albeit in the ordinary context, why should we embrace contextualism? Specifically, in what sense would contextualism be a response to skepticism? After all, contextualism would have to admit that a skeptic cannot grant us any knowledge, not even in the ordinary context.

We can also note the counterintuitive consequence of the contextualist strategy, when it denies the possibility to know that one has hands and that one is not a BIV at least in the philosophical context. For, after all, philosophical reflection is a high epistemic standards situation in which our reflective faculties are exercised with the utmost attention and where presumably our perceptual faculties work at their best. The question is: what else should we do to have knowledge of the external world? The contextualist leaves this mystery unsolved. Yet she is happy to grant knowledge when we presumably operate in a less rigorous way.

Another way to see how unsatisfactory contextualism is as a philosophical response to the Cartesian skeptical paradox depends on noting that for a contextualist both the Cartesian skeptic and the person untouched by it are right from their own point of view. Then, it is not clear how they can disagree with each other since each of them is right by their own standards. If the price to pay for immunity from skepticism is that it is correct based on its standards, it is unclear how contextualism can be a response to Cartesian skepticism. In fact, a skeptic might even agree that in ordinary contexts we have knowledge both of 'I have hands' and of 'I am not a BIV' and yet insist that the paradox she poses remains unsolved when the relevant standards are in place. And, in fact, according to a contextualist, the lay person should agree. Yet, if contextualism is really a way to limit the damage of skepticism by confining it to standards that are only rarely considered operational, then it is not clear how contextualism can account for the philosophical relevance of skepticism. Why, even in philosophy,

should one take seriously a problem that never really touches us? (Similar concerns are raised by Dretske 2005, 45.)

In turn, this last horn of the dilemma can develop into two separate questions. The first is: is it ever possible that skepticism has been felt by so many philosophers as one of the central problems of philosophy, if it really has so little relevance? The second is: why is a contextualist devoting time and reflection to skepticism if, considering her own doctrine, it appears irrelevant?

The contextualist might respond by saying that her account is valuable in order to dispel the mistaken impression that skepticism is a relevant philosophical issue. Once the contextual dependence of 'know' and the dynamics between contexts are clarified, skepticism is finally recognized as one of those characteristic intellectual 'bumps' that, according to the Wittgensteinian metaphor (1953, § 119), we get by running up against the limits of our language.

Nevertheless, we have already seen in Chapter 1, Section 1.2 how in the passage from antiquity to modernity skepticism was put aside with respect to common life, without for this reason being considered domesticated. For this does not detract from the fact that it remains in the form of a paradox, with characteristic premises and inferential principles that lead to a dramatic conclusion. As we have seen, the contextualist response to the skeptical paradox, while no doubt intuitively appealing, actually depends on contentious assumptions about the contextual variability of "know," and is actually impotent to address the paradox, once the relevant standards are considered.

Further reading

On the denial of the Closure principle, see Dretske (1970), Nozick (1981), Ichikawa (2011), and Coliva (2015, Chapter 3). For a criticism of Nozick's account of knowledge, see Kripke (2011). For a criticism of the denial of Closure, see Heller (1999), Vogel (1990b). On epistemic contextualism, see Cohen (1998a; 1999), DeRose (1995; 2005, 2006), Lewis (1996), and Kompa (2002). For an alternative, contrastivist account, see Schaffer (2004; 2005a; 2005b). For an assessment and criticism of epistemic contextualism, see Feldman (2001), Greco (2008), Kornblith (2000), Pritchard (2018b), Rysiew (2001), Schiffer (1996), Williamson (2005), and Wright (2005).

Notes

1 According to Nozick, Gettier (1963) cases (see Chapter 2, Section 2.3) too would flout condition (3). Suppose you form the belief that it is 3 p.m., by watching a clock. Suppose, further, that it is 3 p.m. Thus, your belief is true. However, unbeknownst to you, the clock is broken. In such an event, if it weren't 3 p.m., by watching the clock, you would still form that belief. Hence, (3) would be flouted and this is why, for Nozick, you don't know that it is 3 p.m., even if you have a true belief, supported by evidence.

2 As we will see in Chapter 7, Section 7.5, a different account of Transmission failure is possible and compatible with the failure of Closure.

4

HINGE EPISTEMOLOGY AND CLOSURE-BASED CARTESIAN SKEPTICISM

4.1 Introduction

Hinge epistemology is the name given to a family of epistemological proposals which are inspired by Wittgenstein's remarks in his final notebooks. These notebooks, published posthumously as *On Certainty* (Wittgenstein 1969,= OC), were written right at the end of his life, with the last entry just a few days before his death. They offer a sustained treatment of a range of epistemological issues related to knowledge and certainty. As we will see, part of the challenge facing hinge epistemology is to make sense of what, exactly, Wittgenstein was proposing in this regard, given that all we have to go on are the impressionistic observations found in these notebooks, which were completely unedited by him. The main upshot of these remarks is that we have a range of everyday commitments that we are optimally certain of but which are completely lacking in

DOI: 10.4324/9780429057946-5

rational support. Nonetheless, it is entirely legitimate for us to have these commitments. If that last idea can be made good, then it clearly has anti-skeptical potential. In particular, as we will see, there are ways of unpacking this idea that can make it relevant to both Cartesian and Humean skepticism (and possibly Agrippan skepticism as well).

Being such a family of different views, here we will consider only what one main variety of hinge epistemology, due to Duncan Pritchard (2015), has to say with respect to the Closure-based version of the Cartesian skeptical paradox. Other versions of hinge epistemology that are more directly connected to the Humean skeptical paradox will be considered in Chapter 6, Section 6.4 and Chapter 7.

4.2 Hinge epistemology in outline

Wittgenstein's focus in *On Certainty* is on the everyday certainties that G. E. Moore (e.g., 1925; 1939) identified. These Moorean certainties concern mundane quotidian claims about which, Moore argued, we are optimally certain. This includes such propositions as 'I have two hands,' 'I have never been to the moon,' 'I am speaking English,' and so on. Moore was interested in these everyday certainties because he was trying to articulate a commonsense approach to philosophical questions, whereby one can appeal to commonsense to push back against revisionist philosophical theorizing. On this way of thinking about the Moorean certainties, they possess a kind of special epistemic status in virtue of being commonsense certainties. Wittgenstein agreed that these Moorean certainties were epistemologically significant, but instead drew a radically different conclusion—*viz.*, that their optimal certainty means that we can have no rational basis for their truth at all.

Wittgenstein's idea is that these Moorean certainties play a 'hinge' role in our rational practices, such that by holding fast, they enable rational evaluations to occur. A consequence of this picture, however, is that they are not themselves subject to rational evaluation; indeed, he claims that the very idea of rationally evaluating a hinge commitment is incoherent for just this reason. This means that, on the Wittgensteinian picture, fully generally rational evaluations—i.e., rational evaluations of all our commitments—are simply impossible. This calls into question not just the radical skeptical project of rationally evaluating all our commitments at once and finding them wanting, but also the traditional *anti*-skeptical project of rationally

evaluating all our commitments and finding them in good epistemic order. Both radical skepticism and conventional forms of response to radical skepticism are predicated upon the same fundamental error.

Here is the famous passage where Wittgenstein introduces the hinge metaphor:

> [T]he *questions* that we raise and our *doubts* depend upon the fact that some propositions are exempt from doubt, are as it were like hinges on which those turn.
>
> That is to say, it belongs to the logic of our scientific investigations that certain things are *in deed* not doubted.
>
> But it isn't that the situation is like this: We just *can't* investigate everything, and for that reason we are forced to rest content with assumption. If I want the door to turn, the hinges must stay put.
>
> (OC, §§341–343)

Note that Wittgenstein's claim in this regard is that the hinge commitments are an essential part of our practices of rational evaluation, in that we can make no sense of a system of rational evaluation that is lacking them. This is why they are very different from assumptions, which can be discharged. In particular, it is not as if by being more rigorous or imaginative, or generally smarter, that we could find a way to dispense with our hinge commitments. They are rather nonoptional features of the architecture of rational evaluation.

That there are some fundamental epistemic commitments that we are obliged to have is not in itself radical. What is radical is rather the idea that these commitments are both (1) entirely mundane and everyday propositions, and (2) lacking in any rational status. On (1), notice that the hinge commitments are very different from the kinds of putatively foundational claims that we are familiar with in epistemology, which are usually philosophical theses like the *cogito* rather than quotidian commitments. Relatedly, on (2), the putative foundational claims that we tend to find in epistemology are also meant to have a special kind of rational status, in virtue of their epistemic properties (incorrigibility, self-evidence, being infallibly known, etc.,), but the hinge commitments are held to lack any rational status at all.

In keeping with Wittgenstein's general philosophical strategy in his later writings, we are not offered an argument for the existence of hinge

commitments, at least not in any conventional sense at any rate. We are instead presented with cases and observations on those cases, and invited to draw the relevant conclusions for ourselves. For example, it might be tempting to suppose that one's certainty, in normal conditions, that one has hands is grounded in one's sight of one's hands. Wittgenstein tries to get us to recognize, however, that we can make no sense of the idea of the optimal certainty that we have in a hinge commitment such as our commitment to having hands as being grounded in anything:

> My having two hands is, in normal circumstances, as certain as anything that I could produce in evidence for it.
>
> That is why I am not in a position to take the sight of my hand as evidence for it.
>
> (OC, §250)

The point isn't that this certainty outweighs any possible rational basis one could offer for it, but rather that to understand it as rationally grounded at all is to misunderstand its role:

> If a blind man were to ask me "Have you got two hands?" I should not make sure by looking. If I were to have any doubt of it, then I don't know why I should trust my eyes. For why shouldn't I test my *eyes* by looking to find out whether I see my two hands? *What* is to be tested by *what*?
>
> (OC, §125)

The hinge commitments instead constitute the backdrop of certainty against which rational evaluations—the 'testing,' to follow through on the analogy in this passage—occurs. Compare, in this regard, someone asking you whether you have your car keys. Now it would make perfect sense to look for them, e.g., to search one's pockets. This is thus a normal belief, open to rational evaluation in the usual ways. The contrast between these normal beliefs that we hold and the hinge commitments is meant to get us to see how the latter are playing a very different role in our epistemic practices.

This is an intriguing proposal, albeit one which raises more questions than it answers. If the hinge commitments are lacking in rational support, then does it follow that they are unknown? And if they are unknown, then how is hinge epistemology meant to be a response to radical skepticism? After all, isn't it part of radical skepticism to contend that at root our

system of rational evaluation boils down to arational commitments at its core? In particular, why doesn't it follow from the arational nature of our hinge commitments that all our beliefs are lacking in positive rational standing? A related question in this regard concerns the status of the Closure principle. If our hinge commitments are unknown, then how does that square with our other beliefs amounting to knowledge, particularly when the latter is known to entail the former? For if Closure holds, then it seems that allowing these beliefs to amount to knowledge must mean that there is an inferential route to knowing the hinge commitments too. Alternatively, is the proponent of hinge epistemology obliged to reject the Closure principle? In the midst of all this is also the further exegetical question—possibly unanswerable, given the limited textual evidence that we have to go on—of what Wittgenstein specifically had in mind in this regard (or would have had in mind, had he gone on to develop this idea more fully).

The issue of the status of the Closure principle relates to a more general issue concerning the application of hinge epistemology to the radical skeptical problem. This is because Wittgenstein's remarks clearly suggest an undercutting response to this difficulty. Recall from Chapter 1 that undercutting responses maintain that the putative radical skeptical paradox is in fact not a paradox at all. In particular, rather than there being a deep tension in our epistemological concepts, radical skepticism is instead the result of faulty philosophical theorizing. Wittgenstein's description of how our epistemic practices have hinge commitments at their heart is meant to show that we have been working with a defective picture of how our epistemic practices function, and which is thereby generating the skeptical philosophical puzzle. This is in keeping with the broad approach to philosophical problems that we find in Wittgenstein's later work—see especially Wittgenstein (1953)—where such problems arise because of the philosophical misuse of everyday language. The problem, however, is that if hinge epistemology leads to the denial of a principle like Closure—or, for that matter, a related principle like underdetermination—then it is hard to make this undercutting claim stick. After all, these principles do seem to be genuinely rooted in our commonsense conception of the epistemic landscape, and hence denying them in order to respond to radical skepticism makes that anti-skeptical strategy look revisionary, and hence overriding rather than undercutting.

There are numerous attempts to unpack hinge epistemology, with each unpacking offering a distinct set of answers to questions like this. P.F. Strawson (1985) (see Chapter 7, Section 7.2) argues that the salient feature of hinge commitments is the fact that we are unable to doubt them, and that this 'naturalistic' fact about our human condition provides a kind of insulation against radical skeptical doubt. Danièle Moyal-Sharrock (2005) contends that it is a mistake to think of our hinge commitments as propositional commitments at all, maintaining instead that we should treat them along enactivist lines as fundamentally being ways of acting (see Chapter 7, Section 7.3). Other commentators, like Crispin Wright (2004) and Annalisa Coliva (2015), see Wittgenstein as offering a way of accounting for how our hinge commitments might enjoy positive epistemic support, even despite their lack of rational support, at least to the extent that they play an essential role in the framework of epistemic rationality. And Michael Williams (1991) (see Chapter 7, Section 7.4) regards Wittgenstein as developing a kind of epistemic contextualism, although of a very different kind to that we encountered in Chapter 3, in that rather than presenting a context-sensitive account of 'knows,' it instead offers a contextualized version of foundationalism, with the epistemic foundations shifting from context to context. We will return to consider some of these proposals in their own right later on, but for now we will focus on one specific interpretation, due to Pritchard (2015). As we will see, this offers a reading of hinge epistemology that embraces the claim that our hinge commitments are not in the market for knowledge, while nonetheless contending that this is compatible with also endorsing the Closure principle. Key to this interpretation is a certain understanding of what our hinge commitments involve.

4.3 Hinge commitments and radical skepticism

Wittgenstein puts a great deal of emphasis on the nature of our hinge commitments. We have already seen the point that they are non-optional features of the epistemic architecture, and hence unlike assumptions that can be discharged. Wittgenstein goes further, however, and contends that our hinge commitments, and the certainty that is associated with them, are not just arational but also have a visceral quality: they are "primitive" and "animal" (OC, §359). In particular, the certainty that is characteristic of our hinge commitments is imbued in all our actions and as such is entirely

prior to ratiocination, making it immune to justification: "it is there, like our life" (OC, §559). This is also reflected in the manner of our relationship to them. Wittgenstein contends that our hinge commitments are impervious to genuine doubt, in that even if we become aware of their special arational standing, this doesn't undermine the certainty we have in them. In particular, even if we might claim to doubt the hinges, our certainty in them will nonetheless be manifest in our actions. (The 'doubt' that philosophers profess of their hinge commitments is thus claimed to be in a certain sense fake.)

In the normal run of events, we are not even aware of our hinge commitments as they "lie apart from the route travelled by enquiry" (ibid., §88). This is the manner in which our hinge commitments are hidden from us, at least *qua* hinge commitments, even though they are in plain view. Relatedly, we are not explicitly taught our hinge commitments, but rather "swallow them down" (ibid., §143) along with other things that we are explicitly taught. For example, we are not taught that we have hands, but rather to do things with our hands. Indeed, it seems clear that for Wittgenstein the specific hinge commitments that one might have in any given circumstance are essentially just nodes in a more general backdrop of primitive certainty that he believes provides the necessary framework for a system of rational evaluation. It is not that we gradually become certain of things as we are enculturated into the space of reasons, but rather that we must already have this certainty in place in order to be, and remain, citizens in the space of reasons at all.

The foregoing suggests a particular view of the nature of our hinge commitments. On the face of it, the specific hinge commitments listed above seem like a heterogeneous bunch, in that they are variable in terms of such factors as geography ('I speak English') and epoch ('I have never been to the moon'). Many accounts of hinge commitments emphasize their heterogeneous character (e.g., Williams 1991 and Moyal-Sharrock 2005). But the preceding discussion indicates that there may be instead a common nature to all our specific hinge commitments, which is the overarching primitive certainty just described. In essence, this is the certainty that one is not radically in error; that one's relationship to the world around one, and the people in it, is broadly veridical. Call this the *über hinge commitment*. It is this general commitment that Wittgenstein maintains is presupposed by the space of reasons. Of course, the claim is not that one

must have an occurrent thought with this particular content. Rather the idea is that one's actions are infused with this overarching certainty.

With this idea of the über hinge commitment in play we can account for why we might have hinge commitments with a more specific content in particular circumstances, and also how these narrower hinge commitments could change over time. For example, in normal circumstances our general certainty that we are not radically in error, encapsulated in the über hinge commitment, will be manifest in a commitment to the specific proposition that one has hands. For if one lacked hands in normal circumstances, then that would entail the kind of radical error that is at issue in the über hinge commitment. The same goes for the other hinge commitments that we have noted, such as that one has never been to the moon or that one is speaking English. In changed conditions, however, one will manifest hinge commitments with different contents. For example, if one is awakening in hospital after a serious accident, it may no longer be a hinge commitment that one has two hands; indeed, it may make perfect sense in such conditions to form a belief that one has hands by simply looking for them, just as one might look for one's car keys (unlike in normal circumstances, where looking for one's hands in this way would be bizarre and indeed a sign of neural impairment or mental disorder). Relatedly, which specific hinge commitments reflect the über hinge commitment may change over time. It was a hinge commitment for Moore and Wittgenstein that they had never been to the moon, and surely it is for us too several generations later, but clearly there could be a future epoch where going to the moon is the sort of thing that could easily occur without one being aware of it (just as someone who grew up in England might have been to the city of Norwich without having realized this).

Thinking of the specific hinge commitments as manifestations of this more general über hinge commitment also makes it easier to explain why our hinge commitments have the properties that they do. That the über hinge commitment needs to be in place in order for rational evaluation to occur, and hence that it cannot be itself rationally evaluated, looks very plausible, given the global nature of the commitment involved, as any attempt at rationally evaluating this commitment would be bound to be epistemically circular. It should not be surprising, then, that specific hinge commitments inherit this property insofar as they are simply manifestations of the über hinge commitment. In addition, we can also account for why our hinge commitments extend beyond the everyday Moorean

certainties to encompass the denials of radical skeptical hypotheses, given that the latter bear an obvious connection to the über hinge commitment.

A further advantage of this way of thinking about hinge commitments concerns the issue of epistemic relativism. As the name suggests, epistemic relativism specifically concerns a relativism about epistemic systems. In particular, the idea is that there can be distinct epistemic systems, incorporating different epistemic standards, where each epistemic system is as legitimate as any other. What is especially troubling about epistemic relativism is when the difference in epistemic systems is so radical that epistemic incommensurability results, whereby there is no rational way of resolving disputes between parties employing distinct epistemic systems.

This was certainly a concern that Wittgenstein was alert to, as it seems to be one consequence of embracing hinge commitments. After all, if our hinge commitments need not be shared, then what is to prevent there from being radical divergences between groups of people in terms of their hinge commitments? Moreover, given that hinge commitments effectively provide the framework for rational evaluation, then this seems to suggest that epistemic incommensurability will inevitably result, as each party will have a very different epistemic system in virtue of their very different set of hinge commitments.

Some commentators—such as Williams (1991) and Martin Kusch (2016b)—embrace such epistemic relativism as being a consequence of hinge epistemology. But before we also embrace it, we should recognize that whether this really is a consequence of the view very much depends on how one understands hinge commitments. It is only because Williams and Kusch endorse an account of hinge commitments that emphasizes their heterogeneity that this result follows. In contrast, as explained in Pritchard (2018c), if one adopts a very different account of hinge commitments that views them as manifestations of the über hinge commitment, then the scope for radical divergence in our hinge commitments is very much lessened. For example, while it might be a hinge commitment of mine that I am speaking English, and a hinge commitment of someone on the other side of the world that she is speaking Japanese, we are both manifesting the über hinge commitment by being hinge committed to our speaking our own language. The differences in our specific hinge commitments on this view are often more superficial than substantive.

A further important feature of this account of hinge commitments concerns the kind of propositional attitude that they involve. In particular, they are not plausible candidates for being beliefs, at least as that notion is employed by epistemologists anyway. The everyday notion of belief is notoriously broad and encompassing, capturing just about any propositional attitude of general endorsement of a proposition. Unsurprisingly, then, we find that the notion gets used in theoretical contexts (both within and beyond philosophy) in a range of different ways. As epistemologists, however, our concern is less with the loose folk notion of belief and more with a subspecies of this notion, *viz.*, that specific propositional attitude which is a constituent part of rationally grounded knowledge. Call this *K-apt belief*. K-apt belief has some interesting properties, not least in terms of its close conceptual connections with reasons and truth (conceptual connections that would be hard to sustain on the loose, folk notion of belief). Here is one such conceptual connection that concerns us: K-apt belief that *p* is such that it cannot coexist with the subject recognizing that she has no rational basis for the truth of *p*. The subject might still believe that *p* in the everyday sense, or have some other propositional attitude towards *p* (such as hoping that *p*, or wishfully thinking that *p*), but this would not be a K-apt belief that *p*. This point is important because the certainty that *p* involved in a hinge commitment that *p* is precisely such that it would, and does, survive the recognition that one has no rational basis for the truth of *p*. The upshot is that while hinge commitments might be beliefs, broadly speaking, they are not beliefs in the specific K-apt sense that is relevant for epistemology.

The significance of this claim goes beyond merely capturing the propositional attitude at issue in our hinge commitments, as it is also relevant to how one might apply hinge epistemology to the skeptical puzzle. As noted above, hinge epistemology entails that fully general rational evaluations—i.e., rational evaluations of all one's commitments—are simply incoherent. Indeed, given the necessity of hinge commitments to any system of rational evaluation, to aspire for such a fully general rational evaluation, of either a positive (anti-skeptical) or negative (skeptical) kind, is akin to aspiring to draw a circle-square. This is the sense in which our hinge commitments are neither known nor unknown, since by being outwith the realm of rational evaluation they are simply not in the market for knowledge. It is not as if we are ignorant of them, for example (as if we could have known them, if only we had been cleverer, undertaken different inquiries, and so on).

The anti-skeptical import of this point isn't immediately clear, however. As noted above, it seems to suggest that the Closure principle should be rejected, in that it implies that our everyday knowledge can coexist with our inability to know the hinge commitments, even where the former clearly entails the latter (as with our hinge commitment to the denials of radical skeptical hypotheses). The problem, of course, is that the Closure principle looks to be eminently plausible. Accordingly, it seems to be a count against hinge epistemology if it has this consequence, at least insofar as we think of it, as surely we ought to as a Wittgensteinian proposal, as an undercutting anti-skeptical strategy (though some versions of hinge epistemology are happy to embrace this consequence nonetheless, as we will see in later chapters).

On closer inspection, however, the tension between the Closure principle and hinge epistemology disappears, at least if we adopt the conception of hinge commitments just described. Recall that we pointed out in Chapter 1 that it was important that Closure was formulated diachronically as a competent deduction principle, as follows:

> *Closure principle (diachronic)*
> If one knows that *p*, and one competently deduces *q* from *p*, thereby coming to believe that *q* on this basis while retaining one's knowledge that *p* throughout, then one knows that *q*.

It is only if Closure is formulated in this fashion that it doesn't succumb to straightforward counterexamples, which would deprive it of any legitimacy in the skeptical argument. So construed, however, notice that the principle is explicitly concerned with the acquisition of a belief via a paradigmatically rational process. Moreover, it is clearly belief in the K-apt sense that is at issue, given that it is meant to be belief that is a component part of knowledge acquired via competent deduction (and thus rationally grounded knowledge). This means that one simply cannot plug hinge commitments into Closure-style inferences as the radical skeptical supposes and thereby generate the required paradoxical results, as they are not the kind of propositional attitude that can be acquired in this manner. In particular, one cannot form a K-apt belief in one's hinge commitments at all, much less via a paradigmatically rational process like competent deduction. It is thus not the Closure principle that needs to be denied, but rather that we need to appreciate how this principle cannot be employed to

generate the skeptical puzzle. The result is that the Closure-based formulation of the radical skeptical paradox fails because there is no inconsistent triad in play. Once we embrace the account of hinge commitments described above, then it can be true that one has everyday knowledge, true that one lacks knowledge of the denials of radical skeptical hypotheses (*qua* hinge commitments), and also true that the Closure principle holds.

How might hinge epistemology bear on the other formulations of the radical skeptical problem? There is an ongoing debate in the literature about whether it can help us to resolve underdetermination-based radical skepticism. Pritchard (2015) argues that hinge epistemology is silent with regard this formulation of the skeptical problem, since it is merely concerned with undermining the idea of universal rational evaluations, a point that is specifically relevant to Closure-based skepticism. For underdetermination-based skepticism, in contrast, what is at issue is whether we can have the favoring rational support for our beliefs that is at issue, and nothing in hinge epistemology speaks to that issue. (He contends that what is required in this regard is the alliance of hinge epistemology and epistemological disjunctivism as part of a 'biscopic' treatment of radical skepticism—we will consider this point in Chapter 5.) In contrast, other commentators, such as Genia Schönbaumsfeld (2016), have claimed that a proper development of hinge epistemology will help us dissolve both Closure-based and underdetermination-based radical skepticism, in the latter case by explaining how we might have the required favoring rational support.

Exactly how hinge epistemology responds to Humean skepticism is moot, as it depends on the version of hinge epistemology at issue, and in particular on whether the hinge commitments are known (and, if so, in virtue of what that they are known). What will be common to all such proposals is the idea that there cannot be any epistemic circularity involved in how one comes to know one's hinge commitments either because (1) they aren't in the market for knowledge, or (2) they are known, but not as a result of direct rational support (much less rational support gained from an inference from other non-hinge beliefs that one holds). The first approach is, of course, the one that we have explored in this chapter; we will be considering proposals that take the second approach in later chapters.

We have already seen how our lack of knowledge of hinge commitments can be compatible with Closure-style inferences, in that such inferences cannot lead to K-apt belief in hinge commitments. A similar point applies to the sorts of inferences at issue in Humean skepticism. If one's

hinge commitments cannot be the result of an inference from prior beliefs, then it follows that they cannot be subject to the kind of epistemic circularity highlighted by Humean skepticism. Moreover, while it is true that our everyday knowledge presupposes the truth of our hinge commitments, this is held to be a necessary feature of what it is to have rationally grounded knowledge, rather than revealing an epistemic lack on our parts.

Hinge epistemology might also have application to Agrippa's trilemma, though this is not a point that has been explicitly discussed in any great detail in the contemporary literature. In particular, hinge epistemology seems to offer a way of making sense of a kind of foundationalist response to this trilemma, albeit of a kind that is very different from traditional foundationalism. While traditional foundationalism is concerned to show that there can be legitimate stopping-points in the regress of reasons on account of how some beliefs can be self-justifying, hinge epistemology takes a very different tack. It offers, instead, a proposal according to which we embrace the fact that the regress of reasons comes to an end, but not with claims that are self-justifying, but rather in claims that lack any rational standing at all. What is key to this proposal—and which is meant to ensure that it doesn't merely entail a surrender to Agrippan skepticism—is the contention that it is in the nature of all rational evaluation that there be groundless commitments of this kind. Accordingly, what is being presupposed in Agrippan skepticism is a conception of the nature of rational evaluation which, it has been argued, is erroneous. In particular, that a local regress of reasons ends with a hinge commitment that lacks any rational support does not show that one's beliefs are thereby ungrounded, but rather reveals the essentially local nature of any system of rational evaluation.

Epistemic vertigo

A final point that is relevant to our discussion of hinge epistemology and radical skepticism concerns the notion of *epistemic vertigo* (Pritchard 2015, Part 4; 2020). Epistemic vertigo relates to the phenomenology of our engagement with radical skepticism, and in particular the fact that even when one is in possession of a plausible anti-skeptical response, there is still a tendency to feel a kind of anxiety about one's epistemic position. What could be the source of this anxiety, especially if one is not actually doubtful, on skeptical grounds at least, about the security of one's knowledge?

Hinge epistemology seems uniquely placed to explain what is going on here. The source of the anxiety cannot be a genuine concern for one's epistemic position, as on this view one is optimally certain that one is not the victim of radical skeptical scenarios. Nonetheless, reflecting on the radical skeptical problem does lead to a kind of philosophical detachment from one's ordinary, engaged, existence, and it is this that is generating the anxiety. In particular, recall that we noted earlier how our hinge commitments are essentially hidden from us in our everyday lives, albeit in plain view, in that we simply are not aware of them at all in the normal run of events. They are rather essentially tacit: implicit in our actions as opposed to being explicit in our thoughts. It takes a special kind of inquiry, of a broadly philosophical nature, to become aware of them, and in particular to become aware of them *qua* hinge commitments. This awareness, however, inevitably provokes a kind of detachment from one's everyday practices, since it involves taking a perspective on those practices which is external to them. It is in this sense that our awareness of our hinge commitments brings with it an air of the uncanny, in that their very ordinariness belies the extraordinary role that they play in our system of rational evaluation.

Being aware of one's hinge commitments, even in the context of an anti-skeptical understanding of why they are essential to our rational evaluations, thus tends to provoke anxiety due to how it involves an unnatural perspective on one's practices. The terminology employed to describe this—epistemic vertigo—is meant to capture this perspective. In particular, it depicts the idea that in, as it were, epistemically ascending—as one adopts a detached, philosophical perspective—one becomes aware of one's hinge commitments and the special role that they play in our epistemic practices, and that this elicits an essentially phobic reaction, even if one at the same time recognizes that there is no genuine epistemic risk in play. The analogy here is with the phobic reaction that can be generated from being high-up (the pedantic can substitute 'acrophobia' for 'vertigo' if they wish), a phobic reaction that can perfectly well coexist with the recognition that one isn't in any danger. (Epistemic vertigo, like acrophobia, thus involves an *alief*, in the sense articulated by Tamar Gendler (2008).)

Understanding the phenomenon of epistemic vertigo and its relationship to hinge epistemology is important because it highlights an important feature of the undercutting response to radical skepticism that Wittgenstein is offering. One might naturally suppose that undercutting responses to

philosophical paradoxes inevitably return us back to a state of philosophical innocence, prior to our engagement with the paradox in question, whereby we are able to partake in our everyday practices as before, without any philosophical concern arising. Interestingly, this does not seem to be what a Wittgensteinian approach to radical skepticism provides (arguably, it is not what the later Wittgenstein offers as regards resolving *any* philosophical paradox). In particular, once we are aware of our hinge commitments, there isn't any straightforward way that we can become unaware of them. Some degree of detachment from our everyday epistemic practices is thus sure to remain even once the radical skeptical paradox has been undercut, as we remain struck by the uncanny nature of our hinge commitments. Hinge epistemology may help us resolve skeptical puzzles, but it doesn't return us back to a state of epistemic innocence prior to our engagement with those puzzles.

Further reading

For further discussion of the relationship between Wittgenstein's (1969) hinge epistemology and Moore's (e.g., 1925; 1939) commonsense philosophy, see Coliva (2010a) and Pritchard (2021a). See also Chapter 6, Section 6.2 and Chapter 7, Section 7.3. For some of the main interpretations of hinge epistemology, see Strawson (1985), McGinn (1989), Williams (1991), Moyal-Sharrock (2005), Wright (2004), Coliva (2015), Pritchard (2015), and Schönbaumsfeld (2016). See also Chapter 6, Section 6.4 and Chapter 7. For some key works that explore the issue of epistemic relativism and its relationship to hinge epistemology, see Kusch (2016b), Pritchard (2018c), and Coliva (2020a). For a useful taxonomy of some different notions of belief in use, see Stevenson (2002). See also Gendler (2008) on alief. For a presentation of hinge epistemology that focuses on the specifically non-standard way that it advances a kind of foundationalism, and which thus might be useful for thinking about the potential application of hinge epistemology to Agrippa's trilemma, see Moyal-Sharrock (forthcoming). For further discussion of the notion of epistemic vertigo, see Pritchard (2015, Part 4; 2020). See also Cavell (1979, *passim*) for a discussion of related ideas in the context of the work of the later Wittgenstein. For more on the later Wittgenstein's brand of philosophical quietism, see McDowell (2009). For an overview of recent work on hinge epistemology, see Pritchard (2017b).

5

EPISTEMOLOGICAL DISJUNCTIVISM

5.1 Introduction

The anti-skeptical proposal that we will consider in this chapter—*epistemological disjunctivism*—can be applied to both the Humean and the Cartesian skeptical paradoxes. According to epistemological disjunctivism, at least some of our everyday perceptual beliefs enjoy factive rational support, in the sense that this rational support actually entails the truth of the target proposition. If that's right, then this has the potential to block the charge of epistemic circularity at issue in Humean skepticism, since such factive rational support should suffice for knowledge of the target proposition, regardless of any further ancillary claims obtaining, including anti-skeptical theses. The factivity of the rational support for our perceptual beliefs should also block the Cartesian skeptical argument, especially the underdetermination-based

DOI: 10.4324/9780429057946-6

formulation of this puzzle, which (as will see) essentially turns on a conception of rational support that is non-factive.

5.2 Epistemological disjunctivism in outline

Epistemological disjunctivism is rooted in the work of John McDowell (e.g., 1995), though this is not the terminology that he used for this position. We will focus on the contemporary presentation of this proposal, found in Pritchard (e.g., 2012; 2015, Part 3). This states that in paradigm cases of perceptual knowledge, the subject's belief can enjoy rational support that is both reflectively accessible and factive. In particular, the reflectively accessible rational support for the subject's belief that *p* can be that she *sees that p*, where seeing that *p* entails *p*.

Let's unpack this proposal a little. To begin with, notice that epistemological disjunctivism is contending that paradigmatic perceptual knowledge enjoys internalist epistemic support, as the rational support in question is reflectively accessible. This doesn't mean that epistemological disjunctivism is committed to internalism about knowledge in general—in fact, it is consistent with a more general externalism about knowledge—but it does mean that paradigmatic perceptual knowledge is of an internalist kind.

Nonetheless, this is not a classical form of epistemic internalism. As we saw in our discussion of the Gettier-style cases and the challenge they pose for classical internalism about knowledge in Chapter 2, the reflectively accessible rational support at issue in classical views is non-factive. (Indeed, this factor is crucial to the set-up of the Gettier-style cases, which suggests that epistemological disjunctivism might have the resources to resist these cases, at least where they concern perceptual belief in epistemically paradigm conditions.) Since epistemological disjunctivism maintains that there can be factive reflectively accessible rational support, we thus have a non-classical form of internalism about (a certain kind of) knowledge in play. Indeed, epistemic internalism of this variety is usually thought to be simply unavailable. Either the epistemic support that one's belief enjoys entails substantive facts about the world (e.g., that the belief was reliably formed), in which case it is externalist (and thus non-reflectively accessible) epistemic support; or it is reflectively accessible (and hence meets the epistemic internalist rubric) but is thereby such that it doesn't entail substantive facts about the world (much less entailing the proposition believed). Epistemological disjunctivism thus effectively offers

a kind of middle way between internalism and externalism about knowledge, by offering a conception of epistemic support that is both reflectively accessible (and thus epistemic internalist) and which entails substantive facts about the world (in keeping with epistemic externalism).

Why might internalist epistemic support be thought to be essentially non-factive? Interestingly, the reason lies with scenarios like radical skeptical hypotheses—i.e., scenarios which are indistinguishable from normal experience, but where one is in error. Take the brain-in-a-vat (BIV) skeptical scenario, for example, and compare it to a corresponding scenario where one is enjoying normal experiences, indistinguishable from the deceptive experiences generated in the vat. Call the former a paradigm kind of 'bad' case, with the latter a corresponding 'good' case. Aren't the reasons one has for one's beliefs in this bad case the very same reasons that one would offer if one were having the corresponding 'authentic' experiences in the good case? At the very least, how could the rational support one's belief enjoys in the good case be any better than the rational support it enjoys in the bad case, if one cannot tell them apart? This is what is known as the *new evil genius* intuition. This intuition entails that the rational support one's belief enjoys, even in the good case, cannot be factive. This is because the rational support one's belief enjoys in the bad case clearly cannot be factive, as one's beliefs in the BIV, for example, are mostly false. In particular, in this bad case, the rational support cannot consist of seeing that *p*, since one doesn't see that anything is the case in the vat, but merely *seems to see that p*, where seeming to see is obviously non-factive (one can seem to see that something is the case when it isn't the case). Accordingly, if the rational support one's belief enjoys in the good case can be no better than the rational support available in the corresponding bad case, then it follows that it cannot be a factive reason like seeing that *p* either, but only a non-factive reason like seeming to see that *p*.

Epistemological disjunctivism resists this line of reasoning, and thus rejects the new evil genius intuition. It insists that we shouldn't restrict the strength of the subject's rational standing in the good case to the subject's corresponding rational standing in the bad case. Instead, it argues that the rational standing enjoyed by one's belief in the good case is fundamentally different in kind from the rational standing enjoyed by one's belief in the corresponding bad case. This is the *disjunctivist* aspect of the view. Rather than there being a rational standing that is common to both the good and the bad case, and hence a non-factive rational standing, *either* one is in the good case and one's belief enjoys a factive rational standing *or* one is in the

corresponding bad case and one's belief enjoys (at most) a non-factive rational standing. Moreover, there is no common core to the rational standings in pairs of good and the bad cases either. It is not as if the rational standing one's belief enjoys in the good case includes the rational support it enjoys in the corresponding bad case, albeit supplemented by additional rational support. Instead, the two forms of rational support are completely distinct.

Epistemological disjunctivism here mirrors the metaphysical disjunctivism that has been widely advanced with regard to perceptual experience. According to metaphysical disjunctivism, we should resist the idea that the indistinguishable perceptual experiences in pairs of good and bad cases thereby enjoy a common metaphysical core. In particular, one might be tempted to hold that since one's experiences in the bad case (e.g., the radical skeptical case) don't involve a veridical grip on the nature of reality, then it follows that neither do one's experiences in the good case either, given that they are indistinguishable. Metaphysical disjunctivism rejects this line of reasoning, however, and maintains that the nature of one's perceptual experience can be fundamentally different in the good case as opposed to the corresponding bad case. In particular, metaphysical disjunctivism tends to go hand-in-hand with a kind of realism about perceptual experience such that in the good cases one's experience is in some sense directly of an external world. *Either* one is in the good case, and enjoying a veridical perceptual experience of the world around one, *or* one is in a corresponding bad case and having a perceptual experience of a very different kind. Moreover, just as epistemological disjunctivism rejects the idea of there being a rational common core across good and bad cases, so metaphysical disjunctivism rejects the idea of there being a metaphysical common core across pairs of good and bad cases (at least beyond the merely negative epistemological fact that they are indistinguishable). In particular, the nature of the perceptual experience in the good case is not merely the perceptual experience in the corresponding bad case combined with some further elements; rather, the two experiences are different in kind.

It is an interesting question how metaphysical and epistemological disjunctivism relate to one another. They are certainly dialectically affiliated, in that being committed to one of these theses almost certainly leads to a commitment to the other. In particular, it is natural to explain how one might have factive reasons in the good case by appealing to the metaphysics of perceptual experience along metaphysical disjunctivist

lines. Nonetheless, we can keep the two positions apart. Moreover, we are wise to do so as they are dealing with very different philosophical issues. One way of bringing this to the fore is to consider how they each confront a different concern about the 'veil of perception.' Metaphysical disjunctivism is concerned to resist the metaphysical veil of perception which falls when we grant that the nature of our perceptual experience is such that we can gain no more purchase on the world external to us directly through perception than we do when we are radically deceived about that world (and hence where our perceptual experience isn't properly hooked up to world around us at all). In contrast, epistemological disjunctivism is concerned with the epistemological veil of perception (or, if you prefer, an epistemological consequence of the metaphysical veil of perception). This is such that our rational support for our perceptual beliefs can be no better than the rational support our perceptual beliefs enjoy when we are radically deceived (such that the rational support our perceptual beliefs enjoy is always compatible with those beliefs being radically false).

In any case, our concern is with epistemological disjunctivism specifically. The defense of epistemological disjunctivism essentially comprises three parts (Pritchard 2012, Part 1). First, that epistemological disjunctivism occupies a 'default' status as an account of the rational support enjoyed by our perceptual knowledge in virtue of it being rooted in our everyday epistemological practices. Second, that the reasons why we do not take our epistemological practices at face value in this regard, and hence reject epistemological disjunctivism, are the product of faulty philosophical theorizing. Epistemological disjunctivism is thus presenting an undercutting response to radical skepticism. Third, that epistemological disjunctivism is a position that we would want to endorse, were it to be available, on account of its theoretical benefits, not least in terms of its relevance to the problem of radical skepticism.

We will examine the second point in Section 5.3, but let us remark briefly on the first and third points. The first claim looks plausible, and is often not in dispute in the debate about the status of epistemological disjunctivism (since the main complaint offered by its critics is that the view is indefensible, and hence that we shouldn't take our everyday epistemic practices at face value in this regard). In particular, it does seem natural to offer factive reasons in support of our perceptual knowledge claims in normal conditions. If you dispute that I have perceptual knowledge of something right before me (perhaps because we are speaking on the phone, and you find what I am

claiming independently implausible), then it would be natural to defend my claim by saying that I can see that this is the case. Indeed, it usually takes rather unusual circumstances to prompt us to 'hedge' the rational support we have for our perceptual knowledge by appealing to a non-factive reason. In particular, it would be odd, in normal conditions where there is no special reason for doubt, to defend one's perceptual knowledge by stating that it merely seems to one as if this is the case (in fact, to do so would usually be thought to suggest that one was at least qualifying one's support for the original knowledge claim, if not withdrawing it altogether).

The relevance of epistemological disjunctivism to radical skepticism—and hence the appeal of the view, were it to be available—should be evident. The Humean skeptical concern about epistemic circularity presupposes that the rational support we have for our everyday beliefs must be non-factive, in that it is held to be rational support that only plays a knowledge-supporting role on the supposition that certain anti-skeptical claims (such as the denials of radical skeptical hypotheses) are true. If the rational support enjoyed by these beliefs is factive, however, then that changes things completely, as it would entail that one's perceptual knowledge is not hostage to the truth of these ancillary claims at all. Moreover, if one now infers these anti-skeptical claims from one's perceptual knowledge, it is unclear why this would be epistemically circular, given that the prior knowledge that one is employing in the inference doesn't presuppose the entailed claim. The charge of epistemic circularity at the heart of the Humean skeptical problematic thus seems to be neutralized.

Moreover, one can also apply epistemological disjunctivism to the Cartesian skeptical problem too. This is especially evident with regard to the underdetermination-based formulation of Cartesian skepticism. Recall that this turns on the Underdetermination principle, which we characterized as follows:

> *Underdetermination principle*
> If one knows that *p* and *q* are incompatible propositions, and one knows that *p*, then one has an epistemic basis for believing that *p* that favors *p* over *q*.

The principle looks very plausible, as we noted in Chapter 1. For if one really does lack an epistemic basis that favors one's belief over knowledge to be incompatible alternatives, then how can that belief amount to

knowledge? The skeptical problem arises, however, from the fact that radical skeptical scenarios seem precisely to be known alternatives to our everyday beliefs where we lack the required favoring epistemic basis. Accordingly, it follows that we lack knowledge of these everyday beliefs.

Epistemological disjunctivism offers us a direct line of response to underdetermination-based radical skepticism. This is because it argues that our everyday perceptual beliefs can enjoy a factive rational basis. If that's right, then of course that rational basis will favor those beliefs over all known alternatives, including radical skeptical ones. The skeptical argument is thereby blocked.

If one has a rational basis that favors one's everyday beliefs over radical skeptical hypotheses, then that should also give one a purchase on Closure-based Cartesian skepticism too. Recall that in this case the problem was that the Closure principle (which we saw in Chapter 1 to be highly compelling) appears to entail that insofar as one has much of the everyday knowledge that one thinks one has, then one must also be able to know the denials of radical skeptical hypotheses. Conversely, if one grants that the latter is impossible, then it follows that the former must be impossible too, and hence that we lack the everyday knowledge that we credit to ourselves.

But with epistemological disjunctivism in play, this argument can potentially be avoided. In particular, insofar as we have the rational support that favors our everyday beliefs over radical skeptical hypotheses just noted, then why should we accept that we are unable to know the denials of radical skeptical hypotheses? Indeed, as we have already noted with regard to how epistemological disjunctivism responds to Humean skepticism, there doesn't seem to be anything essentially epistemically circular about using one's factively grounded everyday knowledge as a basis for concluding that one knows the denials of radical skeptical hypotheses. So what then prevents us from biting the bullet in the Closure-based formulation of the Cartesian skeptical argument and embracing this anti-skeptical conclusion of accepting the Closure principle?

Epistemological disjunctivism might also be helpful in responding to Agrippa's trilemma (though this is a possibility that has not been explicitly explored in the contemporary literature). In particular, one might be able to use it to further a foundationalist response to that trilemma, such that one argues that there can be legitimate stopping-points to the regress of reasons. For wouldn't those beliefs that enjoy factive rational support look like excellent candidates in this regard?

In any case, it should be clear that the anti-skeptical ramifications of epistemological disjunctivism are such as to make the view attractive, insofar as it is defensible. With that in mind, let's turn to the issue of whether it is defensible.

5.3 Defending epistemological disjunctivism

The main challenge facing any defense of epistemological disjunctivism is explaining how one could have reflective access to factive reasons. Rather than run through all the main objections in this regard—see Pritchard (2012, Part 1) for a detailed summary—we will focus instead on what is generally considered to be the most serious challenge to the view, what is known as the *discriminability problem* for epistemological disjunctivism (ibid., Part 2).

Even epistemological disjunctivism does not deny that good cases are indistinguishable from the corresponding bad cases. And yet, if one has reflective access to factive reasons in the good case, and only in the good case, then it seems that one is in a position to thereby know, when one is in the good case, that one is in the good case (and hence not in the corresponding bad case). In particular, one can reflect on the nature of the rational support that one has for one's belief and come to know that it is rational support that one would only have in the good case. But if that's right, then why doesn't it follow that one can distinguish between the good and the bad case after all, contrary to intuition (even by the lights of epistemological disjunctivism)?

Note that the point of the argument is not that epistemological disjunctivism must treat pairings of good and bad cases as *completely* distinguishable, since it is not being denied that one cannot distinguish between good and corresponding bad cases when in the bad case. Nonetheless, that one can even distinguish between the two in the good case is surely a *prima facie* worry for the view, given that it is meant to be a point accepted by all sides that these cases are not distinguishable at all.

It will be particularly useful to consider how epistemological disjunctivism responds to this objection as it appeals to an epistemological distinction that is entirely independent of epistemological disjunctivism, and which might thus be useful to other anti-skeptical proposals. The distinction in question is between *favoring versus discriminating epistemic support*. In order to understand this distinction, let's consider again an example we encountered in Chapter 3. Suppose one is at a zoo and looking into the

clearly marked zebra enclosure, where there is a creature that looks very much like a zebra. In normal circumstances, where what looks like a zebra in a zebra enclosure really is a zebra, we would be happy to credit one with knowledge that what one is looking at is a zebra. If the creature is a zebra, however, then it follows that it is not a cleverly disguised mule. But does one know that it isn't a cleverly disguised mule? Seemingly not. For if one is an ordinary individual—rather than, say, a zoologist—one lacks the kind of discriminative powers that would enable one to tell the difference between a zebra and a cleverly disguised mule.

The problem, of course, is that one's failure to know that the creature is a cleverly disguised mule calls into question one's putative knowledge that the creature is a zebra. In particular, insofar as one grants that one knows that zebras are not cleverly disguised mules (as surely one does), then allowing that one could know that the creature is the former without knowing that it is the latter would seem to lead to some form of epistemological revisionism. For example, it would seem to mean denying the Closure principle that we formulated in Chapter 1. It would also seem to entail the rejection of the Underdetermination principle. And note that this puzzle is not confined to this case, as for just about anything we know we can conceive of an error-possibility in the vicinity which is indistinguishable from the scenario depicted in the target proposition.

Now one might think that the challenge posed by such zebra-style cases is simply a variant of the familiar Cartesian skepticism described above, which also turned in substantial part on our commitment to epistemological principles such as Closure and underdetermination. That would be too quick, however, as there is an important difference. Cartesian skepticism makes appeal to *radical* skeptical hypotheses, scenarios that are inconsistent with most of our ordinary beliefs. But the error-possibilities at issue in zebra-style cases are not radical in this sense at all, but rather just call into question a small sub-set of one's beliefs. Relatedly, they are also error-possibilities that one could have good reason to take seriously, such as if someone authoritative tells us to be beware that sometimes the creatures in the zebra enclosure are cleverly disguised mules. This is not the case when it comes to radical skeptical error-possibilities, as we saw in Chapter 1. The radical skeptic is not offering us a reason for thinking that we are BIVs, for example, as that would be self-defeating, but merely pointing out that it is an error-possibility that's inconsistent with our beliefs but which we can't exclude. As we might put the point, radical

skeptical hypotheses are in their nature *merely raised*, whereas local error-possibilities, like the one at issue in the zebra scenario, are usually *rationally motivated* (since otherwise, why raise them at all?).

These differences between radical skeptical error-possibilities and local error-possibilities point to how we might go about responding to zebra-style cases. Since it is essential to the former that it calls one's beliefs into question *en masse*, it also calls into question one's background beliefs that might be salient to epistemically assessing the error-possibility in hand. That's why it would be irrelevant to the BIV scenario to argue that we don't yet have the technology to implement such a deception, as of course one's conception of what the latest technology might be capable of is itself also put into question by this skeptical hypothesis. But this doesn't apply to zebra-style cases. Moreover, it seems that most people will have all sorts of background evidence that would bear on this scenario.

For example, suppose this error-possibility is merely raised, and so no supporting rationale is offered for why it should be taken seriously. Wouldn't it be normally quite straightforward to cite evidence for disregarding it? For instance, that the deception is highly unlikely, that reputable zoos would not be a party to such a deception, that it would be very difficult to disguise zebras as mules to the required level to prevent the deception being easily spotted, that there would be serious problems for the zoo management if the deception came to light, that the high costs involved in such a deception would undermine any potential benefit, and so on. The point is that it seems quite undemanding for the average person to offer reasons in support of their belief that the creature is a zebra that would favor this belief over the known-to-be-incompatible alternative that it is a cleverly disguised mule. So there is no obvious tension with the Underdetermination principle. Furthermore, if favoring reasons are available, then it isn't clear why one couldn't come to know on this basis that the creature is not a cleverly disguised mule as well, which would also remove the putative tension with the Closure principle. Conversely, if one is unable to offer such favoring reasons, then it would follow that consideration of the error-possibility defeats one's knowledge that the creature is a zebra, but in that case the putative tension with the Closure and Underdetermination principles also disappears (as the antecedent in both cases is false).

If a supporting rationale for this error-possibility is offered, then of course that raises the epistemic bar, as one needs grounds that speak to the specific reasons offered for taking it seriously. But the crux of the matter is

that the difference would be merely one of degree; the point remains that there is no in principle reason why one couldn't satisfy the epistemic demands that the Closure and Underdetermination principles raise in this case. In short, regardless of whether the error-possibility is merely raised or rationally motivated, there is still no in principle reason why a normal epistemic agent wouldn't be able to rationally exclude it.

So what happened to our original puzzle? In order to understand what is going on here, we need to look back at the way this puzzle was framed. Recall that the issue was expressed in terms of a discriminative power that the subject lacked—in this case, the ability to tell zebras and cleverly disguised mules apart. This has not been disputed. Most of us would clearly lack such a discriminative power, as it requires special training (or at least a close-up observation of the creatures that is not available in the case as presently described). What we've seen, however, is that there is a way of excluding error-possibilities that doesn't involve possession of such a discriminative power—i.e., by being able to appeal to relevant reasons, in this case salient background evidence. This point gets obscured if we run zebra-style cases alongside the more familiar radical skeptical scenarios, as the latter also call one's background evidence into question, thereby removing any significance it might have for dealing with the problem. But background evidence can be legitimately cited when it comes to local error-possibilities like the zebra case. Even so, possessing such background evidence, while surely relevant to whether one can exclude the local error-possibility in play, doesn't imply that one has any special kind of discriminative power. This highlights the point that there can be ways of knowing that one scenario obtains rather than a competing scenario which don't involve a discriminative power to tell them apart. As we might put the point in sloganizing form: There are ways of knowing the difference which aren't thereby ways of telling the difference. This is the distinction between favoring and discriminating epistemic support, and the zebra puzzle, it seems, trades on a failure to recognize it (along with a related failure to recognize that background evidence can usually be legitimately employed when it comes to dealing with local error-possibilities). In particular, that one is unable to discriminate between two scenarios doesn't mean that one is unable to know that the one scenario obtains rather than the alternative.

The distinction between favoring and discriminating epistemic support is meant to be one that all epistemologists should accept, but it has particular relevance to epistemological disjunctivism. Recall that the challenge

this view faces is to explain how one might be able to know that one is in the good case rather than the bad case, in virtue of the possession of factive reasons, even though one is unable to discriminate between these two cases. With this distinction in play, the proponent of epistemological disjunctivism can account for the former by appealing to favoring epistemic support. In particular, the idea is that factive reasons provide one with a definitive form of favoring epistemic support, which enables one to know that one is in the good case rather than the bad case even though one cannot discriminate between them.

Now one might contend that the distinction is inapplicable here precisely because 'bad' cases can involve not just local error-possibilities but also radical skeptical error-possibilities, like the BIV scenario. Didn't we note above that the distinction between favoring and epistemic support wasn't relevant in the context of radical skeptical error-possibilities because of how they also call into question one's background knowledge? Notice, however, that epistemological disjunctivism is not appealing to background evidence at all in this regard, but rather to the actual rational support for one's belief. Of course, according to epistemological disjunctivism, one only has the factive rational support in question in the good case, and whether one is in fact in the good case is what is in question, so it might seem that this makes little difference. If one can't ordinarily appeal to background evidence in the face of radical skeptic hypotheses because the latter calls the former into question, then how could one appeal to rational support that one would only have given that the radical skeptical hypothesis is false?

There is a key contrast between our background evidence and the factive rational support that epistemological disjunctivism is concerned with, however, which concerns the factivity of the latter. Appealing to background evidence in the face of radical skeptical hypotheses looks dubious precisely because it doesn't speak to the skeptical hypothesis at all. So what if it seems to you as if the current state of technology is inconsistent with being a BIV—what bearing could that possibly have on whether you are BIV, given that even a BIV could possess the very same (apparent) evidence? The factive reasons putatively possessed in the good case are not like that, however, for they actually entail the falsity of the skeptical hypothesis. Accordingly, unlike the reasons constituted by one's background beliefs, they are not reasons that are available to you regardless of whether the radical skeptical hypothesis is true.

The radical skeptic will inevitably dispute that one possesses the factive reasons even in the good case, but that is by the by in this context, given that epistemological disjunctivism has rejected the new evil genius intuition that is meant to motivate the non-factivity of our reasons. Remember too that the radical skeptic is not offering us any reason for thinking that we are the victims of radical skeptical scenarios; these error-possibilities are merely raised rather than rationally motivated. Indeed, as we noted in Chapter 1, Section 1.2, the radical skeptical problem is meant to be such that it doesn't matter which circumstances one is actually in, as whatever the nature of one's circumstances, the skeptical argument will run and wreak its epistemic havoc. But with all that in mind, what is to prevent the proponent of epistemological disjunctivism from insisting that whether the skeptical argument does go through depends crucially on whether one is in fact in the good case, for if one is, then one has rational support available to one which decisively excludes radical skeptical scenarios?

Notice too that the dialectical situation in this regard is very different to the stand-off between radical skepticism and the anti-skeptical responses rooted in epistemic externalism that we examined in Chapter 2, Section 2.3. Recall that the epistemic externalist had a story to tell concerning how, *contra* the radical skeptic, knowledge was possessed insofar as certain external conditions obtained. The problem, however, was that since that knowledge was contingent on external conditions obtaining—i.e., conditions that the subject lacked any good reflectively accessible reasons for thinking had obtained—this strategy seemed at a considerable dialectical disadvantage when it came to dealing with radical skepticism. In particular, it seems plausible that even if one doesn't need to meet an internalist epistemic condition in order to possess knowledge, one does need to meet such a condition if one is to properly claim knowledge, as one must be able to offer good reasons in support of the knowledge that one is self-ascribing. The epistemic externalist is thus in an odd dialectical position, as although their proposal defends the claim that knowledge is possible, *contra* the radical skeptic, they are not in a position to properly claim that they actually have any knowledge.

One might think that epistemological disjunctivism is in the same bind in this regard. After all, one only has the factive reasons in the good case. Isn't that akin to the claim made by epistemological externalism that one's knowledge is contingent on certain external conditions obtaining? There is a crucial difference, however, which is that the factive reasons at issue in epistemological disjunctivism are reflectively accessible to the subject.

In particular, there isn't any barrier to the subject properly claiming knowledge, as they do have supporting reasons that they can offer in defense of their claim to know.

One could imagine the radical skeptic insisting in response to this line that there is something dialectically inappropriate nonetheless about claiming factive reasons when the problem of radical skepticism is at issue. After all, if one were, say, a BIV, one would also cite factive reasons, so how can it be appropriate, in a context where a radical skeptical hypothesis has been raised, to present factive reasons in response to the skeptical challenge? There are at least three ways of understanding this line of reasoning.

Construed one way, it is focused on the nature of one's rational support, with the radical skeptic maintaining that one doesn't have factive rational support once radical skeptical hypotheses are raised. So construed, why should epistemological disjunctivism accept it? Since epistemological disjunctivism rejects the very idea that the strength of the rational support enjoyed by our beliefs is constrained by the rational support available to one in the bad case, why should they concede that merely raising radical skeptical hypotheses—and remember that they are merely raised; the radical skeptic has offered no reason for thinking that one is in the bad case—suffices to limit that rational support?

A second construal of this objection involves arguing that even if the subject has factive reasons available to her, it is nonetheless dialectically inappropriate for her to cite them in a context where radical skeptical hypotheses are at issue because it involves some kind of epistemic circularity. But we have already argued above that there is nothing epistemically circular about using knowledge supported by factive reasons in order to infer the denials of radical skeptical hypotheses. Accordingly, it can't make any difference whether one happens to be in a dialectical context where such error-possibilities are explicitly raised.

This leaves the third construal. Like the second construal, this also argues that even if the subject has factive reasons available to them in the good case, it is nonetheless inappropriate to appeal to those reasons in a dialectical context where radical skeptical hypotheses have been raised. The rationale in this case concerns the propriety of the subject's claim to have knowledge in this regard. In particular, it contends that it is not enough for the propriety of a claim to know that one can adduce rational support. Instead, one must be able to adduce rational support that is dialectically

appropriate. The thought would thus be that appealing to factive reasons in this dialectical context would be inappropriate. Given the decisive nature of factive reasons, the only way to make this claim stick is to insist that in dialectical contexts where a particular error-possibility is at issue, then one can only present supporting reasons for a claim to know that would suggest that one can discriminate between the believed scenario and the alternative described by the error-possibility. If that were right, then epistemological disjunctivism would face a problem, as of course the rational support available to them in the good case merely favors the target proposition over skeptical alternatives (albeit decisively).

Such a constraint on appropriate claims to know would be oddly demanding, however, to say the least. Going back to our zebra case from earlier, for example, couldn't one claim to know that the creature before one in the zebra enclosure is a zebra on the basis of the favoring reasons that one can offer, even if one is in a context where someone has (merely) raised the cleverly disguised mule error-possibility? Why would it matter to the propriety of one's claim to know that one can't discriminate between zebras and cleverly disguised mules? At most, it might be a conversational implicature of one's claim to know that one is able to undertake the relevant discrimination, but if that's true, all it would mean is that one should cancel the implicature (e.g., by making clear one's favoring rational basis for one's belief). If a claim to know is plausible here, then it ought to be even more plausible in the case in hand, given that the factive reasons that epistemological disjunctivism maintains are available to one in the good case actually entail the falsity of the (again, merely raised) radical skeptical hypothesis. It seems, then, that there wouldn't be anything essentially amiss in claiming knowledge on the basis of factive reasons, even if one does so in the context where radical skeptical scenarios are under consideration.

5.4 The 'biscopic' response to radical skepticism

We have already noted the anti-skeptical ramifications of epistemological disjunctivism, so insofar as one grants that it is a viable position, then it is clearly attractive in this regard. In particular, we have seen how epistemological disjunctivism seems able to have the resources to offer a very direct way of dealing with both Cartesian and Humean forms of radical skepticism. We will close, however, by remarking on a recent proposal—like epistemological

disjunctivism, due to Pritchard (2015)—that, rather than viewing epistemological disjunctivism as offering a response to radical skepticism in isolation, it should instead be viewed as operating in concert with the kind of hinge epistemology that we examined in Chapter 4. This is what is known as the *biscopic* response to radical skepticism.

Back in Chapter 1, Section 1.2, we presented the underdetermination-based and Closure-based formulations of Cartesian skepticism as being simply equivalent expressions of the same skeptical puzzle. This is indeed, how these two formulations have often been understood in the contemporary literature. There is, however, reason to doubt this supposed equivalency. In particular, Closure seems like the more demanding principle, at least on closer inspection. In the context of the skeptical puzzle, for example, Closure effectively demands that if one has the everyday knowledge at issue, then one also knows the denials of radical skeptical hypotheses. In contrast, from the same antecedent, the Underdetermination principle only demands that one's epistemic basis favors one's belief in the everyday proposition over radical skeptical alternatives. The latter is a weaker claim, as one could have an epistemic basis that favors everyday beliefs over radical skeptical alternatives without thereby having an epistemic basis that suffices for one to know that the radical skeptical alternatives are false.

More generally, aside from the logical differences between these two formulations, Pritchard argues that they trade on very different underlying theses, what he calls "sources of skepticism." He claims that what is driving the Closure-based formulation of the problem is a thesis that he terms the *universality of rational evaluation*. This is the idea that there is no in principle barrier to extending the scope of rational evaluations. We can see this idea at work in Closure-based skepticism in terms of how plugging radical skeptical hypotheses into Closure-based inferences effectively leads to rational evaluations of one's commitments in general, as opposed to merely local rational evaluations. The universality of rational evaluations thesis can look harmless. For what principled basis could there be to limit the scope of rational evaluations? What limitations there are appear to be merely practical barriers, such as time, imagination, consistency, ingenuity, and so on. Relatedly, while we do not undertake universal rational evaluations in our everyday epistemic practices, this seems to merely reflect these practical barriers, such that we could reasonably regard a commitment to the universality of rational evaluations thesis as being at least implicit in our everyday practices.

As we saw in Chapter 4, however, according to Wittgenstein, we should reject the universality of rational evaluations thesis. Indeed, Wittgenstein is more specifically arguing that not only is this thesis not rooted in our everyday practices, but it is a theoretical claim that should be discarded as erroneous. This is key to his undercutting response to radical skepticism, as he wants to get us to see that there is no fundamental tension in our epistemological concepts as the radical skeptic maintains; the skeptical puzzle is instead arising out of faulty philosophical theses masquerading as commonsense. Related to this, we also saw in Chapter 4 that the Wittgensteinian response to radical skepticism needn't result in the rejection of the Closure principle. In particular, if one accompanies a rejection of the universality of rational evaluation thesis with an account of hinge commitments that doesn't treat them as K-apt beliefs, then one can endorse the Closure principle without this having any import for radical skepticism (which further reinforces the undercutting credentials of this anti-skeptical approach).

Underdetermination-based skepticism doesn't seem to trade on the universality of rational evaluation thesis, however. In particular, what is at issue in this formulation of radical skepticism is not a claim about the scope of rational evaluation, but rather a claim about the nature of the rational support our beliefs might enjoy. Pritchard refers to the source of skepticism in this case as being the *insularity of reasons* thesis. This is the idea that, even in optimal epistemic conditions, the rational support enjoyed by our external world beliefs is consistent with those beliefs being overwhelmingly false. As we saw above, the motivation for this thesis comes not from our everyday epistemic practices, which in fact are suffused with factive reasons, but rather from a theoretical claim, encapsulated in the new evil genius intuition. It is precisely this theoretical claim that epistemological disjunctivism denies. In particular, epistemological disjunctivism maintains that we should take our everyday epistemic practices at face value and embrace the factive reasons that they attribute to us in paradigmatic epistemic conditions. This is key to the undercutting credentials of this anti-skeptical strategy, in that, as with the Wittgensteinian line, the idea is that radical skepticism is not a paradox involving a deep tension in our epistemological concepts, but is rather the consequence of faulty philosophical reasoning.

Hinge epistemology thus seems to be the antidote to Closure-based skepticism, in virtue of its rejection of the universality of rational evaluation thesis, while epistemological disjunctivism seems to be the antidote

to underdetermination-based skepticism, in virtue of its rejection of the insularity of reasons thesis. If Closure-based and underdetermination-based formulations of Cartesian skepticism trade on distinct sources of skepticism, then there need be no guarantee that an adequate response to one formulation is thereby an adequate response to the other. We saw in Chapter 4 that some proponents of hinge epistemology, such as Genia Schönbaumsfeld (2016), argue that it can also respond to underdetermination-based skepticism, and we have noted above that epistemological disjunctivism might also have the means to respond to Closure-based skepticism too. Pritchard (2015) argues, however, that an optimal solution to the Cartesian skeptical problem treats it as two distinct problems requiring an integrated, dual solution. This is where the biscopic response to Cartesian skepticism comes in.

The guiding idea behind the biscopic response is that hinge epistemology and epistemological disjunctivism are stronger as anti-skeptical proposals when integrated with one another, in that attempting to resolve skeptical problems that turn on two distinct sources by appealing to a proposal that is only directly applicable to one of these problems is bound to weaken its effectiveness. In particular, while these anti-skeptical proposals are usually thought of as competing, once we understand their respective anti-skeptical targets, then we can see how they can be not only consistent with one another, but also mutually supporting.

Let's start with hinge epistemology (as presented in Chapter 4, as opposed to other versions of it, such as that defended in Coliva (2015), which we will review in Chapter 7, Section 7.5). Taken by itself, it tells us nothing about the nature of the rational support our beliefs enjoy in good cases, but only that a certain kind of rational evaluation is incoherent. In particular, insofar as we take seriously that our hinge commitments are not known, then it seems that the Wittgensteinian line as applied to underdetermination-based skepticism must involve the denial of the Underdetermination principle. After all, if our everyday knowledge coexists with a failure to know the denials of radical skeptical hypotheses, then why think that we could possess the favoring epistemic support at issue in the Underdetermination principle? But, of course, if we are obliged to deny this principle, then that makes it hard to sustain the undercutting credentials of this anti-skeptical strategy.

Allying hinge epistemology with epistemological disjunctivism removes this difficulty. There is now no need to deny the Underdetermination principle, as in the good case one has the favoring rational support at issue

in underdetermination-based skepticism. Moreover, notice that hinge epistemology is much easier to stomach if it is combined with epistemological disjunctivism. That one's system of rational evaluation has arational commitments at its heart is far less troubling if the rational support that one's perceptual beliefs enjoy can be factive. (Consider, in contrast, the alternative of embracing hinge epistemology along with the idea that the rational support one's perceptual beliefs enjoy even in the best case can be compatible with the widespread falsity of one's beliefs.)

Now consider epistemological disjunctivism. Since this is specifically concerned with the strength of the rational support available for our perceptual beliefs in good cases, it doesn't directly tell us anything about the nature of rational evaluation. In particular, it seems entirely consistent with there being universal rational evaluations. Indeed, the application of epistemological disjunctivism to the Closure-based formulation of Cartesian skepticism suggests that it is part of a traditional anti-skeptical proposal that rationally evaluates our commitments as a whole and returns a positive verdict. If we are convinced by the considerations raised by Wittgenstein with regard to our essential reliance on hinge commitments, however, then we should be immediately suspicious of such an extension of epistemological disjunctivism to Closure-based radical skepticism.

In contrast, if we combine epistemological disjunctivism with a hinge epistemology along the lines set out in Chapter 4, then we have a way of responding to Closure-based radical skepticism that allies the undercutting response to underdetermination-based radical skepticism that epistemological disjunctivism provides with a parallel undercutting response to Closure-based radical skepticism, one which doesn't result in the rejection of the Closure principle. Moreover, epistemological disjunctivism ought to be more palatable when so combined with a hinge epistemology, since it is no longer committed to the idea that one can have knowledge of the denials of radical skeptical hypotheses, given that they form part of our set of hinge commitments. Since it is widely held that such knowledge is counterintuitive, this ought to be a point in favor of epistemological disjunctivism.

Epistemological disjunctivism and hinge epistemology, far from being competing anti-skeptical proposals, might thus be combinable as part of an integrated biscopic treatment of the Cartesian skeptical problem, one that treats this difficulty as effectively two distinct puzzles in disguise.

Further reading

See McDowell (1995) for his key defense of (the view we are describing) epistemological disjunctivism. For an exegetical discussion of McDowell's view and its relevance to contemporary epistemology, see Neta and Pritchard (2007). For a detailed presentation of epistemological disjunctivism, including its compatibility with both accessibilist and mentalist conceptions of epistemic internalism, see Pritchard (2012). See also Pritchard (2008; 2015, Part 3). Pritchard (forthcoming a) defends the claim that epistemological disjunctivism is compatible with externalism about knowledge more generally. For two key defenses of the new evil genius intuition, see Lehrer and Cohen (1983) and Cohen (1984). For an overview of the main issues with regard to this putative intuition, see Littlejohn (2009). For another prominent proposal that rejects this intuition, albeit by appealing to a form of epistemic externalism rather than epistemic internalism, see Williamson (2000). For some of the main defenses of metaphysical disjunctivism, see Hinton (1967a; 1967b; 1973), Snowdon (1980–1981; 1990–1991), McDowell (1982; 1986; 1994), Martin (1997; 1998; 2002), and Brewer (2000; 2011). For a survey of work on metaphysical disjunctivism, see Soteriou (2020). For a general discussion of the relationship between disjunctivism (both metaphysical and epistemological) and radical skepticism, see Pritchard and Ranalli (2016). The zebra case is described in Dretske (1970). The distinction between favoring and discriminating epistemic support is defended in Pritchard (2010), and applied in the service of defending epistemological disjunctivism in Pritchard (2012, Part 2). For further discussion of the logical relationship between Closure-based and underdetermination-based formulations of the radical skeptical puzzle, see Brueckner (1994), Cohen (1998b), and Pritchard (2005a; 2005c). The biscopic response to radical skepticism is primarily defended in Pritchard (2015, Part 4). For an overview, see Pritchard (2018a); for some responses to recent critics, see Pritchard (2021b).

6

MOORE, LIBERALS, AND CONSERVATIVES

6.1 Introduction

Let us take up the Humean paradox again. Following G. E. Moore (1939), we will consider 'There is an external world' as its conclusion (even though it could just as well be 'I am not a BIV'):

(1) I know that I have hands.

(2) I know that 'I have hands' entails 'There is an external world.'

(C) I know that there is an external world.

DOI: 10.4324/9780429057946-7

As we have already seen in Chapter 1, Section 1.4, the Humean skeptical paradox does not call into question the logical validity of this argument. Hence, if we know 'I have hands,' we also know the conclusion 'There is an external world' through correct inference, since hands are physical objects, which are external to our minds, whose existence can be taken to support the conclusion that there is an external world—that is, a world populated by physical objects that exist independently of our minds. Moreover, the Humean paradox can concede, at least for the sake of the argument, that we know 'I have hands.' Still, the paradox arises once the nature of our justification for 'I have hands' is considered (on the assumption that justification is a necessary component of knowledge). To bring that out more clearly, let us consider the Humean paradox in this form:

(1*) I am justified in believing that I have hands.

(2*) I am justified in believing that if I have hands, then there is an external world.

(C*) I am justified in believing that there is an external world.

The problem arises once we consider the nature of our justification for 'I have hands.' This is typically afforded by perception. However, according to Humean skepticism, in order to be perceptually justified, it is not enough to have an experience with a given representational and phenomenal content (in this case as of one's hands). After all, that experience can be subjectively identical even when produced under non-standard conditions, as we saw in Chapter 1, Section 1.4. Hence, that experience can give one a justification for 'I have hands' only on the assumption that there is an external world. Now, according to the Humean skeptic, such an assumption is rational only if there is a justification in its favor. However, this is precisely the conclusion (C*) we would like to achieve by means of this very argument. Therefore, Moore's argument cannot produce a *first* justification for 'There is an external world,' because premise (1*) depends on the conclusion (C*). That is, for one to be justified in believing 'I have hands,' one needs a prior justification to believe 'There is an external world.' Thus, the argument is epistemically circular (or question-begging) and cannot produce a justification for 'There is an external world.'

Now, if there is no other way to obtain such a justification—and that is what the Humean skeptic claims—then such a belief is not epistemically

rationally supported. Indeed, as we saw in Chapter 1, Section 1.4, that was precisely Hume's conclusion (1739, 1.4.2). In his view, that belief comes naturally to us by dint of the operations of the faculty of imagination. Key to its workings is the fact that it obliterates the distinction between various perceptions, and that it projects the properties of perceptual representations onto the objects which (allegedly) caused them. This way, for instance, we form the belief that there is an object, such as a hand, with its characteristic shape, color, and tactile properties, that exists even when we are not perceiving it and that remains identical under variations in our perceptions due to changes in our orientation with respect to it. Yet, these properties, for Hume, are not in the object itself, but are projections of our own minds (ibid.).

Hume's conclusion gives rise to a host of problems. First, there is the issue of whether the account of perception he is relying upon is correct. Thomas Reid, another great Scottish philosopher, a contemporary and major adversary of Hume, already disputed such an account of perception, by arguing in favor of the idea that in perception we directly take in objects. Of course, according to Reid ([1764] 1997, 6.2, 81–82, and [1785] 2002, 2.14, 181, 183), our sensory system operates according to specific laws and therefore we may perceive distant objects as elliptical when they are round. Yet, this is simply the way the object itself is given to us through perception. Hence, there is no veil of ideas—that is, of intermediate mental entities—between us and objects out there, in perception.

Furthermore, as we saw in Chapter 5, disjunctivists nowadays dispute Hume's claim, by holding that in perception we reach all the way down to the physical world. In perception, on their view, we do not represent objects but take them directly into our consciousness. That is, perception does not produce representations of distal objects. Rather, it puts our minds directly in contact with them. This view, however, proposes a contentious account of perception (for poignant criticism, see Burge 2010) and, moreover, gives rise to a second-order skeptical paradox (on this issue, see Wright 2002). For one may concede that if we are perceiving our hands, then we know that we have hands and that we have factive reasons to hold that belief. However, since perceptions could be indistinguishable no matter what their causal origins might be, how could we know that we are in fact perceiving, rather than merely hallucinating our hands? That is, how can we *claim* to have the factive reasons and knowledge that we take ourselves to have? In other words, the Humean

skeptical paradox could take the following form: how can we know and therefore claim, just based on our current experience, that we are in a good case scenario, thereby perceiving our hands, rather than in a bad case one, thereby only hallucinating them? Any further appeal to the phenomenologically salient features of our experience would be compatible with our being in such a bad case scenario that everything would look to us to be identical to what we take to be an ordinary perception, while we are merely hallucinating. Similarly, any appeal to the absence of defeaters—understood as subjectively available reasons to think that our perception may not be accurate—would be compatible with such a bad case scenario. Thus, disjunctivism cannot be a remedy against Cartesian skepticism, at least on its own. (In Chapter 5, we considered some responses to these charges, and also the possibility that disjunctivism ought to be combined with hinge epistemology to make it more palatable as an anti-skeptical strategy.)

Second, even conceding a representationalist account of perception, is it true that we need to presuppose the existence of an external world, and indeed the justification of the corresponding belief, in order to have a perceptual justification for 'I have hands'? After all, it seems that no such collateral assumption is typically entertained by subjects when forming such a belief based on their current perceptual experiences. Nor is it the case that subjects are requested to be able to justify it.

Third, even conceding that we need to presuppose the existence of an external world, as well as the justification of the corresponding belief, in order to have a perceptual justification for 'I have hands,' is it really the case that such an assumption cannot be rationally redeemed independently of this very argument?

Finally, and more generally, is it really the case that philosophical speculations about the alternative etiology of our current experience, or else about the very nature of perception, or indeed about the epistemic structure of perceptual justification can shake our commonsensical intuition that we are seeing our hands? That is, are we under any real pressure to give up our commonsensical view of the world, based on philosophical arguments, the premises of which are of no more secure a standing than our conviction of seeing our hands here and now?

In this chapter, we are going to look at three prominent accounts that explore each of these last three alternatives. That is, G. E. Moore's appeal to common sense (Section 6.2), Jim Pryor's liberal account of perceptual justification (Section 6.3), and, finally, Crispin Wright's conservative account of

perceptual justification and his attendant notion of epistemic entitlement for basic assumptions such as 'There is an external world' (Section 6.4). We will then close the chapter by looking at Ernest Sosa's and Ralph Wedgwood's different ways of defending the idea that we have an a priori justification to believe that there is an external world (Section 6.5).

6.2 G. E. Moore's proof of the external world

In a famous essay entitled "Proof of an external world," George Edward Moore (1939) claimed that he could prove the existence of the external world. Moore conceived of his argument as directed against idealism and not against skepticism, but from the beginning his proof was read also and predominantly as an argument against skepticism. Interestingly, Moore's original proof had a different structure, but today it is mostly rendered as follows:

(1) Here is a hand.

(2) If there is a hand here, then there is an external world.

(3) There is an external world.

According to Moore, the first premise is justified (and even known) based on his occurrent experience as of a hand in front of him. The second premise is known *a priori*, and therefore justified, based on reflection on our concepts of hand and of external world. Finally, the conclusion is correctly obtained through *modus ponens*. Furthermore, Moore did not raise any objection to the Closure principle. Indeed, he claimed that his proof was a good one since the premises were known (and therefore justified), the conclusion was different from the premises, and it was correctly reached from them. Hence, he held that the proof allowed him to *know* the truth of its conclusion. Therefore, the proof has usually been taken to be an overriding anti-skeptical argument, despite Moore's protestations.

Still, Moore did not actually believe that this proof could solve the skeptical paradox. In fact, he did not believe that he had *demonstrated* that he knew premise (1). Yet, for him, to counteract a skeptic, it would have been necessary to do so. That is, Moore took the skeptical argument to be a second-order one, where knowledge of the premises of the proof—particularly premise (1)—is not merely assumed, in the way he did in his proof, but is demonstrated and proved.

Yet, he also claimed that there were many things he knew while being unable to prove them (Moore 1939, 150), chief among them, premise (1) of his proof. That is, Moore held on to a protoexternalist account of knowledge, according to which one can know something even if one cannot prove or demonstrate how one knows it (on this issue, see Coliva 2010a, Chapter 1). Thus, he took himself to know (1), while consistently, from his point of view, rejecting the idea that, by giving his proof and thus reaching knowledge of (3), he would thereby oppose skepticism. For the latter, according to him, required a proof of one's knowledge of the premises of the argument, which he did not provide and had no idea how to provide (Moore 1939, 150).

In other writings (see Moore 1959a), Moore also argued that it seemed to him unequivocally more certain to be seeing a hand in front of him than to be the victim of a dream. Then the conclusion that there was an external world on which that experience depended was more certain than was its skeptical counterpart. However, Moore himself realized that a skeptic would not have been very impressed with such an argument.

A skeptic might say, in fact, that what Moore seems to be certain of—or more certain of than anything else—does not prove that things are as they seem to him. More generally, the skeptic is raising a logical-metaphysical possibility, the subjective probability of which does not necessarily have to be greater than that of its non-skeptical counterpart to give rise to the paradox. Therefore, Moore was always dissatisfied with his proof, understood as an anti-skeptical argument. Not by chance, the papers in which he tried to reach an anti-skeptical conclusion, based on his proof, were all published posthumously, since, while alive, he never authorized their publication.

What is interesting, though, in what has become known as 'Moore's gambit,' which was anticipated in Reid ([1764] 1997; [1785] 2002), is the appeal to common sense. That is, the idea that philosophical arguments, which are meant to discard our commonsensical picture of the world, rely on premises and inferential principles which are of no more secure a standing than that very picture itself. Hence, for a philosopher of common sense like Moore, we are under no real pressure to give up on the commonsensical picture of the world, based on contrary philosophical arguments. (As we saw in Chapter 4, Wittgenstein too noticed that we do not acquire our hinge commitments from experience, let alone philosophical argumentation.)

From the point of view of Humean skepticism, then, Moore's proof is the exemplification of that very paradox. Indeed, according to a Humean skeptic, in order to have a justification for (1), it is not enough merely to have an experience with a representational and phenomenal content as of a hand. It is also necessary to assume that it was produced by interaction with the external world, and, for this assumption to be rational, that assumption must in turn be justified. Thus, the proof cannot give one a first justification for believing that there is an external world. If there is no other way to justify that assumption, and, according to a Humean skeptic there isn't, (3) will be unjustified and unjustifiable.

Indeed, Moore was a sense data theorist and did not think that in perception we directly take in objects. His opinions regarding the nature of sense data changed a lot during the years, ranging from the idea that sense data are just figments of our minds, to the notion that sense data are identical to the perceived surface of physical objects. Yet, even this last idea, which could have been a starting point for rejecting the account of perception on which the Humean skeptical argument rests, was found wanting. For Moore noticed (1925, 54–55) that if sense data are the direct object of perception, then since sometimes we see double, that is, we have two sense data while there is only one object causing them, sense data cannot be identical to the perceived surface of the object.

Finally, there is also an element of commonality between Hume and Moore (as well as Reid). For Hume (like Reid) thought that philosophy is powerless *vis-à-vis* our human nature. That is, even though philosophical arguments could lead us to be suspicious of our most ingrained beliefs, such as that there are physical objects and that we perceive them, we cannot help holding them. (As we saw, however, Reid, contrary to Hume, thought that philosophical arguments leading to skepticism were in fact faulty and that we should trust the deliverances of our senses.) Similarly, for Moore, even if philosophical arguments could raise doubts about our commonsensical view of the world, our commitment to common sense is stronger than any conviction philosophy may produce, or than any premises philosophical arguments could rest upon, at least *prima facie*. Hence, just as for Hume and Reid, nature has always the upper hand to philosophy, so, for Moore, does common sense have the upper hand to philosophy. This was an element of Moore's position that greatly influenced Wittgenstein, as we will see in Chapter 7, Section 7.3.

6.3 Pryor and the liberal account of perceptual justification

In recent years, Moore's proof has been resumed with favor in the context of contemporary epistemology by Jim Pryor. According to Pryor (2000; 2004), it can be argued that we have *perceptually basic beliefs* which we form directly based on experience and which we are justified in holding immediately *with no need for any additional information* (Pryor's position echoes those of Pollock 1974, Alston 1989, and Audi 1993). This is the gist of what Pryor (2000) calls the "dogmatist" or "liberal" account of the justification of empirical beliefs, of which he considers Moore and himself the clearest examples (a similar position has been defended by Michael Huemer (2001), under the label of "phenomenal conservatism"). Pryor's account is called "liberal" because, contrary to skepticism, it does not require any background assumptions for experiences to provide us with a *prima facie* justification to believe that there are specific objects in our surroundings—let alone a justification for them. Furthermore, it is a form of dogmatism because, contrary to skepticism, it claims that we do have knowledge of 'there is an external world.'

According to Pryor, among the basic beliefs there is 'Here is a hand,' when it is formed based on our current perception, absent defeaters—that is, in the absence of any reason we may be aware of to think that our perception may not be accurate. If this is the case—that is, if we do not need any prior justification for (3) to have a perceptual justification for (1)—then we have an immediate justification for (1) that is transmitted to (3) through the correct inference. Moore's argument therefore isn't circular and provides us with a direct answer to Humean-style skepticism, by showing that it rests on a mistaken conception of the structure of perceptual justification. In short, Pryor's strategy is an undercutting one. For, if he is right, there would be no Humean skeptical paradox to respond to in the first place.

However, Pryor does not want to argue that the proof is successful against skepticism. According to him, in fact, it is indeed ineffective but for merely *dialectical* reasons, not for epistemic ones. For the proof is directed at a skeptic who doubts the existence of the external world. That is, a skeptic considers (probably) false that there is an external world, and thus does not concede that there is a perceptual justification for (1). Since (1) is not justified from a skeptic's point of view, there is no justification that can be

transmitted from (1) to (3) through the correct inference. Therefore, according to Pryor's reconstruction, a skeptic does not consider Moore's argument to be circular but simply unable to produce a justification for (3), since it starts with a premise that is not justified, at least from her point of view.

Unlike Moore, Pryor thinks that 'Here is a hand' is a perceptually basic belief. As we mentioned in Section 6.2, Moore believed that we have direct knowledge only of *sense data*, the precise nature of which remained mysterious. He oscillated between identifying them with the visible surfaces of objects and considering them as mental entities. This last conception would make it very problematic to attribute to Moore the thesis that there are perceptually basic beliefs concerning physical objects. In fact, Moore himself argued that in order to have a justification for an empirical belief such as (1) based on one's own sensory evidence, it was also necessary already to take for granted that there is an external world (Moore 1959b, 226). Yet, Pryor is not so much aiming at providing an exegetically accurate reading of Moore as at building on some of his insights to provide an understanding of the proof which could be successful against Humean skepticism. Thus, let us assume that Pryor is right to hold that there are perceptually basic beliefs. Even so, some problems arise.

First, how do we determine the extent of perceptually basic beliefs? For Pryor, in fact, an experience as of a policeman in front of him, or of a zebra, would not give him immediate justification for 'Here is a policeman (or a zebra) in front of me.' The reason is that the person might just be dressed as a policeman but not be a policeman; or, similarly, there might be a mule cleverly disguised as a zebra. Then even the object that, based on one's current visual experience, seems to one to be a hand, could be a piece of plastic cleverly forged to look like a human hand, and so on for many other empirical beliefs.

Perhaps the problem of determining the class of perceptually basic beliefs can be solved by assuming that their content is something very generic such as 'Here is a physical object, with a hand-like shape and color.' This retreat, however, does not fit well with the content of the belief that we should allegedly get directly from sensory experience—that is 'Here is a hand,' and a myriad similar one.

Leaving these preoccupations aside, it is important to note that, for Pryor, perceptually basic beliefs must be about material objects, otherwise they could not serve as a starting point for an argument against Humean

skepticism. Yet, it is obvious at this point that we should have some kind of justification for believing that simply by having an experience as of a physical object with a certain shape and color in front of us, we would get a justification for the belief in the existence of that material object, thereby disconfirming the skeptical counterpart of that belief.

Now, Pryor admits that if we were dreaming that there is a material object here with a certain shape and color, we could have an experience subjectively identical to the one we have when we see it. Hence, the experience as such seems to justify both (1) 'Here is a physical object with a hand-like shape and color' and (1*) 'I am dreaming that there is a physical object with a hand-like shape and color here.' Therefore, experience can only justify the former by excluding the latter unfavorable case. It can therefore serve as a justification for (1) only by already assuming the conclusion of the argument.

A first answer consists in saying that the notion of justification at stake is that of *fallible* justification. Therefore, one can have a non-conclusive justification for (1) 'Here is a physical object with a certain shape and color' while not assuming (3). Therefore, the argument would not be circular. Yet, it must be noted that Pryor takes it for granted that experience gives only a fallible justification for (1). The point is that for it to be a pertinent (albeit fallible) justification for (1) and not for (1*), one must already assume—perhaps implicitly—that there is a world of physical objects with which one is causally interacting. Therefore, (3) would need to be assumed and the argument continues to be circular. That is, we need to assume (3) in order to be within our rights in taking our current experience to bear on a belief about a specific physical object. To use Wright's (2004, 172–174) apt phrase, we need to assume (3) in order to overcome "our cognitive locality."

Moreover, there are those who have argued (Cohen 2002; 2005b; Vogel 2002) that any theory that contemplates forms of "basic knowledge" exposes itself to the so-called problems of "easy knowledge" and "bootstrapping." Let us take a closer look at them.

Those who support the existence of forms of "basic knowledge" say that we can have immediate justification—internalist or externalist—to believe, for example, that there is a hand here based on our experience. Therefore, the objection does not only apply to Pryor's proposal, but also to reliabilist theories of justification. In fact, it was originally formulated against reliability theories, but it was recently extended to Pryor's proposal too. In

these cases, since perception is a generally reliable faculty, if 'Here is a hand in front of me' is a belief based on one's own perception, then it is justified.

Now, suppose that a subject has an experience with a representational and phenomenal content as of a red wall in front of her. Given the liberal conception of the structure of justification, her experience immediately gives her a justification for 'Here is a red wall in front of me.' This also allows the subject to have a justification for 'I am not in front of a white wall lit up in red,' or for 'I am not having the illusion of a white wall lit up in red,' since the first proposition excludes the second and the third. Thus, she can rule out being a victim of a perceptual illusion simply based on her current experience. This, however, seems intuitively wrong. That is, it seems wrong that an experience can, as such, tell us about its actual cause and exclude other possible ones that—indistinguishably from our own point of view—could have given rise to it.

Usually, reliabilists argue that they are not the only ones who incur this objection and that the alternative to their proposal is skepticism, which is a worse evil than "easy knowledge." Pryor, in contrast, bites the bullet and insists on the fallible character of the justification he would obtain for the conclusion of the argument. Thus, he tries to mitigate a little the strength of the intuition, which tells us that following the procedure just described does not seem to produce a justification that excludes being victims of a perceptual illusion.

However, to insist on the fallible character of justification appears once again unconvincing. The point of the objection is not so much that our experience alone cannot exclude *with certainty* that we could be victims of a perceptual illusion in that case. Rather, the point is that experience alone does not even give us more reason to believe 'Here is a red wall' than 'Here is a white wall lit up in red.' To realize this, it is enough to note that the experience could be phenomenologically identical although produced under non-standard circumstances. Therefore, it is not clear how an experience can justify, on its own *though in a fallible way*, the proposition 'Here is a red wall,' without already relying on the collateral information that it is obtained through causal interaction with material objects and through the generally reliable operation of the senses.

At this point, one could admit for the sake of argument that, contrary to the internalism Pryor adheres to, one can have an immediate perceptual justification for (1) without assuming (3), if the experience is reliable.

However, two problems arise. The first is the classical *bootstrapping* problem. One can arrive at the (externalistically) justified belief that one's senses work properly and that there is an outside world, without any collateral information regarding the reliability of the senses, or on the existence of the external world. Indeed, one can simply do that based on experience, deduction,and induction, which are known to be reliable sources of knowledge. For, if as the theory predicts, one's occurrent experience gives the subject a justification (externalistically understood) for 'Here is a red wall,' which in turn can give her by deductive means a justification for 'My senses are working properly' and for 'Here is a physical object.' Repeating the procedure a number of times, a subject would then obtain inductively a justification (externalistically understood) for more general propositions like 'My senses are working properly,' 'There is an external world' and also, for example, 'I am not a victim of perceptual illusions' and 'I am not a victim of lucid and persistent dreams.' This, however, seems intuitively incorrect.

A reliabilist could insist that instead one would have an externalist justification for the general propositions just mentioned (or even that one would know them). However, this would raise the following problem: when one deals with skepticism, as we saw in Chapter 1, Section 1.2, and indeed with Moore's own understanding of the skeptical challenge in Section 6.2, one is not content merely with having a justification or knowledge, but one would also want to be able to claim them.

Let us consider the matter in more detail. In order to be able to claim a justification for (1), it also seems necessary to argue that one's experience is produced by interaction with a universe of physical objects, by means of the operation of senses functioning properly and in the absence of lucid and persistent dreams. Therefore, in order to be able to claim a justification for (1), one must assume (3), i.e. that one is in the favorable case in which there is an external world and that one is in a causal relation with it, or that one's senses function properly and that one is not the victim of lucid and persistent dreams. In turn, these assumptions should be independently justified and be justified in such a way that one is able to claim one's justification for them. Now, according to a reliabilist, given *bootstrapping*, these beliefs would be externalistically justified. Still, it does not seem that a subject could rationally *claim* such justification. This justification would in fact exist, but for reasons external to the individual's reflection alone, and its obtaining wouldn't be known to her (Stroud 1994; Wright 2007).

Finally, Pryor's reconstruction of the skeptical position deserves more careful consideration. For him, skepticism is "a disease" that we shouldn't catch and from which we can cure ourselves by realizing that the correct account of the structure of empirical justification is the liberal one. Moreover, according to Pryor, the skeptic, who doubts (3), believes (most likely) that it is false that there is an external world.

Let us start with this last claim. In this regard, Pryor is clearly misinterpreting the skeptical position. As we saw in Chapter 1, Section 1.2, the skeptic is not an idealist and does not support an ontological position, but an epistemic one. Thus, a skeptic does not claim that it is (probably) false that there are physical objects. Rather, she believes that it is neither justified to believe that there are, nor to believe that there are not. Moreover, a skeptic maintains her agnosticism by means of *arguments*. Of course, there may be something amiss with them, but the whole point of the inquiry is to explain why this is so. A skeptic's position, therefore, cannot be simply reduced to a disease one should not catch.

Surely it may be the case that skeptical paradoxes can be solved and that once they are solved, they appear as "diseases" of the intellect—Wittgenstein would have said "bumps"—that we have caused ourselves by misunderstanding some crucial aspect of the issue (and indeed this is partly also Wittgenstein's position in *On Certainty*, see Chapter 7, Section 7.3). However, it is difficult to cure a disease if the causes are not removed. Moreover, since it is a disease of the intellect, we would need to remove them in a way that can be convincing to the sick person herself. Therefore, simply to reply to the skeptic that despite their subjective indistinguishability from dreams or hallucinations, the experiences we have give us an immediate justification for the corresponding empirical beliefs, does not seem sufficient to achieve this latter purpose and, therefore to really cure us of the intellectual disease that allegedly afflicts us, at least when we are considering skeptical hypotheses.

6.4 Wright's conservative account of perceptual justification and entitlements

Let us now turn to Wright's "conservative" or "skeptical" account of perceptual justification. Wright's account is called "conservative" or even "skeptical" because it agrees with a skeptic that, for a perception-based belief about specific objects in our surroundings to be justified, the general

presupposition that there is an external world needs to be justified. Yet, contrary to skepticism, Wright does not deny that such a presupposition can be justified—albeit non-evidentially, as we will see. Even less does he deny that our perception-based beliefs about specific objects in our surroundings are justified or even known.

Thus, according to Wright, Moore's proof is ironically an exemplification of the Humean skeptical paradox. For, in order to have a justification for 'Here is a hand,' one needs prior justification to believe that there is an external world. Hence, the proof fails to transmit justification from the premises to the conclusion (see Chapter 3, Section 3.2).

The first question that can be asked is why in order to have a perceptual justification for 'Here's a hand' when one has an experience with that given content, one must make collateral assumptions that should also be justified. In favor of this objection, one can also appeal to the phenomenological fact that when we acquire perceptual justifications for specific empirical beliefs about material objects around us, none of us formulate the assumption 'The external world exists.' Furthermore, none of us would ever be able to produce a justification for it. To reinforce the point, one could add that children too have beliefs about the objects around them, but they have neither the conceptual repertoire to formulate that assumption nor any capacity to justify it.

This objection can be answered by noting that for a long time it was thought that only beliefs could serve as a justification for other beliefs and that the structure of justification should be inferential also for empirical beliefs based on perception. According to this model, our justification for believing that there is a hand here would be given by the *belief* 'I have an experience as of a hand here in front of me.' However, a subject could have the same experience even if there really wasn't a hand in front of her and she was just dreaming it. That belief could only serve as justification for 'There is a hand in front of me' if she already assumes in a justified way that there is an external world, correctly represented by her experience. However, since this is what the argument should prove, the proof is circular and cannot serve to give us a (first) justification for believing that there is an external world. More schematically, according to the model of justification considered here, we should think of the structure of justification as given by the following inference in which (1)–(5) are all the contents of *beliefs*:

(1) I have an experience as of a hand here in front of me.

(2) The external world exists.

---

(3) There is a hand here in front of me.

(4) If there is a hand here in front of me, there is an external world.

---

(5) There is an external world.

Note that the argument thus reconstructed is not deductively valid since it could be that both (1) and (2) are true, but (3) could be false. However, (2) is necessary for (1) to be considered at least a *defeasible justification* of (3). Yet, this is entirely in line with Hume's argument (Wright 2004, 169–170). Therefore, the logical non-validity of the argument does not pose any problem to Hume's skepticism, which relies only on the obvious circularity of the inference and it shows that we cannot even have a defeasible justification for (5), let alone a conclusive justification.

This account of the justification of our empirical beliefs has come under attack in recent years. It is nowadays widely held that perception can (at least partly) justify beliefs based on it. For instance, as we saw in Chapter 5, epistemological disjunctivists hold that, but one need not necessarily endorse their account of perception in order to hold that view (on this, see Peacocke 2003). Yet, even so, it should be noted, concerning the phenomenological and child-related arguments in support of the objection we are analyzing, that here we are considering so-called *propositional justifications*. Propositional justifications are those which exist in favor of a certain proposition, such as 'Here is a hand' or 'The external world exists,' in the 'abstract space of reasons' regardless of whether they are ever considered by subjects. Moreover, they are independent of whether there are subjects who entertain those propositions in their mind, namely on whether there are subjects who believe them.

Propositional justifications must be distinguished from *doxastic* ones, which are those propositional justifications that are attached (or not) to the *beliefs* actually entertained by a subject, when they are believed on the basis of those very justifications. Doxastic justifications, in turn, depend on the existence of propositional ones, since they are propositional justifications that attach (or fail to attach) to our beliefs. Thus, since we are philosophers who are considering certain positions, we are also dealing with doxastic justifications. That is, we are wondering whether our beliefs are justified propositionally.

Finally, propositional and doxastic justifications are to be distinguished from those propositional justifications that are *rationally available to subjects*. Let us suppose that there is a subject who believes that there is a hand in front of her based on her sensory experience and let us also admit that her belief enjoys a propositional justification. However, let us imagine that she has the collateral belief that she often suffers from very realistic hallucinations. Given this collateral belief, she may find it impossible to make rational use of the propositional justification there is for 'Here is a hand.'

Therefore, the Humean skeptical paradox concerns the existence or non-existence of propositional justifications for 'The external world exists.' Hence, even if we do not typically entertain such collateral assumptions, and do not require subjects (to be able) to justify them, this does not show that in the 'abstract space of reasons' perceptual justifications do not depend on whether we have to assume in a justified way that there is an external world.

Furthermore, the idea that skeptical paradoxes concern second-order justification and thus represent a "crisis of intellectual conscience" (Wright 2004, 167, 209–211) is the basis of the anti-dogmatic response to the Humean skeptical paradox proposed by Crispin Wright. As we saw in Section 6.2, one might concede that in a good case scenario we have a justification for (1) and that that justification transmits to (3). Yet, just by having a visual experience as of a hand one could not thereby claim to possess such a justification for (1). For, as we saw while discussing Pryor's position in Section 6.3, that experience is compatible with its being produced in deviant causal conditions. To be able to claim one's justification for (1), then, one should already assume (3). Hence, the proof, even if successful at first-order (and this is far from being obvious), it would not allow one to vindicate one's justification for (3).

Another key element in Wright's strategy consists in noticing that at the root of any skeptical paradox, whether Cartesian or Humean, there is a "gap," i.e. an unmotivated assumption that the only kind of justification for very general assumptions like 'I'm not dreaming right now' and 'The external world exists' must be *acquired* through experience, or through reasoning. In the case of the Cartesian paradox, this means that we could never know whether the test to verify if we are dreaming at this moment has been performed and this makes it impossible to have a justification to believe that there is really the external world. In the case of the Humean paradox, in contrast, this assumption means that there is always a need for

the justified collateral assumption that there is an external world, which makes any attempt to prove it with a Moorean argument circular.

Furthermore, Wright crafts his proposal as a development of some of Wittgenstein's ideas in *On Certainty* (OC 1969, §§342, 151 in particular, but see also §§56, 82, 318, 628), in which Wittgenstein says that it belongs to the "logic" and "method" of our inquiries that certain propositions are not doubted. Here we will not address the exegetical issue of whether Wright's reading of Wittgenstein is plausible. We can simply point out that it is debated and questionable and that there are alternatives, as we will see in Chapter 7.

Wright, therefore, argues that it is possible to set up a *unified response* to the two paradoxes (Wright 2004, 174–175), aimed at claiming possession of a non-evidential justification—an *entitlement*, in his terminology—for the relevant assumptions. An entitlement is acquired neither by means of experience nor through reasoning, in his view. According to Wright, in fact, it is certainly not by mere reflection on the concept of a dream, or of the external world, or of a hand (or of any other specific physical object), that we can acquire a justification—in this case *a priori*—for believing that we are not dreaming right now, or for believing that there is an external world or a hand in front of us. Nor can we acquire such justification by empirical means, resting on the deliverances of our senses and arguments such as Moore's proof, if the Cartesian skeptic and the Humean skeptic are right. Therefore, the idea is that we are justified by default—whether we recognize it or not—to believe that we are not dreaming or that there is an external world. More precisely, according to Wright, there is a generic notion of warrant—that is, of epistemic support—that comes in two species: either as evidential support—*a priori* or *a posteriori*—for our beliefs; or else, as an entitlement—that is, a non-evidential support for them. Thus, we are warranted by entitlement, or entitled for short, to accept the relevant assumptions.

A further specification is needed, though. For Wright believes that we have an entitlement to *accept* 'I'm not dreaming right now' and 'There is an external world.' According to him, an acceptance differs from a belief in that only the latter depends on evidence—empirical or *a priori*—and gives rise to either manifest behavior or other beliefs. Instead, the former is not based on any kind of evidence, although it gives rise to the same kind of consequences as beliefs. That is, acceptances have the same output conditions as beliefs, while they differ from the latter in their input conditions.

Thus, according to Wright, we are entitled to *accept* certain propositions by default, but not to believe them.[1]

Before delving deeper into Wright's strategy let us consider a possible objection. One might in fact say that his is no better than those of Dretske and Nozick and the contextualists, because it does not seem to offer a direct response to skepticism, but it is merely a damage limitation strategy. Indeed, the skeptical paradox shows that we cannot acquire a *justification* for *believing* that we are not dreaming right now, or for believing that there is an external world, in a way conducive to the possibility of knowledge. Indeed, Wright concedes that this is a legitimate conclusion.

Yet, according to Wright, once his diagnosis of the skeptical paradox (in either of its forms) is accepted, and it is recognized that it rests on an unduly restrictive conception of justification, this would allow us to solve it. Hence, if correct, Wright's strategy would be an undercutting one, for it would show that the skeptical paradox is based on the *identification* of justification (or warrant) *tout court* with the sole evidential justification (or warrant). Therefore, Wright's response would not just be a damage limitation anti-skeptical strategy, but it would show that the skeptical paradox depends on a too narrow and apparently unmotivated conception of *justification* (or warrant). Thus, it would be a typically 'skeptical,' yet diagnostically illuminating response to the skeptical paradox. For, to reiterate, if successful, it would demonstrate that that paradox is based on a *partial* conception of justification (or warrant). Particularly, in the case of the Cartesian skeptical paradox, the first premise—that we have no justification for 'I'm not dreaming right now'—would be false, if the notion of justification (or warrant) included entitlements in addition to evidential justifications (or warrants). While, in the case of the Humean skeptical paradox, one would have an entitlement for 'The external world exists,' which, in turn, would be necessary to have a perceptual justification for 'Here is a hand.'

However, according to Wright, in order to *claim* our entitlement for 'I'm not dreaming right now' and 'There is an external world,' we need two different arguments. Let us analyze them in that order.

Let us consider a cognitive project such that its failure would be no worse than the costs of not executing it and its success would be better. Let *p* be an assumption of a cognitive project such that doubting *p* would mean doubting the importance, or the possibility of performing that project competently. Let us also assume that (i) we have no sufficient reason to

believe that *p* is false; and that (ii) any attempt to justify *p* would involve other assumptions that are no more certain than *p*, so that if one wanted to justify *p*, one would implicitly commit oneself to an infinite regress of cognitive projects.

Suppose that we need to determine by sight (and counting) how many books there are on a shelf. Let us therefore assume that *p* is 'I am not dreaming right now.' If *p* were revoked into doubt, we could no longer think that any cognitive project like determining the number of books on the shelf by sight (and counting) could be executed competently. However, precisely because if we were dreaming, then we wouldn't know it, we have no reason to believe that *p* is false either. Moreover, any attempt to justify *p* would involve assumptions:, for example, that our senses work properly and that there is an external world that is correctly represented by our experience. Yet, if we wanted to justify these assumptions, we would have to be able to justify *p*, and so we would find ourselves caught up in a regress of cognitive projects.

According to Wright, we must therefore conclude that 'I'm not dreaming right now' is an assumption of all our cognitive projects (empirical and ratiocinative) for which we have an entitlement. Therefore, it is rational to accept that we are not dreaming right now, even though we cannot in fact acquire any evidential justification for it. In fact, and this is the main connection with Wittgenstein's views in Wright's proposal, such an assumption would play a hinge role with respect to our cognitive projects.

Still, if the Cartesian skeptical paradox is blocked, it does not mean that the Humean one is. According to Wright, in fact, this first argument only allows us to claim the entitlement for 'I am not dreaming right now' and, therefore the rational legitimacy of the use of our perceptual and ratiocinative *faculties*, not of our acceptance that there is a world populated by physical objects that such faculties would enable us to know. To this end we need another argument.

By Wright's own admission, this is a much sketchier argument than the one we have just reviewed. It can be summarized as follows: (1) in order to have an objective conception of experience, one must think of it as due to interaction with objects that exist independently of it. (2) We have an objective conception of experience. (3) Therefore, we must have an objective conception of objects that exist independently of our experience.

Wright is perfectly aware that this does not actually mandate our ontology, but only makes it admissible. However, he believes that this is the

best we can do to claim our entitlement for 'There is an external world.' The argument is typically a transcendental one, which does not establish the truth of the conclusion, but merely shows that, if we stick to a certain conception of the objectivity of experience, we would need to presuppose the existence of mind-independent physical objects.

To see this, consider that if someone adhered to a different conception of the objectivity of experience, according to which, for experience to be objective, it must simply be uniform across time and perceptual modalities, say, it would not follow that, in order to have it, we need to presuppose the existence of mind-independent physical objects.

Let us now consider the more general question of the epistemic status of entitlements. It has been argued that entitlements would not be epistemic justifications (or warrants) corroborating the truth of the propositions in question, but *pragmatic* ones (Pritchard 2005a; Jenkins 2007; Williams 2012). In fact, they would give us no reason to believe that it is *true* that we are not dreaming or that there is an external world, but only good pragmatic reasons to act as if things were like that, since this would allow us to keep and pursue those 'cognitive projects' that are useful and important for us.

Wright (2014, 214) has responded by saying that entitlements are "warrant(s) to accept a proposition as true." Yet no argument is given to that effect, save one concerned with the kind of attitude involved, namely, an attitude of trust, which entails trusting in the truth of what is being relied upon. This obviously is no argument for the truth of the proposition trusted on, though. For it merely makes explicit the meaning of "trust"—that is, that to trust that *p* means to hold that *p* is true, even while lacking reasons in favor of *p*, or indeed while having evidence against *p*'s truth. Or, if you will, it makes explicit the norm of trust. That is, what makes trust correct, when complied with, or incorrect otherwise. For our trust in *p* will be correct if *p* is true, incorrect otherwise. Yet, this is not a good response to the challenge of providing epistemic warrants—albeit non-evidential ones—for the relevant hinges. For epistemic warrants are such that the proposition thereby warranted is, as such, more likely to be true than not. Merely noticing that we trust hinges and thus accept them as true, does not make their actual truth any more likely, though. Thus, the take-home message is that entitlements are not genuinely epistemic warrants; for they do not speak to the likely truth of the proposition they are supposed to warrant.

If this is correct, Wright would have provided nothing that supports the epistemically good standing of the relevant assumptions—that is, 'I am not now dreaming' and 'there is an external world.' At most, by means of the two arguments we have just reviewed, he would have provided *a priori* justification for altogether different propositions—that is, 'To accept that an external world exists (or that one is not dreaming) is rational.' Notice, moreover, that if the *a priori* arguments produced by Wright were in fact taken to support 'I am not now dreaming' and 'There is an external world,' respectively, then he would *ipso facto* have evidential, *a priori* reasons to *believe* (if not even know) them, and it is not clear why he should be content with merely accepting them.

At this point the question arises as to whether Wright's arguments have at least this secondary effect. That is, whether they allow us to claim that assuming certain unjustified (and unjustifiable) propositions, like 'There is an external world' or 'I am not dreaming right now,' is *epistemically* rational after all. The feeling is that even from this point of view, they are problematic. As already briefly pointed out, the first argument offered by Wright seems to show that, since forming empirical beliefs on the basis of our experiences is an important cognitive project for us—we could even say indispensable—we can't help but accept that we are not dreaming, etc. However, this is a clearly a pragmatic justification and a skeptic would have no reason to object to the fact that it is useful, or even essential for us to assume those propositions, in order to go about our ordinary epistemic practices.

The second argument tells us that in order to have the conception of the objectivity of experience that we have, we cannot but accept that there are physical objects. We have already pointed out how such a conception of the objectivity of experience may be contested. Yet, the main point is that a skeptic could happily agree that for us to have such a conception of the objectivity of experience, we cannot but accept that there is an external world. Yet, she might argue that this does not exclude that it is metaphysically possible that the descriptions for which we use it are systematically incorrect.

One could appeal to semantic externalism at this point, and say that for our thoughts (and their linguistic counterparts) to have a certain content, the physical objects these thoughts are about must exist. (This argument has been supported in various guises by authors such as Putnam (1981), but also by Wittgenstein (OC, §369), see Chapter 7, Section 7.3).

Yet, as we saw in Chapter 2, Section 2.2, not only is semantic externalism a thesis that cannot be taken for granted, but even if it were correct, it would only require that we have been, or at least sometimes still are, causally related to physical objects to ensure that our thoughts (and assertions) have certain contents. However, this does not mean that we are able to *know* whether we are still or currently interacting with such objects. In other words, the semantic thesis—granted that it is correct—implies a certain ontology but has no effect on the *epistemic* problem raised by skepticism.

In conclusion, it seems that Wright's strategy does not solve the skeptical paradox, because it does not provide a convincing diagnosis of the mistaken conception of justification (or warrant) on which that paradox would allegedly rest. In fact, it does not provide justifications (or warrants) that corroborate the truth of propositions such as 'I am not dreaming' and 'The external world exists,' but only *a priori* justifications for believing that accepting those propositions, which as such remain unjustified, is rational. Finally, on closer inspection, the reasons it provides for believing that accepting these propositions would be rational are at most pragmatic, and hence completely compatible with the skeptical position.

6.5 Between liberals and conservatives: Wedgwood and Sosa's *a priori* strategies

As expected, the dispute between liberals *à la* Pryor and conservatives *à la* Wright has recently given rise to alternative proposals, which seek to support a kind of middle ground between these two perspectives. These responses are still in progress and therefore difficult to fully assess, but nonetheless deserve to be mentioned. However, readers who are not particularly familiar with the debate can—if they wish—leave out this complex section.

The common feature of these strategies is that they try—albeit in different ways—to give an *a priori* justification for very general assumptions such as 'There is an external world,' 'We are not victims of a lucid and sustained dreams,' and 'Our sense organs work mostly reliably.' In a sense, they could be seen as contemporary developments of the reflection on what Reid called "*first principles, principles of common sense, common notions,* [or] *self-evident truths*" ([1785] 2002, 6.4, 452), which are remarkably similar to what several hinge epistemologists, including Wright, would consider hinges properly so-regarded (see Chapters 4 and 7).

According to Reid, common sense first principles are identifiable because they have certain additional traits. For instance, their denial is not only false but absurd (ibid.: 6.4, 462); they are "necessary to all men for their being and preservation, and therefore [they are] unconditionally given to all men by the Author of Nature" (ibid., 4.6, 412); and, finally, they are "no sooner understood than they are believed. The judgment follows the apprehension of them necessarily, and both are equally the work of nature, and the result of our original powers" (ibid., 6.4, 452). The ultimate nature of these first principles was not clear in Reid, for, on the one hand, he claimed that they were instilled in us by God, and therefore would not result from empirical generalizations, or further empirical inquiries. On the other hand, he also claimed that an empirical study of what people believe across cultures, religions, and mores would reveal them. More specifically, he claimed that these principles are psychological generalizations about belief formation applicable to most of our species (ibid., 6.4, 464–465). Furthermore, as we mentioned in Section 6.1, for Reid, these general principles provide evidence for particular beliefs. Thus, the principle "That those things do really exist which we distinctly perceive by our senses, and are what we perceive them to be" (ibid., 6.5, 476) issues in the self-evident belief, for example, that I perceive a computer before me (Van Cleve 1999). As a result, Reid's philosophical method accords with common sense, as we saw, insofar as ordinary perceptual beliefs are evident (or even self-evident).

Other *a priori* strategies in favor of these "first principles" are nowadays present in the epistemological literature. An interesting family of those, which could be brought to bear onto the liberal-conservative dispute about the structure of perceptual justification, even though it preceded that dispute and developed independently of it, is the abductivist one (Russell 1912; 1914; Goldman 1988; Vogel 1990a; 2005; BonJour 2003). Abductivist responses to skepticism come in various fashions, yet their common central claim is that "while both commonsense and skeptical explanations of our sensory experiences are equally consistent with the phenomenological data … there are explanatory reasons for favoring commonsense explanations over skeptical ones" (Beebe 2009, 605). These explanatory reasons vary, ranging from ontological, explanatory, and psychological simplicity to explanatory breadth and depth, as well as coherence with background knowledge, intrinsic plausibility, avoidance of ad hoc elements, fecundity, and more (see ibid., 609–611). In particular, *a*

priori abductivists,[2] such as BonJour (2003, 88), claim that we can be *a priori* justified in believing that certain features of our experiences—such as the continuity among experiences, their coordination across sensory modalities, their regularity, etc.—are caused by the interaction with a world populated by physical objects because such a commonsensical hypothesis provides a better explanation of those features of our experience. Thus, we are *a priori* justified in holding onto the hinge that there is an external world and, consequently, in disbelieving radically skeptical hypotheses. Borrowing here from Beebe's (2009, 624) rendition of BonJour's position, abductivists hold that we can be *a priori* justified in believing that:

(a) the real-world hypothesis (RWH) is the best explanation of our sensory experience;

(b) RWH satisfies the explanatory criteria better than its skeptical rivals;

(c) explanatory criteria such as simplicity and intrinsic plausibility are correct;

(d) the explanatory criteria are truth-linked;

(e) other things being equal, it is justifiable to infer that an explanation is true, on the basis of the fact that it is the best explanation of the relevant phenomena.

Consequently, according to BonJour, there is an abductive *a priori* justification for believing RWH and for disbelieving radically skeptical hypotheses. The detailed defense of BonJour's position hinges on complex issues in probability calculus and on the endorsement of a form of modal frequentism that comprises infinite sets of possible worlds, the ordering of which is unknown. This latter aspect of these sets of possible worlds is problematical, however, for it compromises the idea that the RWH be actually better—that is, more probable—than alternative skeptical ones in accounting for the features of our experience. (For a detailed discussion of BonJour's proposal, the objections to it, and a possible reply, see Beebe 2009).

Be that as it may, the contemporary supporters of "first principles" we will presently consider claim that such principles are *a priori* justified but do not offer an abductivist defense of the *a priori* justification of anti-skeptical hypotheses. This, however, we shall see, also involves a substantial and not

obvious revision of the very notion of *a priori*. In any case, it is clear that, on the one hand, these strategies embrace Wright's idea that these assumptions must be justified in order to be rational and, on the other, they reject Wright's entitlement strategy and set themselves the more ambitious goal of giving *a priori* justifications for believing that there is an external world and that one is not the victim of a lucid and persistent dream.

According to Wedgwood (2013), it can be shown with semi-formal tools derived from probability calculations (thus, following White 2006), that an experience with representational and phenomenal content of a hand alone cannot justify 'I am not dreaming of seeing a hand in front of me now.' As such, in fact, that experience would also raise the probability of the contrary specific skeptical hypothesis (cf. Section 6.3). According to Wedgwood, contrary to the liberal position, this shows that the existence of a justification to exclude being in the specific skeptical scenario compatible with the experience one is having at that moment is a *necessary consequence* of having a perceptual justification for 'Here is a hand.' However, contrary to what the conservative maintains, for Wedgwood, this justification *does not explain* why the experience one is having is a justification for believing 'Here is a hand,' For, precisely because as the liberal claims, one is immediately justified in believing propositions such as 'Here is a hand' based on one's own experience.

According to Wedgwood, the kind of *a priori* justification we have for these propositions depends on a type of *bootstrapping* argument, which leverages on being justified to form empirical beliefs directly on the basis of one's own experience, as long as there is no reason to believe otherwise. However, it is a *bootstrapping* argument that is different from the ones we saw in Section 6.3, because it is *a priori*.

The argument is quite complex, and we are not going to see it in detail. Suffice it to say that, for Wedgwood, the practice of taking experiences at their face value to form the corresponding beliefs in the absence of reasons to the contrary is rational. Moreover, the practice of forming beliefs by deductive means and through inferences to the best explanation, as long as one starts from rationally believed premises, is rational too. Therefore, it is rational to form the belief that I have a hand if I have an experience with that content and there are no contrary reasons. Moreover, I can rationally believe many other specific empirical propositions based on appropriate perceptual experiences. Through inference to the best explanation, I can then rationally conclude that my experiences are mostly reliable, which is

incompatible with skeptical hypotheses such as 'I am the victim of a lucid and persistent dream' and 'There is no external world.' Therefore, this process of rational reasoning, called "*a priori bootstrapping*" can lead any subject with the relevant concepts and cognitive abilities to believe the denial of skeptical hypotheses and therefore it explains why we have *a priori* justification for doing so. Finally, it is important to note that, if Wedgwood is right, then there would be *a priori* justifications for contingent empirical propositions, such as 'I am not dreaming right now,' which, indeed, would be *a priori* justified even if they were false, were one in a skeptical scenario.

As it has been said, this argument rests on the assumption that the practice of taking experiences at their face value in the absence of reasons to the contrary to form the corresponding empirical beliefs is rational. Wedgwood explicitly admits that he has done nothing to prove the legitimacy of this assumption. Clearly a Humean skeptic will deny it: from his point of view, if one does not already have a justification to exclude being in a (however specific) skeptical scenario, one cannot rationally take one's experience at face value to form the corresponding empirical beliefs.

Elsewhere, Wedgwood (2011) gives an argument to establish the rationality of empirical beliefs formed by taking one's experience at its face value in the absence of reasons to the contrary. His argument essentially relies on the fact that in order to have certain concepts, it is necessary to form beliefs that involve the use of those concepts immediately based on experience with a corresponding representational content. (A similar argument can be found in Peacocke 2003.) It is not important to explain this in detail here, or to dwell on the—questionable—fact that the absence of collateral beliefs is deemed enough to call such practices of belief formation '*a priori*.' After all, they are still epistemologically based on experience. Rather, what should be noted is that even if it is true that at least sometimes those practices must have led us to true beliefs, this does not prevent us—as Wedgwood himself allows– from relying on them even in skeptical scenarios where they would not produce true beliefs. Therefore, it is difficult to see in what sense they can be used to give a convincing answer to the skeptical paradox.

There is also a further problem that has to do with the stability of Wedgwood's proposal. Namely, the perplexity that it might fall on one or the other of the positions from which he wants to distance himself. In fact, if there is an immediate justification for 'Here is a hand,' given one's

sensory experience and the absence of reasons to the contrary, why shouldn't this also give one a justification for the denial of skeptical hypotheses? What is wrong with the inference from 'Here is a hand'—which hypothetically is a premise immediately justified by one's own experience in those conditions—to 'There is an external world' or 'I am not dreaming right now'? It is not clear, since Wedgwood does not deny the Closure principle, and does not commit to the fact that the Moorean argument is an example of failure of the Transmission of justification principle. Hence, Wedgwood's intermediate position seems to collapse into Pryor's liberal thesis. Yet, Wedgwood could fall back onto Wright's conservative position too. This would happen if he were to argue that experience alone cannot give one a justification for 'Here is a hand,' because this *ipso facto* would give one a justification to rule out being in a skeptical scenario. And, therefore, if he were to argue that having a perceptual justification for that proposition has the necessary consequence that there must be a justification for 'There is an external world' or 'I am not dreaming right now.'

The fact that it might seem otherwise depends on a confusion between doxastic and rationally available justifications, on the one hand, and propositional justifications, on the other, probably caused by the "robust" internalism subscribed to by Wedgwood. For him, every time we talk about justifications, we are in fact talking about justifications actually possessed by subjects. Yet, as we saw in Section 6.4, the difference between Pryor and Wright is measured in terms of *propositional* justifications. In this perspective, if, like Wedgwood, one is convinced that experiences alone cannot justify the belief in the denial of skeptical hypotheses, but that for this purpose there must be independent justifications, then one is already a conservative. It does not matter if, from the point of view of the justifications available to the subjects these justifications are acquired only after having exercised the ability to form empirical beliefs on the basis of one's own experiences and the cognitive and conceptual capacities necessary to reflect on this practice. Finally, once it is admitted that experiences cannot justify belief in anti-skeptical hypotheses, it is only from the point of view of the subject's available justifications that one can continue to think—somehow in line with the dogmatist—that these directly justify the subject to form the corresponding empirical beliefs. It is only from this point of view, in fact, that it is plausible that subjects do not already have to have a justification for believing in the denial of skeptical hypotheses in order to form rational empirical beliefs on the basis of their own

experience, in the absence of reasons to the contrary. Yet, to repeat, the dispute between Wright and Pryor concerns the nature of propositional justifications and, as things stand, Wedgwood's position is unstable. In fact, it falls now on Pryor's liberal position, now on Wright's conservative one, and does not seem to offer a plausible alternative to either.

Moreover, we have seen that, for Wedgwood, there are contingent empirical propositions such as 'I am not dreaming right now' that are justified *a priori* and would continue to be so even if they were false if, for example, one were in a skeptical scenario.

Now, this is not the place for a thorough discussion of the notion of the *a priori*. However, it should be noted that, if this were the case, it would show how the arguments aimed at giving an *a priori* justification for these propositions would in no way corroborate their truth. It would therefore have the paradoxical consequence that having an *a priori* justification for these propositions would be completely compatible with the truth of the skeptical hypothesis. Hence, it would be doubtful that this strategy could provide a solution to the skeptical paradox.

Another proposal which argues that propositions like 'There is the external world' are *known a priori* (and are therefore true), can be found in Ernest Sosa (2013). According to him, we can explain how we know these propositions by appealing to the notion of *rational competence*, although we are not yet able to explain in detail how this competence works. Here is this account.

A subject S has a rationally competent human *intuition* that P if and only if: (i) understanding (the proposition that) P suffices for her to accept P, independently of any other reason she might have to accept P besides understanding it; (ii) the mechanism that makes her accept P, on the basis of the understanding of P, independently from any other rational basis, is an epistemic competence; and (iii) typically, this mechanism operates in human beings in the habitual course of their development, since it gives subjects the constitutive concepts of P (or the concepts necessary to grasp that proposition) together with the propensity to accept P.

Now, the relevant point for our purposes is that such a rational competence would give rise to *intuitive* knowledge of P—that is, knowledge that is not based on experience. In fact, Sosa believes that having knowledge of specific empirical propositions, like 'Here is a hand,' on the basis of one's own experience, depends in turn on having knowledge of propositions such as 'There is an external world.' This latter piece of knowledge,

however, cannot in turn be obtained through experience. In this sense the knowledge that we would have of such a proposition would be *a priori*, although it would depend on the functioning of our epistemic competence, the nature and functioning of which we do not know (yet).

At this point, we should ask ourselves what this proposal adds to the classical *a priori* idea that we could have knowledge of 'There is an external world' simply through reflection on the concepts involved in the grasp of that proposition. It also raises the question of how it can face the objection that as intuitive as it seems to us to accept 'There is an external world,' we could still find ourselves in a skeptical scenario (unbeknownst to us). In other words, although accepting 'There is an external world' is natural for us, given that we have the concepts at stake, such a spontaneous assent does not seem to be enough to determine that it is true that there is an external world around us and therefore to determine that we know there is.

Obviously, one could say that if a spontaneous assent is the effect of the exercise of a real epistemic competence, this guarantees that it is about a true proposition. However, at this point it would be necessary to say something more about the nature and workings of such a competence to ensure that it 'keeps track' of the truth.

Furthermore, it can be imagined that this competence may have been activated in such conditions as to give rise only to true beliefs (in effect, knowledge), but that subsequently, having been placed into a skeptical scenario, it continues to operate while no longer keeping track of the truth of the propositions believed by means of it and without thereby providing us with knowledge of them. Therefore, we would have reached the counterintuitive conclusion that we could have *a priori* justified beliefs that are false. Once again, this casts doubts on the anti-skeptical effectiveness of the *a priori* strategy.

Sosa (2021) presents a very different strategy in support of hinges—as he himself now is happy to call them—such that there is an external world, or that our sense organs work mostly reliably. In Chapter 2, Section 2.3, the basics of Sosa's latest views regarding knowledge were introduced. Here we will briefly consider what he says about the status of hinges in his more recent work. First off, he claims that hinges are a necessary component of knowledge, understood in virtue-theoretic terms (see Chapter 2, Section 2.3). Second, like Coliva (2015, see Chapter 7, Section 7.5), he thinks that no justification of hinges is needed. Subjects who judge, based

on their current perceptions, that there is a hand in front of them, are tacitly deploying their sensitivity to such hinges and there is no need for them to possess, even less provide a justification for them. Third, the mere possibility that these background assumptions are mistaken does not impugn the possibility of knowing ordinary empirical propositions like 'Here is my hand,' thanks to an apt performance of one's perceptual and cognitive functions on a given occasion (Sosa 2021, 149). Thus, as Sosa himself recognizes (ibid., 160ff.), thanks to default assumptions, a subject can be credited with reflective knowledge, and not just with animal knowledge, when she operates based on them. Yet their knowledge may not be maximally secure in contexts in which a subject could have easily been mistaken about the obtaining of default assumptions. Hence, writes Sosa (ibid., 150):

> Conscientiously enough, without negligence or recklessness, we normally assume ourselves free of skeptical scenarios. And this assumption is proper even on the rare occasions when it is true but not known to be true and even quite unsafe.

Still, the classical problem of "discharging the antecedent," which besets epistemically externalist positions, remains. That is, if it is the case that we are not in a skeptical scenario, then Sosa is right that we can proceed on that basis, by assuming "ourselves free of skeptical scenarios." Yet, a skeptic would likely object that unless we justify that assumption, it is arbitrary and not epistemically safer than a skeptical one and nothing of what has been said so far provides such a justification.

To counter such an objection, Sosa (ibid., 169) claims the following:

> Radical victims [that is, victims of radically skeptical scenarios] are unable to make genuine judgments. To me this seems plausibly connected to the metaphysical detachment from society that comes along with their detachment from their surrounding world. Human beings and human knowledge plausibly require the metaphysical connection with human community and an external world that radical scenarios would tear away from them.

Notice that this is very reminiscent of Wittgenstein's claim, echoed by Putnam, that in a dream or a BIV scenario one cannot make assertions or

even mean what words normally mean. Skeptical scenarios would just be illusions of doubt, rather than possibilities worth taking seriously.

If correct, this claim would be highly therapeutic—i.e., there would no real problem posed by skeptical scenarios because if *per impossibile* we were in them, we could not really assert or mean anything at all, not even 'I may be dreaming/a BIV.' Hence, Sosa's strategy would be an undercutting one, since it would claim that the putative radical skeptical paradox is in fact not a paradox at all. Still, it is not clear that Sosa's claim be based on endorsing semantic externalism *à la* Putnam, or a use theory of meaning *à la* Wittgenstein. His claim, rather, concerns the possibility of making genuine *judgments*, if placed in a radically skeptical scenario. Yet, one may ask what exactly the connection is between being able to judge that *p*, and having a (causal) connection with the world and the human community. Compare with making an assertion. Now, assertions are speech acts in which a speaker typically intends to inform the audience that *p* is the case, while being appropriately (causally) connected with the world. By contrast, a judgment is typically something that goes on in *foro interno*. Sure, it is a judgment that *p*, but why should that entail currently having a connection with the world, let alone a human community? Moreover, as we saw in Chapter 2, Section 2.3, while it may be true that in order to possess the concepts necessary to entertain the judgment that *p* one needs to have been in appropriate causal relations with a physical world and a human community, that does not mean that in the immediate aftermath of a brain envattment one could not still deploy such concepts. Further qualifications about the necessary connection between judging that *p* and the non-obtaining of a radically skeptical scenario are therefore needed to bolster Sosa's more recent anti-skeptical proposal.

Further reading

On Hume, see Stroud (1977), Fogelin (1985), Garrett (1997), Owen (1999), Ainslie (2015), and De Pierris (2015). On Reid, see Daniels (1974, Chapter 4), de Bary (2002), Pappas (1989), Van Cleve (1999), and van Woudenberg (2018). On G. E. Moore's "Proof of an external world," skepticism, certainty, and common sense, see Coliva (2010a, Chapter 1), Clarke (1972), Malcolm (1942; 1949; 1977), Stroud (1984, Chapter 3), Wittgenstein (1969), Wright (1985), and Pritchard (2021a). Pryor has further defended his liberal account in Pryor (2000; 2012). For criticism

of his position, see Coliva (2015, Chapters 1, 2), White (2006), and Wright (2007; 2012). On Wright's views, see Coliva (2012; 2015, Chapters 1–3; 2020b), Williams (2012), and Zalabardo (2012). For Wright's replies, see Wright (2012). Further works on Wright's notion of entitlement are Pritchard (2014b), Smith (2020), Pedersen (2020), Elstein and Jenkins (2020), and Greenough (2020). On abductivist responses to skepticism, besides BonJour (2003), see Vogel (1990a; 2005), and Goldman (1988). Beebe (2009) is an excellent review of most abductivist responses to skepticism since Russell's (1912) defense of it up to the present day.

Notes

1 In fact, if belief were taken to span across evidentially sustained and unsustained belief, one could safely replace Wright's talk in terms of acceptance with belief, provided belief were taken to be of the evidentially unsupported variety.

2 *A priori* abductivists are better off than *a posteriori* ones, as far as the rebuttal of external world skepticism is concerned, because they do not appeal to our current empirical knowledge to claim that it is better explained by the hypothesis that there is an external world than by alternative skeptical ones. Appealing to our current empirical knowledge would clearly beg the question against external world skepticism.

7

VARIETIES OF HINGE EPISTEMOLOGY

Naturalism, contextualism, and constitutivism

7.1 Introduction

As we saw in Chapter 6, in the dispute between Wright and Pryor, it has been noted by several parties how it is problematic to maintain that in the absence of defeaters, experience can be an immediate justification for 'Here is a hand,' without assuming or taking for granted—no doubt mostly implicitly—that there is an external world. However, this does not imply *ipso facto*—contrary to what Wedgwood argues—that there should be an *a priori* justification for that proposition. Nor does it imply—contrary to what Wright believes—that in order for it to be rationally legitimate to believe (or accept) 'Here is a hand' on the basis of one's own experience, there must be a non-evidential justification—an entitlement—for it. In fact, beyond the

DOI: 10.4324/9780429057946-8

specific problems faced by entitlement and *a priori* strategies (analyzed in Chapter 6, Sections 6.4 and 6.5), all that follows from the criticism of the liberal position is just that in order to have an empirical justification based on one's own experience for ordinary empirical propositions like 'Here is a hand,' in the absence of defeaters, one has to make collateral assumptions and, especially, assume that there is an external world, that one's senses are functioning properly, and that one is not the victim of a lucid and persistent dream. Therefore, it remains to be explored more thoroughly whether the structure of empirical justification is neither the one proposed by liberals, nor the one put forward by conservatives, but is in fact the one supported by so-called 'moderates.' Accordingly, to have a justification for specific propositions such as 'Here is a hand,' in the absence of defeaters, one needs an appropriate experience, as well as a collateral assumption concerning one or more of the 'heavyweight assumptions' just mentioned. In this way, one's own occurrent experience can constitute a justification for 'Here is a hand,' rather than for its skeptical counterpart 'I am dreaming of seeing a hand,' while obviously remaining a fallible (or defeasible) justification.

On closer inspection, the moderate position is not new, even though it can be unprecedented to call it thus. Hume, Wittgenstein, Strawson, and pragmatists, such as Charles S. Peirce, can all be considered moderates regarding the nature of empirical justification. After all, they all believed that our specific empirical justifications are based not only on experience, but also on collateral assumptions which is natural for us to make—in some sense of 'natural'—or which are only pragmatically rational (see Boncompagni 2016).

More precisely, all moderates believe that in Moore's proof:

(1) Here is a hand.

(2) If there is a hand here, there is an external world.

(3) There is an external world.

Premise (1) is justified based on one's course of experience, absent defeaters, only if the conclusion (3) is assumed. Yet, contrary to Wright and a Humean skeptic, that conclusion can be assumed even in the absence of a justification (evidential or otherwise) for it. At the heart of the moderate account, therefore, lies an alternative way of conceiving of the structure of perceptual justification, with respect to both liberalism and conservativism.

Yet, moderates offer different accounts of the status of (3) itself. That is, granted that (3) is neither justified nor justifiable, they disagree over its provenance and overall rational status.

In this chapter, we will look at four prominent ways of developing this intuition and therefore of bringing hinge epistemology to bear on Humean skepticism in its connections with Cartesian skepticism, instead of merely on the latter, as in Pritchard's version of hinge epistemology (see Chapter 4). That is: (1) Peter Strawson's (1985) naturalism; (2) Wittgenstein's (1969) normativism; (3) Williams' (1991) contextualism; and (4) Coliva's (2015) constitutivism. All in all, these views are variations on Wittgenstein's account and belong to the growing trend of hinge epistemology. (See also Chapter 4 for Pritchard's variety of hinge epistemology and Chapter 6, Section 6.4 for Wright's.)

7.2 Strawson's naturalism

In an influential essay published in 1985, Sir Peter Strawson claimed—believing in this way to be exposing Wittgenstein's position in *On Certainty* (OC 1969)—that the Humean skeptical paradox cannot find a direct answer. Thus, we cannot provide a justification for 'There is an external world.' However, according to Strawson, not all doubts are worth taking seriously. In particular, those concerning beliefs that we find ourselves having by dint of nature would not be so.

Now, Hume had argued that 'There is an external world' is a belief we have because of certain psychological connections which, given our *cognitive structure*, are as much involuntary as inevitable for us. According to Strawson, Wittgenstein would have argued that the belief in the existence of the external world is obvious and inevitable for us because it is part of our *form of life* to consider it as such. That is, we are brought up within a community that takes the existence of physical objects for granted, and we have 'swallowed' that belief by being trained to take part in practices that presuppose the existence of mind-independent physical objects. Hence, according to Strawson's reading of *On Certainty*, by growing up in our community, we are trained to take for granted those beliefs that serve as the foundation of our language-games. Among these beliefs, there would be the belief in the existence of the external world, which plays an essential role in our language-games since it allows us to transcend "our cognitive

locality" (Wright 2004, 172–174)—that is, it allows us to take our sensory experiences to bear onto beliefs about mind-independent objects.

For Strawson, Wittgenstein proposed a form of naturalism that, with John McDowell (1994), we could call of "second nature." For it is the product of training and acculturation within the human community, as opposed to a kind of naturalism of "first nature" that is the result of our cognitive structures that are largely biologically determined, of which Hume was a main supporter.

The doubts raised by a skeptic would then be *unnatural*, for they go against what is natural—in either of these senses—for us to believe. Hence, they go against what we naturally believe and, although legitimate—that is, admissible from a logical and metaphysical point of view—they would be completely implausible and ineffective.

Strawson's position is problematic on several fronts. First, as a rendition of Wittgenstein's claims in *On Certainty* and, second, in its own right. We will look at an alternative account of Wittgenstein's views in Section 7.3. Here we will just confine ourselves to note that the naturalist response is quite impotent against skepticism. Indeed, as just remarked, it is the very outcome of Hume's skepticism. After all, once persuaded by philosophical arguments that no justification for 'There is an external world' would be forthcoming, Hume himself embraced the idea that it was simply part of our psychological constitution, and therefore of our nature, to uphold that belief, which would thereby be *arationally* held. Far from providing a solution or a dissolution of skepticism, insisting on the natural origin of that belief was a way of explaining why we find ourselves saddled with it, even if we have no epistemic reasons in its favor. If the goal were to trace the *rationally* secure foundations of our epistemic practices, then it would simply be an unattainable aim, because at the foundation there would be only assumptions that are not rationally well-founded.

Moreover, Hume himself insisted on the fact that nature is stronger than philosophy. Hence, even a philosopher, and indeed a skeptic, in virtue of being a human being, cannot help believing in the existence of an external world. Yet, from the fact that a belief is unavoidable for us it does not follow that it is true, let alone epistemically justified and therefore knowledge-apt, or at least rational. By considering Wittgenstein the heir to Humean naturalism, whose main contribution would have been to substitute "second-nature" to "first nature"—that is, our culturally determined nature to our psychologically or biologically determined one—Strawson not only missed the real import of

Wittgenstein's position, but he exposed it to the objection that naturalism is indeed a consequence, rather than a solution or a dissolution, of skepticism.

7.3 Wittgenstein's normativism

As mentioned in Section 7.2, Strawson's reading of *On Certainty* is problematic from an exegetical point of view. This is obviously a very complex issue that we cannot address here with the attention it deserves. However, it is useful to note that there is no agreement that in his last work Wittgenstein supported a form of naturalism, albeit *sui generis*. Furthermore, it should be considered that there is at least another reading of *On Certainty* that is much more accredited today—namely, the so-called "framework reading."[1]

Accordingly, Wittgenstein put forward a more radical position than the one Strawson attributes to him. Key to this reading is the claim that, for Wittgenstein, propositions such as 'There is an external world' would not be empirical propositions liable to be verified and controlled, but they would be *rules*, or, more tentatively, they would play a *rule-like role* (see OC, §95). More precisely, they would be norms of linguistic (and conceptual) representation and of evidential significance. That is, it is a rule of our language that the class of physical objects—namely, of objects that exist independently of the fact that they are perceived—is exemplified and therefore is not empty. Thus, we can say things like '*a* is a physical object'; '*a* has existed before human beings existed', or '*a* exists even when it is unperceived', etc. (OC, §35–37). Moreover, it is also a rule of our epistemic practice of acquiring perceptual evidence for or against ordinary empirical propositions, that the sensory experiences we have are taken to bear on the truth or falsity of beliefs about objects that exist independently of those same sensory experience. Skeptical doubts would therefore be not only unnatural, but completely *senseless*—that is, illegitimate—since it makes no sense to doubt a rule. That is, one can argue whether it is useful, or otherwise pragmatically appropriate, to have a given rule, but it certainly makes no sense to wonder whether a rule correctly represents a state of affairs.

More specifically, according to framework readings, doubting of these hinges means to call into question either the very meaning of our words, or the very possibility of acquiring perceptual evidence for or against specific empirical propositions. That is, if it is doubted that there is a hand where I seem to see it right now, then it becomes doubtful that I understand or assign a meaning to the word 'hand' (OC, §§369–370, 456).

For, according to the later Wittgenstein, meaning is determined not only by means of definitions, but also thanks to an agreement in judgments. Thus, if it cannot be agreed that there is indeed a hand where I seem to see it right now, absent defeaters, the very meaning of the word 'hand' is jeopardized. Yet, if its meaning is in jeopardy, then no sense can, literally, be made of the alleged doubt 'Maybe there is no hand here' (in what appear to be optimal perceptual conditions). Similarly, when it comes to 'there are physical objects,' if there is no agreement regarding the possibility for the category of physical objects to be exemplified correctly, then the very meaning of 'physical object' would be jeopardized as well. Therefore, even in this case, no sense could really be made of the alleged doubt 'Maybe there are no physical objects' (OC, §§35–37, 57).

Moreover, if it is doubted that there are physical objects, such as my hand here and now, then our sensory experience cannot be taken to bear on beliefs about mind-independent objects (OC, §§125, 446). Yet this is essential to the very practice of providing evidence for or against ordinary empirical propositions. If that practice is endangered, then allegedly skeptical doubts would turn out not to be supported by any reason. And, if that is the case, they would fail to be rational and would pose no real challenge to our ordinary empirical beliefs.

Of course, a skeptic could reply that the kind of doubt she raises is not like any ordinary empirical doubt. Rather, it is a mere logical and metaphysical possibility. Yet, Wittgenstein thought that philosophical doubts were illegitimate. For, in his opinion, philosophy is not a *sui generis* language game alongside ordinary ones. Rather, it is an illusion of a language game brought about by the conflation of the rules that govern ordinary ones. Thus, either a doubt is an ordinary one, and it must therefore be based on reasons, which in turn presuppose certain hinges, or else it is no doubt at all, but it is only an appearance or an illusion of doubt (see Coliva 2010a, Chapter 3; Schönbaumsfeld 2016). Indeed, with respect to Cartesian skepticism, Wittgenstein was adamant that it would be senseless, as the following quotes make clear:

> The argument "I may be dreaming" is senseless for this reason: if I am dreaming, this remark is being dreamed as well – and indeed it is also being dreamed that these words have any meaning. (OC, §383)

> "But even if in such cases I can't be mistaken, isn't it possible that I am drugged?" If I am and if the drug has taken away my consciousness, then I

> am not now really talking and thinking. I cannot seriously suppose that I am at this moment dreaming. Someone who, dreaming, says "I am dreaming", even if he speaks audibly in doing so, is no more right than if he said in his dream "it is raining", while it was in fact raining. Even if his dream were actually connected with the noise of the rain. (OC, §676)

In these passages, which are strikingly reminiscent of Putnam's idea that, if entertained by a brain-in-a-vat (BIV), words would not retain their ordinary meaning (see Chapter 2, Section 2.2), Wittgenstein puts forward the even more radical idea that they would not retain any meaning whatsoever. For he subscribed to the view that meaning can only be conferred on words by their ordinary use, with the criteria that are in force within it. In its ordinary use 'I may be dreaming' is employed to express one's incredulity *vis-à-vis* a given situation, or else it is used in the past tense—that is, in the form 'Maybe I was dreaming'—to indicate that a given memory may not be veridical. In this event, we also normally have various ways to ascertain whether we are remembering correctly.

Several objections can be raised against Wittgenstein's account. First, that it relies on a conception of linguistic meaning that can be and has been disputed. In the most general way, it may be objected that once words have their meaning, they carry it on their sleeves, as it were, no matter in which context they are used. Nowadays, many 'literalists' about linguistic meaning would endorse this view. Thus, once the word 'dream,' say, has been conferred a meaning, it retains it even if the contexts of its utterance are not the ordinary ones, and one is dreaming of uttering that word.[2]

Furthermore, one may notice that the word 'witch' did not lose its meaning simply because it became apparent that it could not be correctly predicated of anyone. Similarly, 'physical object' could still retain the meaning of 'mind-independent entity' even if it could never be applied to correctly describe a state of affairs, and 'hand' would still retain its meaning even if it were no longer taken to be exemplified by anything (on the assumption that 'hand' did entail 'physical object'). (As we saw in Chapter 2, David Chalmers disputes this claim.)

A supporter of Wittgenstein's position could reply by saying that if there is an agreement now in taking these words to be correctly exemplified, then disputing that this that I am seeing in front of me is a hand, or is a physical object would be tantamount to not using those words with their

current ordinary meaning. Given the primacy of ordinary use, for Wittgenstein, that would be equivalent to not using those words with *any* meaning. Conversely, if use changed and there was an agreement, as in the case of 'witch,' over the fact that it is not exemplified by any entity, then the meaning of 'hand' and 'physical object' would have changed as well. (See Chapter 2 for a discussion of the simulation hypothesis, whereby 'hands' and 'physical object' would have a different meaning.) Yet, from this it would not follow that now, while retaining their ordinary meaning, it could be legitimately doubted that I am seeing my hand right now, or that there are physical objects, of which my hand is an instance.

Nonetheless, even if the account of meaning endorsed by Wittgenstein were correct, it would not safeguard against the possibility, which we already explored in connection with Putnam (see Chapter 2, Section 2.2), that we might just have been 'envatted.' In such a scenario we would retain the same concepts and linguistic meanings (yet we could not speak), at least for a while, whereas systematically offering false descriptions like 'Here is my hand' or 'I am not a BIV.'

Second, there is, in Wittgenstein, a "hinge demarcation problem" (Coliva and Palmira 2021). For example, 'Here is my hand' seems to be an ordinary empirical proposition, supported by plenty of evidence. Or else, take some of his other examples, such as 'My name is L.W.' or 'Nobody has ever been on the moon.' The last one turned out to be false, since in 1969 the moon landing occurred; and if it had turned out that Wittgenstein was mistaken about his name, this would have been a disturbing discovery for him, but it would not have made it impossible for him to take part in many ordinary epistemic practices. By contrast, propositions like 'There are physical objects' or 'the Earth has existed for a very long time' seem to be constitutive of the very possibility of participating in inquiries about objects in our surroundings or events that have taken place in the distant past like the ones normally investigated by doing history or geology. Yet, Wittgenstein does not seem to see any principled difference between these diverse propositions and considers all of them hinges.[3]

Third, depending on the details of the kind of framework reading one endorses (see Moyal-Sharrock 2005), saying that hinges are rules may lead one to hold that they have no descriptive content whatsoever and are in fact not propositions at all. Yet this seems problematic, for they can be embedded in suppositions—such as, 'If this is my hand/there is an

external world, then I am not a BIV/we don't live in the Matrix'—and in meaningful negations—such as, 'It is false that there are no physical objects.' Moyal-Sharrock holds that in these occurrences we would be considering the doppelgängers of the relevant hinges—that is, sentences that are identical to the ones with which we express the pertinent hinges, but that would not express them—and not those hinges themselves.

Fourth, following on from the non-propositional reading of hinges, one may even hold that *qua hinges*, hinges are never said, but only manifested in action (see ibid.). Yet again, on the face of it, this seems problematical. After all, we do say them in contexts that Wittgenstein himself would have considered perfectly legitimate, like explaining to someone—either a pupil or someone learning a second-language—the meaning of the relevant words, or explaining the role of hinges within the context of a philosophical discussion.

Fifth, Wittgenstein's idea (1953, §119) that (most) philosophy originates in linguistic and conceptual confusions would no doubt be rejected by most practitioners of the discipline nowadays. Similarly, his according pride of place to ordinary linguistic and epistemic practices, such as the ones we engage in when we raise ordinary doubts about the existence of specific physical objects, while leaving intact the assumption that there are physical objects in general and that our senses are reliable guides in finding out about them, would no doubt be challenged by many philosophers.

Finally, one could grant Wittgenstein's points about meaning and the primacy of ordinary practices and language over philosophy and yet consider the following. If we don't have a justification for 'There is an external world', and the principle of Closure (understood either synchronically or diachronically, as we saw in Chapter 3, Section 3.3) applies, we don't have a justification for 'Here is a hand' either (or take another genuinely empirical proposition that may serve as a premise of Moore's proof and that wouldn't be regarded as a hinge by Wittgenstein). From this perspective, skeptics not only would show that there is no justification for one of our basic beliefs, but they would also show that *all* of our empirical beliefs are unjustified (and therefore, *a fortiori* that we have no knowledge, if knowledge requires justified true belief). To block such a disastrous consequence, one could insist that, within Wittgenstein's framework, hinges make it possible for one to acquire justification for ordinary empirical belief. Hence, in a 'moderate' vein, their being unjustified and unjustifiable does not prevent the possibility of having knowledge and

justification for ordinary empirical propositions, the epistemic good standing of which depends on taking hinges for granted. Yet, this would entail a (partial) rejection of the Closure principle, and, as we saw in Chapter 3, Section 3.3, this is an outcome that several contemporary epistemologists find uncomfortable.

7.4 Williams' contextualism

There is another sense in which skeptical doubts are unnatural, as Michael Williams (1991) has emphasized: not because they are entirely legitimate, yet impotent against what it is natural for us to believe, as Strawson has it (see Section 7.2), but because they trade on a conception of the structure and nature of justification which is problematical and in fact entirely unlike the way justification and knowledge operate in our epistemic practices. As Williams puts it: "sceptical arguments depend essentially on theoretical commitments that are *not* forced on us by our ordinary ways of thinking about knowledge, justification, and truth" (1991, 31–32). Yet, like Strawson, Williams too thinks that Wittgenstein's *On Certainty* contains a powerful antidote against skepticism.

According to Williams, skepticism, in either of its Cartesian or Humean varieties, is rooted in foundationalism, and, moreover, in the idea that at the foundations lie either mind-dependent truths, which enjoy a maximum degree of certainty, like the Cogito—"I think, therefore I am"—or impressions, ideas or sense data—that is, mind-dependent entities, hopefully formed by means of the causal interaction with mind-independent entities like physical objects, such that, while being infallibly known, make the inference to their alleged causes only probable. Whatever form of foundationalism is endorsed, for Williams, it depends on a realist—and, therefore, mind-independent—view of the structure of justification and knowledge. That is to say, when it comes to the metaphysics of epistemology, foundationalists—let them be rationalists, like Descartes, or empiricists, like Hume and his contemporary heirs in analytic philosophy, like Russell, Ayer, and Moore, to some extent—hold that the kind of justificatory structure their respective systems posit is objective and absolute, that is independent of human interests and practices.

According to Williams, Wittgenstein has in fact launched a powerful attack against epistemological realism and against any form of foundationalism. For, according to Wittgenstein, knowledge and justification

stem from our epistemic practices, which are premised on specific, and contextually dependent hinges, or "methodological necessities," as Williams calls them. Such methodological necessities, moreover, are not the ones identified by empiricists—that is, some basic sensory experiences from which we should then derive justification for our beliefs about ordinary physical objects. Nor are they *a priori* truths, like the Cogito, from which the rest of our knowledge should be derived. Rather, they are ordinary empirical propositions like "The Earth has existed for a very long time" (OC, §316), which is a methodological necessity of doing history and geology; or "Instruments do not appear and disappear out of their own accord," which is a hinge of doing scientific experiments; or "Figures do not appear and disappear from paper out of their own accord" (OC, §337), which is again a methodological necessity of doing math, etc. (see OC, §§53, 57, 447–448).

Furthermore, what is a methodological necessity in one of our epistemic practices could become the object of further investigation in a different one. For instance, 'I have hands' could be a hinge in ordinary circumstances, while becoming subject to verification after a car accident (cf. OC, §§23, 372). Thus, according to Williams, Wittgenstein opposes traditional forms of foundationalism by identifying hinges that are not mind-dependent. Moreover, he is also dismantling epistemological realism by endorsing a form of epistemological contextualism, such that context alone determines what, within it, and within it only, plays a hinge role or not. Connectedly, whether a proposition is a hinge is not something that we can read off its content. Rather, we can determine it only by attending to the role it plays within a specific context. To be clear, epistemological contextualism should not be confused with the kind of semantic contextualism we discussed in Chapter 3, Section 3.3. The key claim of the former is that the structure of justification varies from context to context, such that what stays put and allows one to gain justification for other propositions in one context can itself become subject to verification in a different context. By contrast, semantic contextualism holds that the standards for justification can be more or less stringent depending on the context, thus making the same knowledge attribution—'S knows that P,' where P is the content of S's original assertion, and is based on a given kind of evidence—either true or false. To exemplify the difference: epistemological contextualism holds, as we have seen, that 'I have hands' can be a hinge in normal circumstances and be subject to verification after a car accident. In such a case, the same sentence can be uttered, but it would be

uttered in two different epistemic circumstances in which two different epistemic methods would be in force: in the former circumstances 'I have hands' stays put, is not known or verified through experience and can serve to determine, for instance, if one's sight is good. Whereas in the latter circumstances, 'I have hands' is no longer taken for granted and is verified through experience, and, to the extent that experience supports it, it is known. By contrast, semantic contextualism holds that whereas 'I have hands' is known in ordinary circumstances as it is supported by our sensory experience, it is no longer known, in the same circumstances and by application of the same epistemic methods, just by raising the standards and by imposing, for instance, that one should be able to rule out skeptical hypotheses.

Moreover, Williams claims that:

> To adopt contextualism, however, is not just to hold that the epistemic status of a given proposition is liable to shift with situational, disciplinary and other contextually variable factors: it is to hold that, independently of all such influences, a proposition has no epistemic status whatsoever. There is *no fact of the matter* as to what kind of justification it either admits or requires. (Williams 1991, 119)

Contrary to Wittgenstein's position, however, for Williams, philosophical doubts are intelligible even though they rest on the two questionable assumptions we have just reviewed—namely, foundationalism and epistemological realism. Yet, once their questionable assumptions have been exposed, we should no longer be bothered by them.

More precisely, according to Williams, it is only if we are committed to foundationalism and epistemological realism that we find skeptical arguments threatening at all. Yet, if neither of these commitments is 'natural,' and is indeed entirely questionable once further analyzed, we may avoid engaging with skepticism all together. By revising these underlying assumptions, while showing them to be illegitimate, Williams is thus proposing an undercutting response to skepticism. For he does not grant that there is a legitimate philosophical paradox to be addressed. Rather, according to him, we are under the illusion of there being one because we are—no doubt mostly implicitly—committed to foundationalism and epistemological realism. Once freed from that commitment, we are also freed from skeptical worries and preoccupations.

Several objections have been raised against Williams' account both in point of exegesis of Wittgenstein's writings and *per se*. As we saw in Section 7.3, Wittgenstein arguably holds a much more radical view concerning the status of philosophical doubts. On this interpretation of his views, they are senseless and not just based on dubious theoretical commitments.

Furthermore, some Wittgenstein scholars (see Wright 2004; Moyal-Sharrock 2005) hold that Wittgenstein did propose a form of foundationalism, albeit different from traditional foundationalism of either a rationalist or of an empiricist variety. For he did hold that at the foundations of our epistemic practices and ways of operating with language, there are hinges which are discontinuous with respect to what is epistemically dependent on them—either because they are non-evidentially justified and yet necessary to possess evidential justifications, or else because they are not justified at all and yet are necessary for the whole of our knowledge to be possible.

Now, though in many ways instructive, such an objection ultimately rests on a terminological dispute. That is, whether it is legitimate to call 'foundationalism' a position which, like Wittgenstein's, poses hinges at the foundations of knowledge and justification, while not identifying hinges in terms of their content, let alone with beliefs about our own minds, or with mental entities. Moreover, there are passages in Wittgenstein that appeal to the foundational image, while other ones clearly and unequivocally don't (OC, §§152, 248). Thus, Williams might respond by saying that even if one wanted to reserve the term 'foundationalism' for Wittgenstein's position, it would be so different from its more traditional counterpart that it would not make Wittgenstein a foundationalist in this traditional sense at all. Thus, Williams could insist that he was right to maintain that Wittgenstein was not a foundationalist *in that sense*.

Furthermore, it has been argued (see Pritchard 2018c; 2019b; and Chapter 4) that while Wittgenstein was no epistemic relativist because in his conception hinges were not as contextually variable as Williams'—in that they are all ultimately the expression of the *über hinge* commitment that we cannot be systematically and massively mistaken—Williams' position is hospitable to relativism. Indeed, if we take seriously his idea that no hinge is ever absolutely immune from the possibility of being demoted, it quite straightforwardly follows that relativism, if not a *de facto* possibility,

remains a possibility *de jure*. Indeed, for Williams, there is *de facto* widespread disagreement over hinges. Whereas, as we saw in Chapter 4, for Pritchard, this is not even a possibility.

Again, while interesting and instructive, such an objection rests on a specific understanding of Wittgenstein's hinges which is not widely shared among Wittgenstein scholars, as well as on an anti-relativist reading of *On Certainty* that, while quite common to many interpreters of *On Certainty* (such as Coliva, Moyal-Sharrock, and Pritchard, and interestingly Williams himself, just to mention a few), is opposed by several others (such as Kusch, but also Kuhn, Kripke, Rorty, and more).

In recent years, however, Williams has put forward a reading of Wittgenstein's hinges which inherits some elements from Wright's and has argued in favor of an epistemic account of them. Indeed, according to Williams, we may have entitlements, i.e. non-evidential justifications, and perhaps even knowledge, of contextually determined hinges, which could be regarded as minimally truth-apt propositions. As he writes:

> Holding some particular propositions fast need not be a matter of credulity: to hold them fast is reasonable. So while they are (in one sense) ungrounded, we are justified in cleaving to them. In this way, they can be the objects of beliefs that are true (in a deflationary sense of 'true') and justified (though not derived from evidence). Thinking of 'justified true belief' in this second way, even basic certainties can amount to knowledge.
>
> (Williams 2004b, 280)

That is, according to Williams, if we consider hinges true not in virtue of their correspondence with a pre-ordinate mind-independent reality, or of their being supported by all available evidence, but rather just in virtue of the fact that they can be embedded in the disquotational schema:

> 'P' is true if and only if P

and if, moreover, we consider them justified, though non-evidentially, we can even say that we know them. For we would believe them, while they are true and justified, albeit non-evidentially. Hence, according to Williams, they would be known, at least in an externalist sense which does not require subjects to have subjectively available reasons in favor of their true beliefs.

Yet, it has to be noted that Wittgenstein is not univocal about the possibility for hinges to be true—albeit in a minimal sense[4]—and in several places, he actually denies that they are either true or false (OC, §§94, 162, 204–206).

Moreover, in the very passage, which, ironically, is often appealed to by supporters of the epistemic reading, Wittgenstein denies the very idea that hinges enjoy any kind of positive epistemic status. Here is the passage in its context:

> The questions that we raise and our doubts depend on the fact that some propositions are exempt from doubt, are, as it were hinges on which those turn.
>
> (OC, §341)

> That is to say, it belongs to the *logic* of our scientific investigations that certain things are in deed not doubted.
>
> (OC, §342; our emphasis)

> But it *isn't* that the situation is like this: we just can't investigate everything, and for that reason we are forced to rest content with assumption. If I want the door to turn, the hinges *must* stay put.
>
> (OC, §343; our emphasis in both places)

Here Wittgenstein is introducing the idea of hinges and is making clear that it belongs to the *logic*—i.e. to the *norms* that regulate our investigations—that certain propositions are in fact not doubted. This, however, isn't due to the fact that since we cannot—pragmatically—investigate everything, if we ever want to start investigating anything, we have to content ourselves with assumptions, as Wittgenstein explicitly acknowledges in OC, §343. By contrast, on Wright's view, as we saw in Chapter 6, Section 6.4, it is precisely the fact that certain cognitive projects are valuable or even indispensable to us that forces us, in order to get them started and to pursue them, to rest content with their characteristic assumptions. Here is a telling passage:

> If a cognitive project is indispensable, or anyway sufficiently valuable to us ... and if the attempt to vindicate (some of) its presuppositions would raise presuppositions of its own of no more secure an antecedent status,

> and so on ad infinitum, then we are entitled to ... the original presuppositions without specific evidence in their favour.
>
> (Wright 2004, 192)

Yet, for Wittgenstein, saying that it belongs to the *logic* of our investigations that certain propositions aren't doubted simply means that it is a constitutive element of our empirical inquiries that, whenever evidence is collected for or against any empirical proposition, certain other propositions have to be kept fixed. Hence, it is constitutive of empirical inquiries, and not just a fact about our nature or pertaining to pragmatics, that certain propositions are taken for granted.

However, saying this does nothing to corroborate the idea that hinges are therefore reasonable, and they are reasonable because they would be somewhat epistemically justified, albeit non-evidentially, as they would be required in order to collect evidence for or against ordinary empirical propositions. Indeed, Wittgenstein repeatedly points out that they are neither reasonable nor unreasonable (OC, §559). Yet he does say that they are required in order to collect evidence for or against ordinary empirical propositions. What that means, rather, is that they are rationally mandated. Yet a rational mandate is not a *sui generis* kind of justification. It is a requirement for reasons and reasoning to be possible in the first place, not a product of reason, which, like any proper epistemic good, makes certain propositions somehow justified and therefore reasonable. Hence, hinges must stay put, and thus behave like rules, for us to be in a position to acquire evidence that may add up to justification and sometimes to knowledge, or that can motivate rational doubts and eventually the dismissal of a previously held belief. Thus, there is nothing in the key passages usually appealed to by supporters of entitlements to corroborate their view. Nor is there anything which suggests that knowledge, for Wittgenstein, is to be understood somewhat in an externalist spirit and that it can be attained even if one is in no position to produce reasons for one's own true beliefs.

In fact, the idea that Wittgenstein would allow for the possibility of knowledge even when no reasons can be produced in favor of one's claim is based on disregarding the distinction between empirical (ordinary) and grammatical uses of 'I know' he draws in *On Certainty*. As Wittgenstein clearly writes:

> If "I know etc." is conceived as a grammatical proposition, of course the "I" cannot be important. And it properly means "There is no such thing as a doubt in this case" or "The expression 'I do not know' makes no sense in this case". And of course it follows from this that "I know" makes no sense either. "I know" is here a logical insight. Only realism can't be proved by means of it. (OC, §§58–59)

It should be kept in mind that, for Wittgenstein, only grammatical uses of 'I know' are not backed by reasons. Furthermore, and as weird as that may be, they don't express, for him, an epistemic relation between a subject and a proposition. Rather, they express a kind of certainty, which would be more perspicuously expressed by saying 'Here a doubt is (logically) impossible' (see Coliva 2021a and Schönbaumsfeld 2021, on the grammatical use of 'I know'). Thus, we do not have anything in *On Certainty* resembling the externalist idea that, even if we cannot produce reasons in favor of knowledge claims regarding hinges, we would know them, nonetheless. Rather, we have the suggestion that the use of 'I know' in their connection is a misleading expression of the kind of certainty, which characterizes them. Namely, a certainty that has nothing subjective or psychological about it, but which depends on the role hinges play in the relevant epistemic practices. That is, all we get from these and other key passages in *On Certainty* is the idea that justification and knowledge do not take place in a vacuum. They always depend on there being certain hinges, which, as such, cannot themselves be justified or known, yet allow us to acquire evidence for or against ordinary empirical propositions. They are therefore constitutive of the practice that in turn determines what being epistemically rational amounts to. Some of them can change in time—'Nobody has ever been on the Moon' is no longer a hinge for us—and according to context—sometimes 'Here is my hand' is an empirical proposition we subject to verification and control. Yet, there are several of them we simply cannot revise, e.g. 'The Earth has existed for a very long time,' 'There are physical objects,' etc. on pain of giving up all our system of beliefs. This is not to say that they are metaphysically necessary but only that they play such a fundamental role in our world-picture that, from within it—as we in fact are—we can't actually find any reason to doubt them.

Turning now to some objections to Williams' account considered *per se*, a disjunctivist (see Pritchard 2019b) could claim that one need not embrace Williams' contextualism in order to get rid of foundationalism. For,

according to disjunctivism, we do have factive reasons in favor of our ordinary empirical beliefs, whenever we are perceiving objects in our surroundings.

Furthermore, it has been argued (see DeRose 1993) that Williams' account depends on a problematic understanding of the source of skeptical doubts. For instance, if one were convinced by Nozick's idea that we cannot know the denials of Cartesian skeptical scenarios because, even if they obtained, we would believe that they don't, one would find oneself at the mercy of skeptical hypotheses without thereby endorsing traditional foundationalism. (See also Pritchard 2019b, for a related objection that Williams misunderstands the source of skepticism.)

Finally, it must be noted with Pritchard that

> [a] key part of Williams's diagnosis of radical scepticism is the observation that there is a sense in which scepticism is correct, albeit not (as it purports to be) as an contextual thesis, but only relative to a particular set of methodological necessities. (2019b, 230)

Here is Williams:

> The sceptic takes himself to have discovered, under the conditions of philosophical reflection, that knowledge of the world is impossible. But in fact, the most he has discovered is that knowledge of the world is impossible under the conditions of philosophical reflection.
>
> (1991, 130)

Thus, there is a sense in which a skeptic wins, by Williams' lights. The victory is much more limited than a skeptic normally thinks of it since it holds only within the philosophical context. Yet, it is a victory, nonetheless. More importantly, even if contextualism holds and epistemological realism is dismantled, once it is further noted that skeptical hypotheses need not necessarily depend on embracing foundationalism, it turns out that skepticism not only isn't defunct, but it also remains undiagnosed.

7.5 Coliva's constitutivism

In Chapter 6, we saw that in the dispute between liberals and conservatives it is problematic to maintain that, in the absence of defeaters, experience

can be an immediate justification for 'Here is a hand,' without assuming, trusting or taking for granted—most often implicitly—that there is an external world. As we noted in Section 7.1, however, this does not *ipso facto* imply that there should be an *a priori* justification for that proposition or an entitlement for it, contrary to what Wedgwood and Wright respectively maintain (see Chapter 6, Section 6.4). In fact, all that follows from the criticism of the liberal position is just that in order to have an empirical justification based on one's experience, in the absence of defeaters, the collateral assumption that there is an external world must be in place (possibly together with the assumption that one's senses are functioning properly and that one is not the victim of a lucid and persistent dream). As we saw in Section 7.1, this is the gist of the so-called "moderate" account of perceptual justification, as Coliva (2015) calls it. To repeat: to have a justification for specific propositions such as 'Here is a hand,' in the absence of defeaters, one need not only have an appropriate experience, but also make a collateral assumption concerning one or more of the 'heavyweight' presuppositions just mentioned. In this way, one's own occurrent experience can constitute itself as justification for 'Here is a hand' rather than for its skeptical counterpart 'I am dreaming of seeing a hand.' Still, it remains that this would be a fallible justification, which could be defeated if contrary evidence came in—for instance, if it turned out that the object in front of me is not a human but a plastic hand instead.

Notice that Sosa's (2021) position too, which was briefly discussed at the end of Chapter 2, qualifies as a form of moderatism. Sosa's version of it, however, is of an externalist kind. For it posits that if it is in fact the case that there is an external world, then we can acquire a justification for our ordinary empirical beliefs based on the deliverances of perception. Yet, like Coliva's version of it, and being a form of virtue-reliabilism, it also requires subjects to be somewhat sensitive to the obtaining of that presupposition.

There are three main challenges the moderate position must face. The first is to clarify the nature and role of those assumptions on which the existence of empirical justifications depends. The second is to give a verdict on the cogency of Moore's argument and thus determine why an empirical justification cannot be acquired for its conclusion. The third is to find a convincing answer to the skeptical challenge: if naturalist (and pragmatist) positions embrace moderatism but are not able to solve the skeptical paradox, as we have seen in this chapter, can a moderate position be developed such that it would do better in this regard?

7.5.1 Hinge assumptions

As to the first challenge, it is important to note that, like Wright and Pryor (see Chapter 6, Sections 6.3 and 6.4), Coliva is interested, first and foremost, in propositional justification. Thus, the claim is that, 'in the abstract space of reasons,' the proposition 'There is an external world' has to be countenanced for perceptual justification to be possible at all (together with certain experiences, in the absence of defeaters). Yet, subjects can appropriate these justifications, when they exist, and when they entertain the appropriate beliefs. That is, they can have doxastic justifications for their beliefs. Thus, the question is: besides having the relevant experiences, absent defeaters, what else is needed of them to appropriate those propositional justifications? Certainly not to consciously entertain such presuppositions, for, otherwise, not many would ever be able to appropriate them. For example, children and many adults who lacked the concepts relevant to entertaining those presuppositions could not appropriate those propositional justifications. Nor would it be necessary for subjects to be able to make those presuppositions explicit, if required to do so, even if it would certainly be sufficient. Rather, subjects should be committed to the truth of P and to acting rationally and practically accordingly. This is something that both children and adults who are otherwise unable to entertain or articulate those presuppositions can and ordinarily do. In this sense, assuming that P is different from making the hypothesis that P, with respect to truth of which the subject may be uncertain, agnostic, or even doubtful. It is indeed more like trusting in and being committed to the truth of P in one's verbal and non-verbal actions.[5]

More specifically, we could say of a subject that she *implicitly* assumes (trusts in, or is committed to) P—even if she does not have the concepts to entertain that proposition—if she is able to take part in a practice that has that assumption as a precondition of its intelligibility. For example, since using one's own experience as a justification for 'Here is my hand'—rather than for some skeptical counterpart in the absence of defeaters—is only intelligible on the basis of the assumption that there is an external world, that our senses are mostly functioning properly, and that we are not victims of lucid and persistent dreams, then one can say that the subject (implicitly or tacitly) assumes that proposition even if she does not have the concepts to entertain it. Similarly, since going to fetch an object in another room, upon being requested to do so, without questioning the

existence of the object currently unperceived, is rational only if one takes for granted that there are mind-independent physical objects, being able to participate in this practice is enough for being granted with the relevant assumption.

Thus, while both Coliva and Pritchard deny that hinges are knowledge-apt beliefs, Coliva, contrary to Pritchard, does not deny that they are the content of an attitude of acceptance (or trust).

7.5.2 Moderatism and Moore's proof

Let us now move on to the second problem. Namely whether it is possible to acquire a justification for 'There is an external world' through Moore's argument once we have embraced the moderate conception of justification. One might argue that it is by reasoning as follows. According to a moderate, in order to have a perceptual justification for 'Here is a hand,' it is enough to have an experience with an appropriate representational content and to assume that there is an external world, in the absence of defeaters. Furthermore, 'Here is a hand' knowledgeably entails 'There is an external world.' Hence, by running Moore's argument, one would also acquire a justification for 'There is an external world.' Moderatism would therefore have the same outcome as liberalism with respect to the epistemological cogency of Moore's proof (see Chapter 6, Section 6.3).

Obviously, it is possible to support moderatism with respect to the structure of empirical justification while separating it from the issue of the epistemic cogency of Moore's proof. However, this does not appear desirable from a more general philosophical point of view. Here is why. Let us ask how we have come to the moderate position. We have seen that it stems from dissatisfaction with Pryor's liberalism and the consequent admission that, in order to have perceptual justifications for specific empirical propositions, one must make collateral assumptions as such unjustified (unless one resorts to *a priori* strategies or to entitlements, which seem problematic for independent reasons, as we saw in Chapter 6, Sections 6.4 and 6.5). At this point the following perplexity arises: how can an argument, in which the justification for one of the premises depends on assuming the conclusion, provide a justification for that very conclusion? It seems problematically circular to first assume q, without justification for doing so, and then run an argument, which starts from justified premises only if q is assumed, to thereby derive a justification for q

itself. Indeed, this kind of reasoning would resemble *bootstrapping* arguments whose cogency is far from obvious (as we saw in Chapter 6, Section 6.3). Or else, think of a mathematical proof. It would clearly be circular if, during the derivation, we assumed the very conclusion—that is the theorem—we would like to prove by means of it. Even worse, it would clearly be circular if we tried to prove one of the theory's axioms by means of a proof whose premises could be justified only as long as that very axiom is taken for granted.

Similarly, it seems that there is something illegitimate in trying to obtain a justification for 'There is an external world' through Moore's reasoning, which could not start with a justified premise unless that proposition were already assumed—without any justification to do so, according to moderates. Again, think of the analogous case in which one wanted to justify 'The Earth has existed for a very long time' by means of an argument such as 'This fossil dates back millions of years ago; if this fossil dates back millions of years ago, then the Earth has existed for a very long time; therefore the Earth has existed for a very long time.' It is important to note that for the first premise of the argument to be justified, not only must there be evidence in favor of it, but also the conclusion must already be taken for granted. If the conclusion were not taken for granted, the evidence at one's disposal would no more justify the argument's premise than it would justify its skeptical counterpart—that is, 'here is a fossil that has just popped into existence with everything else that makes it look as if it were millions of years old.' Therefore, such an argument cannot produce a justification for its conclusion.

At the end of the day, Wright argued the same thing when he noted that Moore's argument does not transmit justification. However, his notion of Transmission failure appeals to a slightly different idea, namely, that the assumption that guarantees justification for the first premise is itself *justified* (Chapter 6, Section 6.4). Therefore, the epistemic circularity of Moore's argument results from the fact that it is already necessary to have a *justification* for its conclusion in order to have a justification for its premise.

However, circularity remains even if the *mere assumption* of the conclusion of the argument is necessary for it to start with justified premises, as we have just seen. This means that, in fact, there are two types of Transmission failure: Wright's and the one just introduced. We can define them as follows:

(1) *Transmission failure 1*: A (logically valid) argument does not transmit justification from its premises to its conclusion if and only if it is

necessary to already have a *justification* for its conclusion in order to have a justification for one or more of its premises.

(2) *Transmission failure 2*: A (logically valid) argument does not transmit justification from its premises to its conclusion if and only if it is necessary already to *assume* its conclusion in order to have a justification for one or more of its premises.

It should be noted that this second type of Transmission failure occurs when the assumption of the conclusion is one of the *constitutive* conditions of the justification for one or more of the premises and not simply one of the conditions, the obtaining of which may contingently be necessary in the process of acquiring those justifications.

These two notions of Transmission failure (see Coliva 2015, Chapter 3) are not alternative to one another. In fact, it seems that the second applies to all cases—which are the most interesting philosophically—in which the assumptions cannot be justified in any way. Instead, the first notion of Transmission failure would apply to all those arguments in which the conclusions are—at least in principle—justifiable independently of the argument in question. An example of Transmission failure 1 would therefore be the zebra example we already encountered in Chapter 3, Section 3.2:

(1) Here is a zebra.

(2) If this is a zebra, it is not a cleverly disguised mule.

(3) This is not a cleverly disguised mule.

In fact, it is possible at least in principle to get a justification for (3) independently of our current visual experience, which is what justifies us to believe (1). For example, one could run a DNA test that shows that the animal in front of us is not a mule, without yet telling us that it is a zebra. Therefore, the diagnosis of Transmission failure 1 would tell us that the foregoing argument cannot give us a first justification to believe (3) because the justification for (1) depends on already having a justification for (3). The latter is at least in principle obtainable independently of that very argument (or of other ones where the justification for the premises depends on already having a justification for (3)).

By contrast, arguments such as Moore's, or the argument in favor of the long existence of the Earth, or for the existence of other minds, such as:

(1) Here is a person in pain.

(2) If this person is in pain, then there other minds.

(3) There are other minds.

Where the justification for the first premise is based on the observation of a person crying and moaning would be cases of Transmission failure 2. For we cannot provide independent justification for (3), on the assumption that no *a priori* justification can be provided and that there is no entitlement for it. Thus, arguments aimed at giving us empirical justifications for (3) would in turn be based on the assumption of (3) itself and therefore could not give us a justification for it.

This has a notable consequence: arguments such as Moore's are also cases of failure of the Closure principle for justification, in this case, and to the extent that justification is a necessary component of knowledge, for knowledge as well. In contrast, the zebra argument usually used to illustrate Closure failure, wouldn't, as long as we understand Closure synchronically (see Chapter 3, Section 3.2). More generally, we can see how Closure failure is a consequence of Transmission failure 2. Alternatively, if one opted for diachronic Closure, Transmission failure 2 would explain why we cannot acquire a justification for the conclusion of Moore's argument and hence why there isn't one for its conclusion, even though, by moderates' lights, there is a justification for its premises.

This consequence could discredit the very notion of Transmission failure 2, but thinking this would be a mistake. In fact, we have explained why the Closure principle fails (in either of its possible readings) and we have seen that it does so in very few cases and for motivated reasons. To repeat: it would fail only when the conclusion of a logically valid argument, which proceeds from justified or known premises, through known entailment, was a constitutive presupposition of the justification of the premises and could not be independently justified. Therefore, Closure remains valid for most of our ordinary reasonings and admitting its limited failure does not affect our usual epistemic practices. As we noted in Chapter 3, Section 3.2, sometimes the best we can do in order to defend and protect a rule is to note that it admits of motivated exceptions.

Notice, moreover, that compared to Pritchard's variety of hinge epistemology, Coliva's version takes hinges to represent a counterexample to the universal validity of Closure. By contrast, Pritchard thinks (as we saw in Chapter 4) that there are no counterexamples to the Closure principle. For he reads it as involving knowledge-apt belief. Since hinges, according to Pritchard, are not the contents of such an attitude, they simply fall outside the domain of application of that principle.

As to the charge that this would open the way to "abominable conjunctions" (cf. Chapter 3, Section 3.2), we are now able to understand why we cannot really have a justification for heavyweight assumptions (and we cannot have knowledge of them), while these are what makes it possible to have a justification for (or even knowledge of) our ordinary empirical propositions. Yet we can do more, as we shall see in a moment. In particular, we can show how they are epistemically rational, although unjustified and unjustifiable. This will allow us to assert something far less "abominable" than 'We have a justification for believing (or, even, we know) that there is a hand, but we are not justified in believing (or, even, we don't know) that there is an external world.' We will return to this issue shortly.

7.5.3 The constitutive response to skepticism

We thus come to the last challenge that moderatism must face in order to be plausible. Moderatism needs to provide a more convincing answer to skepticism than the ones given by naturalists *à la* Strawson (and perhaps *à la* Wittgenstein, for some interpreters) and pragmatists. For these latter kinds of moderate, no epistemic justification can be attained for 'There is an external world' and other heavyweight presuppositions, but, at most, only pragmatic ones.

Now, it is important to say right away that the solution put forward by moderatism is not a direct answer to the skeptical paradox, but an indirect one. It therefore belongs to the family of overriding, or skeptical responses to skepticism. In other words, it does not provide epistemic justification—that is to say, justifications capable of corroborating the truth of a proposition—for 'There is an external world.' Hence, it does agree with a skeptic that such a proposition is also unknowable. Instead, it offers a diagnosis of what can be considered the ungrounded and in fact mistaken assumption on which the skeptical paradox rests. Namely, that epistemic rationality falls short of the

presuppositions which make it possible to acquire perceptual justifications for and against ordinary empirical beliefs.

The starting point of the constitutive strategy is the analogy with games and their constitutive rules. Just as constitutive rules are part of the game, even if they are not moves within the game, so the constitutive presuppositions of epistemic rationality are part of epistemic rationality even if they are not justified or justifiable, precisely because they allow us to acquire (perceptual) justifications for or against ordinary empirical propositions, and thus to play the 'game' of epistemic rationality itself.

Hence, it is true that we simply take for granted that there is an external world, without any kind of epistemic justification to support it. Yet, doing so is *epistemically* rational, nonetheless. For 'There is an external world' is a *hinge*, that is a *constitutive* assumption of our practice of producing and assessing epistemic justifications for or against any empirical proposition. To say that it is a constitutive assumption of that practice is to say that it is presupposed by every justification and every justified doubt of an empirical kind. A first notable consequence is that this presupposition can neither be justified nor legitimately doubted, because the reasons that could hypothetically be put forward to justify, or doubt it, would all in fact take it for granted.

However, there is more: taking for granted that there is an external world is constitutive of the epistemic practice of producing and assessing epistemic justifications (based on perceptual experience) for or against any empirical proposition. In turn, this practice is what determines our notion of epistemic rationality itself. This notion, in fact, does not come to us from nothing or from armchair reflection, but precisely from our way of operating.

Now, if this is the case, to revoke into doubt one of its constitutive presuppositions would mean that the practice in question, and with it the very notion of epistemic rationality, would be lost. Indeed, what Humean skepticism has taught us is precisely that epistemic rationality exists only *within a system of assumptions*. To put it another way, it has taught us that epistemic rationality is not absolute—that is, independent of any assumption—but is in fact *bound* by (or even *situated* within) some very general assumptions.

However, this observation should not be understood as a pragmatic defense of the very notion of epistemic rationality. Therefore, the idea is not to defend oneself against skepticism by saying that, if our *aim* is to

preserve the epistemic practice that is dear and useful to us and with it the notion of epistemic rationality that follows from it, then we must assume 'There is an external world.' Rather, the idea is that to the extent that the very notion of epistemic rationality itself depends on that practice, and a skeptic too is appealing to it by requesting epistemic justifications for our beliefs, calling into doubt the existence of the external world would not be the most rational of moves, but rather the violation of epistemic rationality itself. This violation, in fact, depends on not realizing that epistemic rationality pertains not only to ordinary empirical propositions for which we can produce perceptual justifications, but it also extends to those presuppositions, such as 'There is an external world,' which, though unjustified and unjustifiable, make it possible to acquire perceptual justifications for specific empirical propositions.

Thus, the skeptic's mistake would not so much lie in ignoring the width and breadth of our concept of *justification* (or warrant) as Wright thinks (see Chapter 6, Section 6.4), as in not realizing the width and breadth of our concept of *epistemic rationality*, which is itself grounded in a practice that rests on certain background assumptions. By fixating only on one aspect of it—that is, the need for a justification to rationally accept something—skeptics fail to recognize the epistemic rationality of assuming what lies at the basis of the epistemic practice of producing and assessing empirical reasons and which cannot therefore be justified by it.

Therefore, the skeptic's mistake consists in adhering to the following principle:

> *Epistemic rationality 1*: It is epistemically rational to believe only specific empirical propositions for which there is perceptual justification.

When instead the following principle is valid:

> *Epistemic rationality 2*: It is epistemically rational to believe those specific empirical propositions for which there are perceptual justifications, and to assume those propositions that, by being constitutive of them, make the acquisition of perceptual justifications possible.

The latter notion of rationality must in turn be distinguished from the one Wright subscribes to:

> *Epistemic rationality 3*: It is epistemically rational to believe those specific empirical propositions for which there are perceptual justifications and to assume those propositions, which make it possible to acquire perceptual justifications, for which we have an entitlement.

Armed with this notion of epistemic rationality, we are now able to fully address the challenge raised by DeRose, regarding the fact that the denial of Closure would commit one to embrace "abominable conjunctions." For, really, the commitment incurred would be 'I know that there is a hand here, and although I don't know that there is an external world, I am epistemically rationally mandated to assume it.' This is by far a much less abominable conjunction than the one proposed by DeRose—that is, 'I know that there is a hand here, but I don't know that there is an external world.' As Harman and Sherman (2011) have noted, "abominable conjunctions" are bound to arise any time there are subtle philosophical notions that are not easy to tease apart from one another, such as, in our case, knowledge (or justified belief) and epistemically rationally mandated assumptions.

Furthermore, it should be noted that we are not at all arguing here that our *practice* of producing reasons based on perceptual experience for or against specific empirical propositions is as such rational (such a view is supported by Wedgwood 2013). Our practice is indeed what it is—an extension "of our natural history," as Wittgenstein would say—and it could have been different from what it is, even if from the inside it is difficult for us to imagine how. What is epistemically rational is to accept its assumptions since they contribute to the determination of epistemic rationality itself.

Finally, one could object that it is a practice, for some people, to predict the future by reading cards or coffee grounds. Then, what would prevent us from believing that the assumptions of these practices are rational, even though they are not?

In response to this objection, it should be noted that a distinction must be made between basic and non-basic epistemic practices (see Baghramian and Coliva 2020, Chapter 7). The basic ones are those which do not presuppose other epistemic practices, or the use of theories in order to operate. Divining the future by reading the cards or coffee grounds obviously isn't a basic epistemological practice, because it presupposes the formation of beliefs about the cards and the arrangement of the coffee grounds based on one's own sensory experience. It also requires the

application of some theory to derive conclusions about the future. By contrast, the practice of forming empirical beliefs based on one's own sensory experience does not require other epistemic practices, or the application of a theory. Now, according to the present proposal, only basic epistemic practices are constitutive of the respective types of rationality. Therefore, for example, the practice of forming (and revising) beliefs about the objects around us based on sensory experience is constitutive of epistemic rationality. Similarly, it can be argued that the practice of forming beliefs based on certain logical principles (such as *modus ponens*) is constitutive of what we might call 'deductive rationality.' Again, the practice of acting based on means-ends reasoning could be said to be constitutive of practical rationality, and so on. It is only of the presuppositions of these basic epistemic practices that one can say that they are intrinsically rational, *sub specie* epistemic, deductive or practical rationality.

If at the basis of all our operating with reason and justification there is neither an arational "form of life," or an instinctive "something animal" (OC, §§ 358–359), nor propositions that we assume by virtue of their usefulness, or because they are justified *a priori* or by default, but there is the intrinsically rational acceptance from an epistemic point of view of the assumptions of that very practice, then we have a diagnosis of the skeptic's mistake, which allows us to side-step the Humean skeptical challenge. That is, while it is indeed true that we have no justifications, and therefore knowledge of 'There is an external world,' it is epistemically rational to assume it. Finally, thanks to such an epistemically rational assumption, given the moderate account of perceptual justification, we are also in a position to justifiably believe and know many things about physical objects in our surroundings.

Further reading

On Strawson's naturalism, see Bell and McGinn (1990) and Callanan (2011). On Wittgenstein's *On Certainty*, besides the papers in Moyal-Sharrock and Brenner (2005), see Coliva (2010), Schönbaumsfeld (2016), and Pritchard (2017b). On Wittgenstein's (alleged) relativism, see Coliva (2010b; 2020a), Pritchard (2011), and Williams (2007). On Michael Williams' contextualism, see Williamson (2001b) and Pritchard (2019b). On Coliva's constitutivism, see Wright (2012), Coliva (2017a; 2017b), Millar (2017), Avnur (2017), Volpe (2017), Baghramian (2017), Ashton (2021), Moruzzi (2021), Zanetti

(2021), and, for a reply, Coliva (2021b; 2021c). For further discussion of the attitude we bear to hinges, see Coliva (2021c). For further comparison with Pritchard's position, see Coliva (2021d; 2021e) and Pritchard (2021b).

Notes

1 Another reading of *On Certainty*, alternative to Strawson's naturalist one and to the framework reading, is the epistemic reading initiated by Wright (2004) we discussed in Chapter 6, Section 6.4. Other prominent supporters of the epistemic reading are Williams (2004a; 2004b), whom we will be discussing in Section 7.4, and Kusch (2016a). Notice that Wright (1985) supported the framework reading instead. Other prominent supporters of the framework reading are McGinn (1989), Moyal-Sharrock (2005), and Coliva (2010a).

2 This response is indeed in the spirit of Moore's (1949) reply to Norman Malcolm (1949) who was pressing this Wittgensteinian line against him.

3 As we saw in Chapter 4, for Pritchard, this can be explained by conceiving of specific hinges as encoding the *über hinge* commitment that we cannot be massively mistaken. While ingenious, it must be stressed that Wittgenstein nowhere proposed such an idea.

4 He clearly objects to a correspondence and to an evidentialist account of their truth, though. Cf. OC, §§191, 192.

5 Notice that, like Williams (see Section 7.4), Coliva (2018) thinks that truth here should be understood neither as correspondence nor as evidentially-constrained, but rather in a minimalist fashion.

GLOSSARY

Abductivism: The view according to which the best explanation—whether *a priori* or *a posteriori*—of our experience is provided by denying the obtaining of skeptical scenarios. (See also *Inference to the Best Explanation*; *Radical Skeptical Hypotheses*.)

Abominable Conjunctions: These are conjunctions that contain a proposition and the negation of what seems to be directly entailed by it, at least *prima facie* (e.g., 'I know I have hands, but I don't know that I am not a handless BIV').

Agrippa's Trilemma: According to Agrippa's trilemma, there are only three ways of responding to a challenge to a belief that one holds, all of them unsatisfactory. One can offer an infinite chain of new reasons (*infinite regress*), or one can offer a reason in support of itself (*circularity*), or one can offer no reasons at all (*groundlessness*). (See also *Coherentism*; *Foundationalism*; *Infinitism*.)

***A posteriori*:** A proposition is known *a posteriori* if, for its justification, appeal is made to the deliverances of the senses (or of testimony).

***A priori*:** A proposition is known *a priori* if no appeal is made, for its justification, to the deliverances of the senses (or of testimony).

Assertibility Conditions: The conditions within a language that allow a sentence to be meaningfully asserted and understood as such.

Biscopic Response to Skepticism: This is the view that the Cartesian radical skeptical problem consists of two logically distinct, but related, difficulties, and

hence that an integrated, two-pronged response to this form of skepticism is required. In particular, it holds that closure-based and underdetermination-based formulations of Cartesian radical skepticism trade on subtly different skeptical presuppositions, and hence that what is required is an integrated combination of hinge epistemology (to combat closure-based skepticism) and epistemological disjunctivism (to combat underdetermination-based skepticism). (See also *Cartesian Skeptical Paradox*; *Closure Principle*; *Epistemological Disjunctivism*; *Hinge Epistemology*; *Underdetermination Principle*.)

Bootstrapping Argument: A bootstrapping argument aims at establishing the general reliability of a faculty (or an instrument) starting with premises the justification (or knowledge) of which depends on the exercise of that faculty (instrument). It is based on performing an inductive generalization on a few cases in which the faculty (instrument) has performed reliably and concludes that the faculty (instrument) is generally reliable. (See also *Easy Knowledge*; *Epistemic Circularity*; *Reliabilism*.)

Brain-in-a-Vat Hypothesis: The brain-in-a-vat (BIV) hypothesis is a radical skeptical hypothesis. Like all radical skeptical hypotheses it describes a scenario that is indistinguishable from ordinary life, but where one's beliefs are radically in error. In this scenario, one is a BIV who is being 'fed' one's experiences by supercomputers such that one believes that one is experiencing a normal life of engagement with the world around one, when in fact one is merely a disembodied brain floating in a vat. (See also *Radical Skeptical Hypotheses*.)

Cartesian Skeptical Paradox: The Cartesian skeptical paradox is a form of radical skepticism that makes essential appeal to radical skeptical hypotheses. In its contemporary formulation, it is usually thought to essentially appeal to an epistemic principle like the Closure principle or the Underdetermination principle. (See also *Closure Principle*; *Radical Skeptical Hypotheses*; *Underdetermination Principle*.)

Closure Principle: In its simplest form, the Closure principle demands that if one knows one proposition (e.g., that one is sitting down), and one knows that it entails a second proposition (e.g., that one is not standing up), then one also knows that second proposition. One can also formulate versions of the Closure principle that are concerned with justification rather than knowledge. Moreover, there are more complex versions of the Closure principle that construe it as a type of reasoning where it is made explicit that the inferred knowledge is based on a competent deduction from the entailing knowledge. (See also *Transmission of Justification Principle*.)

Coherentism: In broad outline, coherentism is an account of the structure of epistemic justification that treats epistemic justification as arising out of the way in which a subject's set of beliefs coheres with one another. This means that epistemic justification holistically arises out of the subject's set of beliefs

as a whole, as opposed to being attributable to particular foundational beliefs. (See also *Agrippa's Trilemma*; *Foundationalism*; *Infinitism*.)

Common Sense: Against skeptical arguments and hypotheses appeal is often made to truths that are generally widely believed, and that are therefore said to be part of commonsense. Examples include 'There are physical objects,' 'There are people (and therefore other minds),' 'I am a human being,' or 'The Earth has existed for a very long time.' According to G. E. Moore, these truths are also known with certainty.

Conservatism: The view according to which, in order to have perceptual justification for a proposition P about physical objects in one's surroundings, one needs independent justification for the proposition that there is an external world (among other *heavyweight assumptions*), besides having an appropriate course of experience absent defeaters. (See also *Heavyweight Assumptions*; *Liberalism*; *Moderatism*.)

Constitutivism: The view according to which one's hinge commitments (or heavyweight assumptions) can be epistemically rational, even if unjustifiable, because they are constitutive of epistemic rationality. (See also *Heavyweight Assumptions*; *Hinge Epistemology*.)

Content Externalism: Content externalism is the thesis that the content of one's mental states (one's thoughts, beliefs, and so forth) can be dependent on one's relationship to an external environment (including one's social environment). It stands in opposition to content internalism, which insists that the content of our thoughts can be determined completely independently of our interactions with an external environment.

Contextualism: The view that knowledge attributions of the form 'S knows that P' vary with the variability of the standards that are operative at a given context. (See also *Relativism*.)

Dialectically Ineffective Argument: An argument that, while valid and sound, and unaffected by epistemic circularity, is incapable of convincing an opponent who doubts or disbelieves one (or more) of its premises.

Easy Knowledge: Reliabilist views of knowledge and justification, as well as liberal accounts of perceptual justification, are often challenged to block the seemingly unwanted consequence that by their lights just by having a perceptual experience as of P ('Here is a red wall'), one acquires knowledge of a proposition Q that would give rise to the same perceptual experience while being incompatible with P ('The wall is not white bathed in red light'). (See also *Reliabilism*.)

Entitlement: A form of justification or warrant for a proposition that does not depend on either *a priori* or *a posteriori* evidence for that proposition.

Epistemic Circularity: An argument is epistemically circular if the justification for one or more of its premises depends on assuming or having already a

justification for its purported conclusion. (See also *Transmission of Justification Principle*.)

Epistemic Externalism/Internalism: In its simplest form, epistemic internalism about knowledge holds that a necessary condition of knowledge is that it is supported by reflectively accessible reasons, whereas epistemic externalism about knowledge denies this. So, for example, an epistemic externalist about knowledge might claim that a belief can amount to knowledge even though the subject is unable to offer any supporting reasons for it, just so long as it satisfies an external epistemic condition (i.e., the obtaining of which is not reflectively accessible to the subject), such as that the belief was formed via a reliable belief-forming process, or that it meets the Safety condition on knowledge. (See also *Reliabilism*; *Safety Condition*.)

Epistemic Rationality: It is standardly held that a belief that *p* is epistemically rational only if supported by evidence that corroborates the truth of *p*. According to some views, however, an assumption that *q* may be epistemically rational even if it is not supported by evidence that corroborates the truth of *q*, if *q* is constitutive of the possibility of providing evidence that corroborates the truth of *p*, where *p* is a content of a belief. (See also *Constitutivism*.)

Epistemic Relativism: Epistemic relativism is the view that different epistemic systems can return opposite and equally valid verdicts concerning the justifiedness or knowledge of one and the same proposition. (See also *Epistemological Contextualism*.)

Epistemic Vertigo: Epistemic vertigo is an anxiety about one's overall epistemic position that arises even once one is in possession of an intellectually satisfying response to radical skepticism. As such, it is meant to capture a kind of phobic reaction to one's epistemic situation. Just as it can be natural to feel afraid when high up even though one knows there is no danger, so there might be a natural epistemic anxiety that arises when one engages in philosophical reflection on one's overall epistemic position, even if one knows that the skeptical problem is illusory.

Epistemological Contextualism: Epistemological contextualism is the view that takes the structure of justification not to be a fixed, universal matter. In particular, what is legitimately taken for granted in one context of inquiry may be properly subject to investigation in a different context of inquiry. (See *Epistemic Relativism*.)

Epistemological Disjunctivism: Epistemological disjunctivism states that in paradigm cases of perceptual knowledge, the subject's belief can enjoy rational support that is both reflectively accessible and factive. In particular, the reflectively accessible rational support for the subject's belief that *p* can be that she *sees that p*, where seeing that *p* entails *p*.

Foundationalism: Foundationalism is a view about the structure of epistemic support. It claims that all epistemic support can be traced back to foundational

beliefs that have a special kind of epistemic status such that they don't depend for their epistemic support on any further beliefs that one holds. Classical forms of foundationalism regard these foundational beliefs as self-justifying, such as by being self-evident. (See also *Agrippa's Trilemma*; *Coherentism*; *Infinitism*.)

Gettier Examples: Gettier examples are cases of justified true belief where the belief in question is only true as a matter of luck, and hence doesn't amount to knowledge. As such, these cases are counterexamples to the classical account of knowledge as justified true belief.

Heavyweight Assumptions: These are assumptions about entire areas of discourse. For instance, that there are physical objects, that there other minds, etc. Only by taking for granted such assumption it is then possible to speak of specific physical objects or minds, such as hands, tables, humans, and animals.

Hinge Commitment/Proposition: Hinge propositions or commitments are a central element of hinge epistemology. According to this proposal, due to the later Wittgenstein, all rational evaluation presupposes a prior 'hinge' certainty that enables this rational evaluation to occur, and which is thus not itself subject to rational evaluation. This hinge certainty is manifest in a subject's certainty in particular propositions, called hinge propositions or commitments. Examples include that one has hands (in normal circumstances), that one's name is such-and-such, and that one has never been to the moon, that the Earth has existed for a very long time, or even, on some readings, that there are physical objects. Some contemporary hinge epistemologists prefer to equate hinges with heavyweight assumptions and others distinguish between hinges with a specific content (like those just listed) and an overarching *über hinge* that we cannot possibly be massively mistaken. (See also *Hinge Epistemology*.)

Hinge Epistemology: Hinge epistemology is the proposal, first articulated in Wittgenstein's later work, and nowadays developed in various ways, that our everyday certainties, such as that one has hands (in normal conditions), or that one has never been to the moon, play a fundamental role in our epistemic practices. In particular, they represent the 'hinge' certainty that must stand fast in order for our rational practices of doubting and offering reasons to function and, as such, are not themselves rationally grounded (at least directly). (See also *Hinge Commitment/Proposition*.)

Humean Skeptical Paradox: Contrary to the *Cartesian Skeptical Paradox*, the Humean Skeptical Paradox does not necessarily appeal to radical skeptical scenarios. The problem it raises primarily concerns the apparent *epistemic circularity* involved in arguments designed to produce knowledge (or justified belief) that there is an external world. In particular, the claim is that any form of reasoning we might employ to derive this anti-skeptical conclusion would already presuppose that this conclusion was known or assumed and hence would be epistemic circular. If that's right, then it seems there simply is no way of reasoning our way to the knowledge of the external world. (See also *Epistemic Circularity*.)

Inference to the Best Explanation: Inference to the Best Explanation (or IBE) is an inference that, while not deductively valid, provides a possible explanation of an observed phenomenon. Inasmuch as that explanation satisfies further criteria such as by being the simplest explanation available, it is considered the best explanation of the said phenomenon.

Infinitism: Infinitism is a controversial view about the structure of epistemic support, such that an infinite regress of supporting beliefs can nonetheless provide epistemic support. (See also *Agrippa's Trilemma*; *Coherentism*; *Foundationalism*.)

Iterativity Principle: In its simplest form, the iterativity principle, or the KK Principle as it is sometimes called, is the claim that when one knows, one also knows that one knows.

Justification: A belief is epistemically justified (or rational) if there is *a priori* or *a posteriori* evidence that corroborates its truth. Justifications are defeasible if new and contrary evidence can come in and speak against the truth of the belief; indefeasible otherwise. A belief is pragmatically justified (or rational) if there are considerations that speak to its usefulness and advisability or even inescapability for creatures like us. On some accounts at least some reasons or justifications, like knowledge, are factive: they are such that if a belief is so justified (or rationally grounded), then the belief is true.

K-Apt Belief: Knowledge-apt, or K-apt, belief is concerned with the particular type of propositional attitude that is a constituent part of rationally grounded knowledge. The thought is that this is a more specific notion of belief than our folk conception of belief, which often covers a range of quite different propositional attitudes. Accordingly, someone could believe a proposition in the folk sense of the term and yet nonetheless lack a K-apt belief in this proposition, and so not be in a position to know this proposition.

Liberalism: According to liberalism, in order to have perceptual justification for a proposition P about physical objects in one's surroundings, one only needs an appropriate experience absent defeaters. (See also *Conservatism*; *Moderatism*.)

Moderatism: According to moderatism, in order to have perceptual justification for a proposition P about physical objects in one's surroundings, one needs no independent justification for assuming the proposition that there is an external world (among other *heavyweight assumptions*), besides having an appropriate course of experience absent defeaters. (See also *Conservatism*; *Heavyweight Assumptions*; *Moderatism*.)

Overriding/Undercutting Responses to Paradoxes: Genuine philosophical paradoxes expose real and fundamental tensions in our ordinary ways of thinking about a certain domain, such as free will or (in the case of radical skepticism) knowledge. As a result, the only way of resolving a genuine paradox is to substantially revise one's ordinary ways of thinking about that domain: this is an *overriding* response to a paradox. Alternatively, if one can show

that the paradox is illusory, in that it in fact trades on contentious theoretical claims that are merely masquerading as being part of our ordinary ways of thinking about the target domain, then one wouldn't be obliged to advance any kind of revisionism in order to dispose of the paradox. This would be an *undercutting* response to a paradox.

Paradox: A paradox is an argument that starts with highly intuitive premises, and employs a seemingly correct form of reasoning, and yet arrives at conclusions that are at odds with our widely shared beliefs. A paradox therefore manifests a deep tension within our ordinary concepts within a certain domain, such that our ordinary commitments in that domain appear to be in conflict with each other. Examples of philosophical paradoxes include the problem of free will and radical skepticism.

Propositional/Doxastic Justification: A propositional justification is the justification there is for a proposition *p* in the abstract space of reasons. Subjects who then believe *p* based on such a justification then possess a doxastic justification for it. (See also *Rationally Unavailable Justification*.)

Pyrrhonian Skepticism: Pyrrhonian skepticism is an ancient form of skepticism which employs skeptical techniques (or 'modes') to motivate a thoroughgoing agnosticism. This particular form of skepticism was an ethical position, in that the guiding idea was that by avoiding commitments as much as possible, one thereby lived an intellectually tranquil life.

Radical Skeptical Hypotheses: A radical skeptical hypothesis is a logically and metaphysically possible scenario that is indistinguishable from ordinary life but where most of one's beliefs are false. A popular example of a radical skeptical hypothesis is the possibility that one is a disembodied brain-in-a-vat (BIV) being 'fed' deceptive experiences about the world. (See also *Brain-in-a-Vat Hypothesis*.)

Radical Skepticism: Radical skepticism is a form of skepticism that calls our beliefs into question *en masse*. In particular, it challenges the possibility that we may know the very existence of physical objects in our surroundings, including the existence of our own bodies.

Rationally Unavailable Justification: A rationally unavailable justification is a propositional justification that cannot be converted to a doxastic justification in one's particular case due to one's further beliefs. For instance, there may be a perceptual propositional justification for the proposition that one has a hand such that, if a subject formed that belief on that basis, it would be doxastically justified. Yet, due to my antecedent belief that I have been drugged, I may not rationally avail myself to that justification. (See also *Propositional/Doxastic Justification*.)

Relativism: The view that the truth of knowledge attributions of the form 'S knows that p' varies with the either the context of use (i.e., who makes them and the standards they bring to bear on the attribution) or the context of

assessment (i.e., who assesses the truth of the knowledge attribution). (See also *Contextualism*.)

Reliabilism: In its simplest form, reliabilism about justification (or knowledge) holds that a belief is justified (or known) in virtue of being formed via a reliable belief-forming process (i.e., a process that tends to produce true rather than false beliefs). (See also *Virtue Reliabilism*.)

Safety Condition: The safety condition is a condition on knowledge which states that when one knows, one has a true belief that couldn't have easily been false (i.e., insofar as one forms a belief on the same basis in close possible worlds, then that belief continues to be true). (See also *Sensitivity Condition*.)

Sense Data: Sense data are the unstructured and immediate deliverances of our senses, such as colored patches or sensory impressions.

Sensitivity Condition: The Sensitivity condition is a condition on knowledge which states that when one knows, then one has a true belief such that, had what one believed not been true, then one wouldn't have believed it. Expressed in terms of possible worlds, this means that in the closest possible world where what one believes is no longer true, one no longer believes it. (See also *Safety Condition*.)

Transmission of Justification Principle: This is the principle that an argument transmits a justification from its premises to its conclusion, thereby allowing one to acquire a first justification for the conclusion, if (and only if) having a justification for the premises does not presuppose that there is already a justification for its conclusion. (See also *Closure Principle*.)

Truth-Conditions: The conditions that would have to obtain for a sentence of a language to be true.

Underdetermination Principle: The Underdetermination principle holds that in order to have knowledge of a proposition (such as that one has hands), one must have epistemic support for one's belief in that proposition which favors it over known alternatives (such as that one is not a handless brain-in-a-vat).

Virtue Reliabilism: Virtue reliabilism is the view that knowledge is the product of specific skill-like reliable belief-forming processes that are stable and integrated features of a subject's cognitive character. (See also *Reliabilism*.)

BIBLIOGRAPHY

Ainslie, D.C. (2015). *Hume's True Scepticism*, Oxford: Oxford University Press.

Alfano, M., Greco, J., and Turri, J. (2017). 'Virtue epistemology,' in E. Zalta (Ed.), *Stanford Encyclopedia of Philosophy*. Available at: https://plato.stanford.edu/entries/epistemology-virtue/

Alston, W. (1989). *Epistemic Justification*, Ithaca, NY: Cornell University Press.

Ashton, N. (2021). 'Extended rationality and epistemic relativism,' in N. J. L. L. Pedersen, and L. Moretti (Eds.), *Non-Evidentialist Epistemology*, Leiden: Brill, pp. 55–72.

Audi, R. (1993). *The Structure of Justification*, Cambridge: Cambridge University Press.

Avnur, Y. (2017). 'On what does rationality hinge?,' *International Journal for the Study of Skepticism* 7, 246–257.

Baghramian, M. (2017). 'Comments on Annalisa Coliva's *Extended Rationality: A Hinge Epistemology*,' *International Journal for the Study of Skepticism* 7, 272–280.

Baghramian, M. and Coliva, A. (2020). *Relativism*, London: Routledge.

Beckwith, C. I. (2015). *Greek Buddha: Pyrrho's Encounter with Early Buddhism in Central Asia*, Princeton, NJ: Princeton University Press.

Beebe, J. (2009). 'The abductivist reply to skepticism,' *Philosophy and Phenomenological Research* 79(3), 605–636.

Beebe, J. (2011). 'A priori skepticism,' *Philosophy and Phenomenological Research* 83(3), 583–602.

Beebee, H. (2010). 'David Hume,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 730–740.

Bell, M. and McGinn, M. (1990). 'Naturalism and scepticism,' *Philosophy* 65(254), 390–418.

Bett, R. (2010). 'Pyrrhonian scepticism,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 403–413.

Boghossian, P. (1989). 'Content and self-knowledge,' *Philosophical Topics* 17, 5–26.

Boncompagni, A. (2016). *Wittgenstein and Pragmatism. On Certainty in Light of Peirce and James*, Basingstoke: Palgrave Macmillan.

BonJour, L. (1985). *The Structure of Empirical Knowledge*, Cambridge, MA: Harvard University Press.

BonJour, L. (2003). 'A version of internalist foundationalism,' in L. BonJour and E. Sosa (Eds.), *Epistemic Justification: Internalism vs. Externalism, Foundations vs. Virtues*, Malden, MA: Blackwell, pp. 3–96.

BonJour, L. and Sosa, E. (Eds.) (2003). *Epistemic Justification: Internalism vs. Externalism, Foundations vs. Virtues*, Malden, MA, Blackwell.

Brady, M. and Pritchard, D. (2005). 'Epistemological contextualism: Problems and prospects,' *The Philosophical Quarterly* 55(219), 161–171.

Brewer, B. (2000). *Perception and Reason*, Oxford: Oxford University Press.

Brewer, B. (2011). *Perception and its Objects*, Oxford: Oxford University Press.

Brueckner, A. (1994). 'The structure of the skeptical argument,' *Philosophy and Phenomenological Research* 54, 827–835.

Brueckner, A. (2010). 'Skepticism and semantic externalism,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 500–510.

Burge, T. (1986). 'Cartesian error and the objectivity of perception,' in J. McDowell and P. Pettit (Eds.), *Subject, Thought, and Context*, Oxford: Oxford University Press.

Burge, T. (1988). 'Individualism and self-knowledge,' *Journal of Philosophy* 85, 649–663.

Burge, T. (1993). 'Content preservation,' *Philosophical Review* 102, 457–488.

Burge, T. (2003). 'Perceptual entitlement,' *Philosophy and Phenomenological Research* 67, 503–548.

Burge, T. (2007). 'Descartes, bare concepts, and anti-individualism: Reply to Normore,' in M. Hahn and B. Romberg (Eds.), *Reflections and Replies: Essays on the Philosophy of Tyler Burge*, Cambridge, MA: MIT Press, pp. 291–334.

Burge, T. (2010). *Origins of Objectivity*, Oxford: Oxford University Press.

Burnyeat, M. (1982). 'Idealism and Greek philosophy: What Descartes saw and Berkeley missed,' *Philosophical Review* 91, 3–40.

Callanan, J. (2011). 'Making sense of doubt: Strawson's anti-scepticism,' *Theoria* 77(3), 261–278.

Cassam, Q. (2007). *The Possibility of Knowledge*, Oxford: Oxford University Press.

Cavell, S. (1979). *The Claim of Reason: Wittgenstein, Skepticism, Morality, and Tragedy*, Cambridge, MA: Harvard University Press.

Chalmers, D. (2005). 'The matrix as metaphysics,' in C. Grau (Ed.), *Philosophers Explore the Matrix*, Oxford: Oxford University Press, pp. 132–176.

Chalmers, D. (2017). 'The virtual and the real,' *Disputatio* 9, 309–352.

Chalmers, D. (2018). 'Structuralism as a response to skepticism,' *Journal of Philosophy* 115, 625–660.

Chisholm, R. M. (1977). *Theory of Knowledge* (2nd edn.), Englewood Cliffs, NJ: Prentice-Hall.

Clarke, T. (1972). 'The legacy of skepticism,' *Journal of Philosophy* 69(20), 754–769.

Cohen, S. (1984). 'Justification and truth,' *Philosophical Studies* 46, 279–295.

Cohen, S. (1998a). 'Contextualist solutions to epistemological problems: Skepticism, Gettier, and the lottery,' *Australasian Journal of Philosophy*, 76(2), 289–306.

Cohen, S. (1998b). 'Two kinds of sceptical argument,' *Philosophy and Phenomenological Research* 58, 143–159.

Cohen, S. (1999). 'Contextualism, skepticism, and the structure of reasons,' in J. Tomberlin (Ed.), *Philosophical Perspectives 13: Epistemology*, Chichester: John Wiley & Sons, Ltd, pp 57–89.

Cohen, S. (2002). 'Basic knowledge and the problem of easy knowledge,' *Philosophy and Phenomenological Research* 65(2), 309–329.

Cohen, S. (2005a). 'Is knowledge contextual? Contextualism defended,' in M. Steup and E. Sosa (Eds.), *Contemporary Debates in Epistemology*, Oxford: Blackwell, pp. 56–71.

Cohen, S. (2005b). 'Why basic knowledge is easy knowledge,' *Philosophy and Phenomenological Research* 70(2), 417–430.

Coliva, A. (2010a). *Moore and Wittgenstein: Scepticism, Certainty, and Common Sense*, Basingstoke: Palgrave Macmillan.

Coliva, A. (2010b). 'Was Wittgenstein an epistemic relativist?,' *Philosophical Investigations* 33(1), 1–23. Reprinted with additions in Coliva (2022).

Coliva, A. (2012). 'Moore's proof, liberals and conservatives. Is there a (Wittgensteinian) third way?,' in A. Coliva (Ed.), *Mind, Meaning and Knowledge: Themes from the Philosophy of Crispin Wright*, Oxford: Oxford University Press, pp. 323–351.

Coliva, A. (2015). *Extended Rationality: A Hinge Epistemology*, London: Palgrave Macmillan.

Coliva, A. (2017a). 'Précis of *Extended Rationality: A Hinge Epistemology*,' *International Journal for the Study of Skepticism* 7, 217–234. Reprinted in Coliva (2022).

Coliva, A. (2017b). 'Replies to commentators,' *International Journal for the Study of Skepticism* 7, 281–295.

Coliva, A. (2018). 'What anti-realism about hinges could possibly be,' in C. Kyriacou and R. McKenna (Eds.), *Metaepistemology: Realism and Anti-Realism*, London: Palgrave, pp. 267–288. Reprinted with additions in Coliva (2022).

Coliva, A. (2020a). 'Relativism and hinge epistemology,' in M. Kusch (Ed.), *Routledge Handbook of Philosophy of Relativism*, London: Routledge, pp. 320–328. Reprinted in Coliva (2022).

Coliva, A. (2020b). 'Against neo-Wittgensteinian entitlements,' in P. J. Graham and N. J. L. L. Pedersen (Eds.), *Epistemic Entitlement*, Oxford: Oxford University Press, pp. 327–343. Abridged version reprinted in Coliva (2022).

Coliva, A. (2021a). 'I Know, I *k*now, I k̲now. Hinge epistemology, invariantism and skepticism,' in C. Kyriacou and K. Wallbridge (Eds.), *Skeptical Invariantism Reconsidered*, London: Routledge, pp. 213–234. Reprinted in Coliva (2022).

Coliva, A. (2021b). 'Hinges, radical skepticism, relativism and alethic pluralism,' in N. J. L. L. Pedersen and L. Moretti (Eds.), *Non-evidentialist*

Epistemology, Leiden: Brill, pp. 97–117. Abridged version reprinted in Coliva (2022).

Coliva, A. (2021c). 'Propositional and doxastic hinge assumptions,' in P. Silva and L. Oliveira (Eds.), *Propositional and Doxastic Justification*, Oxford: Oxford University Press, forthcoming. Reprinted in Coliva (2022).

Coliva, A. (2021d). 'Strange bedfellows: On Pritchard's disjunctivist hinge epistemology,' *Synthese* 198, 3521–3532. Abridged version reprinted in Coliva (2022).

Coliva, A. (2021e). 'Which hinge epistemology between animal, biscopic and constitutivist?,' in C. Sandis and D. Moyal-Sharrock (Eds.), *Extending Hinge Epistemology*, London: Anthem. Reprinted in Coliva (2022).

Coliva, A. (2022). *Wittgenstein Rehinged: The Significance of Wittgenstein's* On Certainty *for Contemporary Epistemology*, London: Anthem, forthcoming.

Coliva, A. and Palmira, M. (2021). 'Disagreement unhinged, constitutivism-style,' *Metaphilosophy* 52(3–4), 1–14.

Conee, E. (2005). 'Is knowledge contextual? Contextualism contested,' in M. Steup and E. Sosa (Eds.), *Contemporary Debates in Epistemology*, Oxford: Blackwell, pp. 47–56.

Conee, E. and Feldman, R. (2004). *Evidentialism*, Oxford: Oxford University Press.

Daniels, N. (1974). *Thomas Reid's 'Inquiry*,' Stanford, CA: Stanford University Press.

Davidson, D. (1975). 'Thought and talk,' in S. Guttenplan (Ed.), *Mind and Language*, Oxford: Oxford University Press, pp. 7–23; and reprinted in Davidson (1980), essay 11.

Davidson, D. (1977). 'The method of truth in metaphysics,' *Midwest Studies in Philosophy* 2, 244–254; and reprinted in Davidson (1980), essay 14.

Davidson, D. (1980). *Inquiries into Truth and Interpretation*, Oxford: Clarendon Press.

Davidson, D. (1983). 'A coherence theory of truth and knowledge,' in D. Henrich (Ed.), *Kant oder Hegel?*, Stuttgart: Klett-Cotta, pp. 428–438, and reprinted in Davidson (2001), Chapter 10.

Davidson, D. (1990). 'Afterthoughts, 1987,' in A. Malichowski (Ed.), *Reading Rorty*, Oxford: Blackwell, pp. 136–168.

Davidson, D. (1991). 'Three varieties of knowledge,' in A. P. Griffiths (Ed.), *A. J. Ayer Memorial Essays*, Cambridge: Cambridge University Press, pp. 153–166, and reprinted in Davidson (2001), Chapter 13.

Davidson, D. (2001). *Subjective, Intersubjective, Objective*, Oxford: Oxford University Press.

de Bary, P. (2002). *Thomas Reid and Scepticism*, London: Routledge.

DePaul, M. (2010). 'Foundationalism,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 235–244.

De Pierris, G. (2015). *Ideas, Evidence, and Method: Hume's Skepticism and Naturalism concerning Knowledge and Causation*, Oxford: Oxford University Press.

DeRose, K. (1993). 'Review of Michael Williams, *Unnatural Doubts: Epistemological Realism and the Basis of Scepticism*,' *The* Philosophical Review 102(4), 604–607.

DeRose, K. (1995). 'Solving the skeptical problem,' *The Philosophical Review* 104(1), 1–52.

DeRose, K. (1999). 'Contextualism: An explanation and defence,' in J. Greco and E. Sosa (Eds.), *Epistemology*, Oxford: Blackwell, pp. 187–205.

DeRose, K. (2005). 'The ordinary language basis for contextualism and the new invariantism,' *The Philosophical Quarterly*, 55(219), 172–198.

DeRose, K. (2006). '"Bamboozled by our own words": Semantic blindness and some arguments against contextualism,' *Philosophy and Phenomenological Research* 73(2), 316–338.

Descartes, R. ([1641–1642] 1996) *Descartes: Meditations on First Philosophy: With Selections from the Objections and Replies*, (Ed.)J. Cottingham, Cambridge: Cambridge University Press.

Dretske, F. (1970). 'Epistemic operators,' *Journal of Philosophy* 67, 1007–1023.

Dretske, F. (1971). 'Conclusive reasons,' *Australasian Journal of Philosophy*, 49, 1–22.

Dretske, F. (2005). 'Is knowledge closed under known entailment? The case against closure,' in M. Steup and E. Sosa (Eds.), *Contemporary Debates in Epistemology*, Oxford: Blackwell, pp. 13–26.

Dutant, J. (2015). 'The legend of the justified true belief analysis,' *Philosophical Issues* 29, 95–145.

Elstein, D. Y. and Jenkins, C. (2020). 'The truth fairy and the indirect epistemic consequentialist,' in P. J. Graham and N. J. L. L. Pedersen (Eds.), *Epistemic Entitlement*, Oxford: Oxford University Press, pp. 344–360.

Feldman, R. (2001). 'Skeptical problems, contextualist solutions,' *Philosophical Studies* 103, 61–85.

Fogelin, R. (1985). *Hume's Skepticism in the Treatise of Human Nature*, London: Routledge.

Fumerton, R. (2010). 'Skepticism and epistemic externalism,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 488–499.

Garrett, D. (1997). *Cognition and Commitment in Hume's Philosophy*, Oxford: Oxford University Press.

Gaukroger, S. (2010). 'René Descartes,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 678–686.

Gendler, T. S. (2008). 'Alief and belief,' *Journal of Philosophy* 105, 634–663.

Genova, A. C. (1999). 'The very idea of massive truth,' in L. E. Hahn (Ed.), *The Philosophy of Donald Davidson*, Chicago: Open Court, pp. 167–191.

Gettier, E. (1963). 'Is justified true belief knowledge?,' *Analysis* 23, 121–123.

Goldman, A. I. (1979). 'What is justified belief?,' in G. Pappas (Ed.), *Justification and Knowledge*, Dordrecht: D. Reidel, pp. 1–23.

Goldman, A. I. (1986). *Epistemology and Cognition*, Cambridge MA: Harvard University Press.

Goldman, A. I. (1988). *Empirical Knowledge*. Berkeley, CA: University of California Press.

Goldman, A. I. (1992). 'Epistemic folkways and scientific epistemology,' in A. I. Goldman, *Liaisons: Philosophy Meets the Cognitive and Social Sciences*, Cambridge, MA: MIT Press, pp. 155–178.

Gopnik, A. (2009). 'Could David Hume have known about Buddhism? Charles François Dolu, the Royal College of La Flèche, and the Global Jesuit Intellectual Network,' *Hume Studies* 35(1&2), 5–28.

Greco, J. (1999). 'Agent reliabilism,' *Philosophical Perspectives* 13, 273–296.

Greco, J. (2000). *Putting Skeptics in Their Place: The Nature of Skeptical Arguments and Their Role in Philosophical Inquiry*, Cambridge: Cambridge University Press.

Greco, J. (2008). 'What's wrong with contextualism?,' *The Philosophical Quarterly* 58(232), 416–436.

Greco, J. (2010). *Achieving Knowledge: A Virtue-Theoretic Account of Epistemic Normativity*, Cambridge: Cambridge University Press.

Greenough. P. (2020). 'Knowledge for nothing,' in P. J. Graham and N. J. L. L. Pedersen (Eds.), *Epistemic Entitlement*, Oxford: Oxford University Press, pp. 361–388.

Harman, G. and Sherman, B. (2011). 'Knowledge and assumptions,' *Philosophical Studies* 156(1), 131–140.

Hawthorne, J. (2005). 'The case for closure,' in E. Sosa and M. Steup (Eds.), *Contemporary Debates in Epistemology*, Oxford: Blackwell, pp. 26–43.

Heller, M. (1999). 'Relevant alternatives and closure,' *Australasian Journal of Philosophy*, 77(2), 196–208.

Hinton, J. M. (1967a). 'Visual experiences,' *Mind* 76, 217–227.

Hinton, J. M. (1967b). 'Experiences,' *Philosophical Quarterly* 17, 1–13.

Hinton, J. M. (1973). *Experiences: An Inquiry into Some Ambiguities*, Oxford: Clarendon Press.

Huemer, M. (2001). *Skepticism and the Veil of Perception*, Lanham, MD: Rowman & Littlefield.

Hume, D. (1739). *A Treatise of Human Nature*, Oxford: Clarendon Press.

Ichikawa, J. J. (2011). 'Quantifiers, knowledge, and counterfactuals,' *Philosophy and Phenomenological Research*, 82(2), 287–313.

Jenkins, C. (2007). 'Entitlement and rationality,' *Synthese* 157(1), 25–45.

Klein, P. (2010). 'Infinitism,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 245–256.

Kompa, N. (2002). 'The context sensitivity of knowledge ascriptions,' *Grazer Philosophische Studien* 64, 11–18.

Kornblith, H. (2000). 'The contextualist evasion of epistemology,' *Philosophical Issues* 10, 24–32.

Kripke, S. A. (2011). 'Nozick on knowledge,' in S. A. Kripke, *Philosophical Troubles: Collected Papers*, vol. 1, Oxford: Oxford University Press, pp. 162–224.

Kusch, M. (2016a). 'Wittgenstein on mathematics and certainty,' in A. Coliva and D. Moyal-Sharrock (Eds.), *Hinge Epistemology*, Leiden: Brill, pp. 48–71.

Kusch, M. (2016b). 'Wittgenstein's *On Certainty* and relativism,' in S. Rinofner-Kreidl and H. A. Wiltsche (Eds.), *Analytical and Continental Philosophy: Methods and Perspectives*, Berlin: De Gruyter, pp. 35–58.

Kuzminski, A. (2010). *Pyrrhonism: How the Ancient Greeks Reinvented Buddhism*, Lanham, MD: Lexington Books.

Kvanvig, J. (2010). 'Virtue epistemology,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 199–207.

Lehrer, K. and Cohen, S. (1983). 'Justification, truth, and coherence,' *Synthese* 55, 191–207.

Le Morvan, P. (2017). 'Knowledge before Gettier,' *British Journal for the History of Philosophy* 25, 1216–1238.

Lewis, D. (1996). 'Elusive knowledge,' *Australasian Journal of Philosophy* 74, 549–567.

Littlejohn, C. (2009). 'The new evil demon problem,' in B. Dowden and J. Fieser, (Eds.), *Internet Encyclopaedia of Philosophy*. Available at: www.iep.utm.edu/evil-new/.

Ludwig, K. (1992). 'Skepticism and interpretation,' *Philosophy and Phenomenological Research* 52, 317–339.

Luper, S. (2010). 'Cartesian scepticism,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 414–424.

MacFarlane, J. (2005). 'The assessment sensitivity of knowledge attributions,' in T. S. Gendler and J. Hawthorne (Eds.), *Oxford Studies in Epistemology*, vol. 1, Oxford: Oxford University Press, pp. 197–233.

MacFarlane, J. (2014). *Assessment Sensitivity: Relative Truth and Its Applications*, Oxford: Oxford University Press.

Machuca, D. and Reed, B. (Eds.) (2018). *Skepticism: From Antiquity to the Present*, London: Bloomsbury.

Malcolm, N. (1942). 'Moore and ordinary language,' in P. A. Schillp (Ed.), *The Philosophy of G. E. Moore*, Evanston, IL: Northwestern University Press, pp. 345–368.

Malcolm, N. (1949). 'Defending common sense,' *Philosophical Review* 58(3), 201–220.

Malcolm, N. (1977). 'Moore and Wittgenstein on the sense of "I Know",' in N. Malcolm, *Thought and Language*, Ithaca, NY: Cornell University Press, pp. 170–198.

Martin, M. G. F. (1997). 'The reality of appearances,' in M. Sainsbury (Ed.), *Thought and Ontology*, Milan: Franco Angeli, pp. 77–96.

Martin, M. G. F. (1998). 'Setting things before the mind,' in A. O'Hear (Ed.), *Contemporary Issues in the Philosophy of Mind*, Cambridge: Cambridge University Press, pp. 157–179.

Martin, M. G. F. (2002). 'The transparency of experience,' *Mind and Language* 17, 376–425.

McDowell, J. (1982). 'Criteria, defeasibility and knowledge,' *Proceedings of the British Academy* 68, 455–479.

McDowell, J. (1986). 'Singular thought and the extent of inner space,' in P. Pettit and J. McDowell (Eds.), *Subject, Thought and Context*, Oxford: Clarendon Press, pp. 137–168.

McDowell, J. (1994). *Mind and World*, Cambridge, MA: Harvard University Press.

McDowell, J. (1995). 'Knowledge and the internal,' *Philosophy and Phenomenological Research* 55, 877–893.

McDowell, J. (2009). 'Wittgenstein's 'quietism',' *Common Knowledge* 15, 365–372.

McGinn, M. (1989). *Sense and Certainty: A Dissolution of Scepticism*, Oxford: Blackwell.

McKinsey, M. (1991). 'Anti-individualism and privileged access,' *Analysis*, 51, 9–16.

Mercer, C. (2017). 'Descartes' debt to Teresa of Ávila, or why we should work on women in the history of philosophy,' Philosophical Studies 174(10), 2539–2555.

Millar, A. (2017). 'On extended rationality,' *International Journal for the Study of Skepticism* 7, 235–245.

Moore, G. E. (1925). 'A defence of common sense,' in J. H. Muirhead (Ed.), *Contemporary British Philosophy* (2nd series), London: Allen & Unwin.

Moore, G. E. (1939). 'Proof of an external world,' *Proceedings of the British Academy* 25, 273–300.

Moore, G. E. (1949). 'Letter to Malcolm,' in G. E. Moore, *Selected Writings*, New York: Routledge, pp. 213–216.

Moore, G. E. (1959a).'Certainty,' in G. E. Moore, *Philosophical Papers*, London: George Allen & Unwin, pp. 227–251.

Moore, G. E. (1959b). 'Four forms of skepticism,' in G. E. Moore, *Philosophical Papers*, London: George Allen & Unwin, pp. 196–226.

Moruzzi, S. (2021). 'Hinge epistemology and alethic pluralism,' in N. J. L. L. Pedersen and L. Moretti (Eds.), *Non-Evidentialist Epistemology*, Leiden: Brill, pp. 73–96.

Moyal-Sharrock, D. (2005). *Understanding Wittgenstein's On Certainty*, London: Palgrave Macmillan.

Moyal-Sharrock, D. (forthcoming). 'Wittgenstein's hinge certainty,' in D. Moyal-Sharrock and C. Sandis (Eds.), *Hinge Epistemology*, London: Anthem.

Moyal-Sharrock, D. and Brenner, W. H. (Eds.) (2005). *Readings of Wittgenstein's On Certainty*, London: Palgrave.

Neta, R., and Pritchard, D. H. (2007). 'McDowell and the new evil genius,' *Philosophy and Phenomenological Research* 74, 381–396.

Normore, C. (2007). 'Burge, Descartes, and us,' in M. Hahn and B. Romberg (Eds.), *Reflections and Replies: Essays on the Philosophy of Tyler Burge*, Cambridge, MA: MIT Press, pp. 1–14.

Nozick, R. (1981). *Philosophical Explanations*, Cambridge, MA: Harvard University Press.

Nuccetelli, S. (Ed.) (2003). *New Essays on Semantic Externalism and Self-Knowledge*, Cambridge, MA: MIT Press.

Olsson, E. (2010). 'Coherentism,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 257–268,

Owen, D. (1999). *Hume's Reason*, Oxford: Oxford University Press.

Pappas, G. S. (1989). 'Sensation and perception in Reid,' *Nous* 23, 155–167.

Pappas, G. (2014). 'Internalist vs. externalist conceptions of epistemic justification,' in E. Zalta (Ed.), *Stanford Encyclopedia of Philosophy*. Available at: https://plato.stanford.edu/entries/justep-intext/

Peacocke, C. (2003). *The Realm of Reason*, Oxford: Oxford University Press.

Pedersen, N. J. L. L. (2020). 'Pluralist consequentialist anti-scepticism,' in P. J. Graham and N. J. L. L. Pedersen (Eds.), *Epistemic Entitlement*, Oxford: Oxford University Press, pp. 297–326.

Plantinga, A. (1993). *Warrant and Proper Function*, Oxford: Oxford University Press.

Pollock, J. (1974). *Knowledge and Justification*, Princeton, NJ: Princeton University Press.

Popkin, R. (2003). *The History of Scepticism: From Savonarola to Bayle*, Oxford: Oxford University Press.

Pritchard, D. (2005a). *Epistemic Luck*, Oxford: Oxford University Press.

Pritchard, D. (2005b). 'Contextualism, scepticism and warranted assertibility manœuvres,' in J. Keim-Campbell, M. O'Rourke, and H. Silverstein (Eds.), *Knowledge and Skepticism*, Cambridge, MA: MIT Press, pp. 85–104.

Pritchard, D. H. (2005c). 'The structure of sceptical arguments,' *Philosophical Quarterly* 55, 37–52.

Pritchard, D. H. (2005d). 'Wittgenstein's *On Certainty* and contemporary anti-skepticism,' in D. Moyal-Sharrock and W. H. Brenner (Eds.), *Readings of Wittgenstein's On Certainty*, London: Palgrave, pp. 189–224.

Pritchard, D. H. (2008). 'McDowellian neo-Mooreanism,' in A. Haddock and F. Macpherson (Eds.), *Disjunctivism: Perception, Action, Knowledge*, Oxford: Oxford University Press, pp. 283–310.

Pritchard, D. H. (2010). 'Relevant alternatives, perceptual knowledge and discrimination,' *Noûs* 44, 245–268.

Pritchard, D. H. (2011). 'Epistemic relativism, epistemic incommensurability, and Wittgensteinian epistemology,' in S. Hales (Ed.), *A Companion to Relativism*, Oxford: Blackwell, pp. 266–285.

Pritchard, D. H. (2012). *Epistemological Disjunctivism*, Oxford: Oxford University Press.

Pritchard, D. H. (2013). 'Davidson on radical skepticism,' in E. LePore and K. Ludwig (Eds.), *Blackwell Companion to Donald Davidson*, Oxford: Blackwell, pp. 521–533.

Pritchard, D. H. (2014a). 'Sceptical intuitions,' in D. Rowbottom and T. Booth (Eds.), *Intuitions*, Oxford: Oxford University Press, pp. 213–231.

Pritchard, D. H. (2014b). 'Entitlement and the groundlessness of our believing,' in D. Dodd and E. Zardini (Eds.), *Contemporary Perspectives on Scepticism and Perceptual Justification*, Oxford: Oxford University Press, pp. 190–213.

Pritchard, D. H. (2015). *Epistemic Angst: Radical Skepticism and the Groundlessness of Our Believing*, Princeton, NJ: Princeton University Press.

Pritchard, D. H. (2017a). 'Skepticism and disjunctivism,' in D. Machuca and B. Reed (Eds.), *Skepticism: From Antiquity to the Present*, London: Continuum, pp. 652–667.

Pritchard, D. H. (2017b). 'Wittgenstein on hinge commitments and radical scepticism in *On Certainty*,' in H.-J. Glock and J. Hyman (Eds.), *Blackwell Companion to Wittgenstein*, Oxford: Blackwell, pp. 563–575.

Pritchard, D. H. (2018a). 'Epistemic *Angst*,' *Philosophy and Phenomenological Research* 96, 70–90.

Pritchard, D. H. (2018b). 'Contextualism and radical scepticism,' *Synthese* 195, 4733–4750.

Pritchard, D. H. (2018c). 'Wittgensteinian hinge epistemology and deep disagreement,' *TOPOI*, [Online First]. DOI: 10.1007/s11245-018-9612-y.

Pritchard, D. H. (2019a). *Skepticism: A Very Short Introduction*, Oxford: Oxford University Press.

Pritchard, D. H. (2019b). 'Unnatural doubts,' in G. Anthony Bruno and A. C. Rutherford (Eds.), *Skepticism: Historical and Contemporary Inquiries*, London: Routledge.

Pritchard, D. H. (2020). 'Epistemic vertigo,' in B. Brogaard and D. Gatzia (Eds.), *The Philosophy and Psychology of Ambivalence: Being in Two Minds*, London: Routledge.

Pritchard, D. H. (2021a). 'Scepticism and certainty: Moore and Wittgenstein on commonsense and philosophy,' in R. Peels and R. van Woudenberg (Eds.), *Cambridge Companion to Common Sense Philosophy*, Cambridge: Cambridge University Press, pp. 247–264.

Pritchard, D. H. (2021b). 'Scepticism and *Epistemic Angst*, redux,' *Synthese* 198, 3635–3664.

Pritchard, D. H. (forthcoming a). 'Moderate knowledge externalism,' in L. R. G. Oliveira (Ed.), *Externalism About Knowledge*, Oxford: Oxford University Press.

Pritchard, D. H. (forthcoming b). 'Putnam on radical skepticism: Wittgenstein, Cavell, and occasion-sensitive semantics,' in J. Conant and S. Chakraborty (Eds.), *Engaging Putnam*, Berlin: De Gruyter.

Pritchard, D. H. and Ranalli, C. (2016). 'Radical scepticism and brains in vats,' in S. Goldberg (Ed.), *The Brain in a Vat*, Cambridge: Cambridge University Press, pp. 75–89.

Pryor, J. (2000). 'The skeptic and the dogmatist,' *Noûs* 34, 517–549.

Pryor, J. (2004). 'What's wrong with Moore's argument,' *Philosophical Issues* 14, 349–378.

Pryor, J. (2012). 'When warrant transmits,' in A. Coliva (Ed.), *Mind, Meaning and Knowledge: Themes from the Philosophy of Crispin Wright*, Oxford: Oxford University Press, pp. 269–303.

Putnam, H. (1975). 'The Meaning of "Meaning",' *Minnesota Studies in the Philosophy of Science* 7, 131–193.

Putnam, H. (1981). 'Brains in a Vat,' in H. Putnam, *Reason, Truth and History*, Cambridge, Cambridge University Press, pp. 1–21.

Quine, W. V. O. (1960). *Word and Object*, Cambridge, MA: MIT Press.

Reid, T. ([1764] 1997). *An Inquiry into the Human Mind on the Principles of Common Sense*, (Ed.) Derek R. Brookes, University Park, PA: Pennsylvania State University Press.

Reid, T. ([1785] 2002). *Essays on the Intellectual Powers of Man*, (Ed.) Derek R. Brookes, University Park, PA: Pennsylvania State University Press.

Russell, B. (1912). *The Problems of Philosophy*, Oxford: Oxford University Press.

Russell, B. (1914). *Our Knowledge of the External World*, LaSalle, IL, Open Court.

Rysiew, P. (2001). 'The context-sensitivity of knowledge attributions,' *Noûs* 35(4), 477–514.

Schaffer, J. (2004). 'From contextualism to contrastivism,' *Philosophical Studies* 119(1–2), 73–103.

Schaffer, J. (2005a). 'What shifts? Thresholds, standards, or alternatives?,' in G. Preyer and G. Peter (Eds.), *Contextualism in Philosophy*, Oxford: Oxford University Press, pp. 115–130.

Schaffer, J. (2005b). 'Contrastive knowledge,' in Tamar Szabo Gendler and John Hawthorne (Eds.), *Oxford Studies in Epistemology*, vol. 1, Oxford: Oxford University Press, pp. 235–270.

Schiffer, S. (1996). 'Contextualist solutions to skepticism,' *Proceedings of the Aristotelian Society* 96, 317–333.

Schönbaumsfeld, G. (2016). *The Illusion of Doubt*, Oxford: Oxford University Press.

Schönbaumsfeld, G. (2021). '"Logical" and "epistemic" uses of "to know" or "hinges" as logical enabling conditions,' in C. Kyriacou and K. Wallbridge (Eds.), *Skeptical Invariantism Reconsidered*, London: Routledge, forthcoming.

Smith, M. (2020). 'Full blooded entitlement,' in P. J. Graham, and N. J. L. L. Pedersen (Eds.), *Epistemic Entitlement*, Oxford: Oxford University Press, pp. 281–296.

Snowdon, P. (1980–1981). 'Perception, vision and causation,' *Proceedings of the Aristotelian Society* (new series) 81, 175–192.

Snowdon, P. (1990–1991). 'The objects of perceptual experience,' *Proceedings of the Aristotelian Society* (suppl. vol.) 64, 121–150.

Sosa, E. (1991). *Knowledge in Perspective: Selected Essays in Epistemology*, Cambridge: Cambridge University Press.

Sosa, E. (2000). 'Skepticism and contextualism,' *Philosophical Issues* 11, 1–18.

Sosa, E. (2007). *A Virtue Epistemology: Apt Belief and Reflective Knowledge*, Oxford: Oxford University Press.

Sosa, E. (2009). *Reflective Knowledge: Apt Belief and Reflective Knowledge*, Oxford: Clarendon.

Sosa, E. (2011). *Knowing Full Well*, Princeton, NJ: Princeton University Press.

Sosa, E. (2013). 'Intuitions and foundations: The relevance of Moore and Wittgenstein,' in A. Casullo and J. Thurow (Eds.), *The A Priori in Philosophy*, Oxford: Oxford University Press, pp. 186–200.

Sosa, E. (2015). *Judgment and Agency*, Oxford: Oxford University Press.

Sosa, E. (2021). *Epistemic Explanations: A Theory of Telic Normativity, and What it Explains*, Oxford: Oxford University Press.

Soteriou, M. (2020). 'The disjunctive theory of perception,' in E. Zalta (Ed.), *Stanford Encyclopaedia of Philosophy*. Available at: http://plato.stanford.edu/entries/perception-disjunctive/.

Steup, M. and Sosa, E. (2005). *Contemporary Debates in Epistemology*, Malden, MA: Blackwell.

Stevenson, L. (2002). 'Six levels of mentality,' *Philosophical Explorations* 5, 105–124.

Strawson, P. F. (1985). *Skepticism and Naturalism: Some Varieties*, New York: Columbia University Press.

Stroud, B. (1968). 'Transcendental arguments,' *Journal of Philosophy* 65, 241–256.

Stroud, B. (1977). *Hume*, London: Routledge & Kegan Paul.

Stroud, B. (1984). *The Significance of Philosophical Scepticism*, Oxford: Oxford University Press.

Stroud, B. (1994). 'Skepticism, externalism and the goal of epistemology,' *Aristotelian Society, Supplementary Volume* 68, 291–307.

Vahid, H. (2010). 'Externalism/internalism,' in S. Bernecker and D. H. Pritchard (Eds.), *Routledge Companion to Epistemology*, London: Routledge, pp. 144–155.

Van Cleve, J. (1999). 'Reid on the first principles of contingent truths,' *Reid Studies* 3, 3–30.

Van Woudenberg, R. (2018). 'Thomas Reid's engagement with skeptics and skepticism,' in D. Machuca and B. Reed (Eds.), *Skepticism: From Antiquity to the Present*, London: Bloomsbury, pp. 395–406.

Vogel, J. (1990a). 'Skepticism and inference to the best explanation,' *Journal of Philosophy* 87, 658–666.

Vogel, J. (1990b). 'Are there counterexamples to the closure principle?,' in M. Ross and G. Ross (Eds.), *Doubting: Contemporary Perspectives on Skepticism*, Dordrecht: Kluwer.

Vogel, J. (2002). 'Reliabilism levelled,' *Journal of Philosophy* 97, 602–623.

Vogel, J. (2005). 'The refutation of skepticism,' in M. Steup and E. Sosa (Eds.), *Contemporary Debates in Epistemology*, Malden, MA: Blackwell, pp. 72–84.

Volpe, G. (2017). 'Extended rationality: Some queries about warrant, epistemic closure, truth and skepticism,' *International Journal for the Study of Skepticism* 7, 258–271.

Wedgwood, R. (2011). 'Primitively rational belief-forming practices,' in A. Reisner and A. Steglich-Petersen (Eds.), *Reasons for Belief*, Cambridge, Cambridge University Press, pp. 180–200.

Wedgwood, R. (2013). 'A priori bootstrapping,' in A. Casullo and J. Thurow (Eds.), *The A Priori in Philosophy*, Oxford: Oxford University Press, pp. 226–246.

White, R. (2006). 'Problems for dogmatism,' *Philosophical Studies* 131, 525–557.

Williams, M. (1988). 'Scepticism and charity,' *Ratio* 1, 176–194.

Williams, M. (1991). *Unnatural Doubts: Epistemological Realism and the Basis of Scepticism*, Oxford: Blackwell.

Williams, M. (2004a). 'Wittgenstein's refutation of idealism,' in D. MacManus (Ed.), *Wittgenstein and Scepticism*, New York: Routledge, pp. 76–96.

Williams, M. (2004b). 'Wittgenstein, truth and certainty,' in M. Kölbel and B. Weiss (Eds.), *Wittgenstein's Lasting Significance*, New York: Routledge, pp. 249–284.

Williams, M. (2007). 'Why (Wittgensteinian) contextualism is not relativism,' Episteme 4(1), 93–114.

Williams, M. (2010). 'Descartes's transformation of the skeptical tradition,' in R. Bett (Ed.), *Cambridge Companion to Ancient Scepticism*, Cambridge: Cambridge University Press, pp. 288–313.

Williams, M. (2012). 'Wright against the sceptics,' in A. Coliva (Ed.), *Mind, Meaning and Knowledge: Themes from the Philosophy of Crispin Wright*, Oxford: Oxford University Press, pp. 352–376.

Williams, M. (2015). 'The Agrippan problem, then and now,' *International Journal for the Study of Skepticism* 22, 80–106.

Williamson, T. (2000). 'Scepticism and evidence,' *Philosophy and Phenomenological Research* 60, 613–628.

Williamson, T. (2001a). *Knowledge and its Limits*, Oxford: Oxford University Press.

Williamson, T. (2001b). 'Comments on Michael Williams' "Contextualism, externalism, and epistemic standards",' *Philosophical Studies*, 103, 25–33.

Williamson, T. (2005). 'Contextualism, subject-sensitive invariantism and knowledge of knowledge,' The Philosophical Quarterly 55(219), 213–235.

Wittgenstein, L. (1953). *Philosophical Investigations*, (Eds.)G. E. M. Anscombe and R. Rhees, (trans.) G. E. M. Anscombe, Oxford: Blackwell.

Wittgenstein, L. (1969). *On Certainty*, (Eds.) G. E. M. Anscombe and G. H. von Wright, (trans.) D. Paul and G. E. M. Anscombe, Oxford: Blackwell. Referred to in the text as OC.

Wright, C. (1985). 'Facts and certainty,' *Proceedings of the British Academy* 71, 429–472.

Wright, C. (2002). 'Anti-sceptics simple and subtle: G. E. Moore and John McDowell,' *Philosophy and Phenomenological Research* 65(2), 330–348.

Wright, C. (2004). 'Warrant for nothing (and foundations for free?),' *Aristotelian Society, Supplementary Volume* 78, 167–212.

Wright, C. (2005). 'Contextualism and skepticism: Even-handedness, factivity, and surreptitiously raising the standards,' *Philosophical Quarterly*, 55, 219–236.

Wright, C. (2007). 'The perils of dogmatism,' in S. Nuccetelli and G. Seay (Eds.), *Themes from G. E. Moore: New Essays in Epistemology and Ethics*, Oxford: Oxford University Press, pp. 25–48.

Wright, C. (2012). 'Replies – warrant, transmission and entitlement,' in A. Coliva (Ed.), *Mind, Meaning and Knowledge: Themes from the Philosophy of Crispin Wright*, Oxford: Oxford University Press, pp. 451–486.

Wright, C. (2014). 'On epistemic entitlement (II): Welfare state epistemology,' in D. Dodd and E. Zardini (Eds.), *Scepticism and Perceptual Justification*, Oxford: Oxford University Press, pp. 213–247.

Wright, C. J. G. (1991). 'Scepticism and dreaming: Imploding the demon,' *Mind* 397, 87–115.

Zagzebski, L. T. (1996). *Virtues of the Mind: An Inquiry into the Nature of Virtue and the Ethical Foundations of Knowledge*, Cambridge: Cambridge University Press.

Zalabardo, J. (2012). 'Wright on Moore,' in A. Coliva (Ed.), *Mind, Meaning and Knowledge: Themes from the Philosophy of Crispin Wright*, Oxford: Oxford University Press, pp. 304–322.

Zanetti, L. (2021). 'Inescapable hinges: Steps towards a transcendental hinge epistemology,' in N. J. L. L. Pedersen and L. Moretti (Eds.), *Non-Evidentialist Epistemology*, Leiden: Brill, pp. 27–54.

INDEX